
Pooling and Drift in Delayed Bandits

Melika Baghi

Georgia Institute of Technology
mbaghi3@gatech.edu

Abstract

A system often has to act long before it learns whether the act worked: a recommender sees a click in seconds and a purchase in days. For K actions and a delay of d rounds, methods that see only the action already attain the minimax-optimal dependence on the delay (Zimmert and Seldin, 2020), and with an intermediate state visible the rate is $\tilde{O}(\sqrt{(K+d)T})$ over T rounds (Esposito et al., 2023). Both are governed by the number of actions, so a longer menu is always more expensive to learn from, and neither adapts to instances in which many actions are effectively interchangeable. They need not be. If the outcome depends on the action only through the state it produced, a click or a browse or a cart, then one late outcome informs every action that could have produced the state observed, and the price is set by how many genuinely different states the actions produce. We charge each delayed outcome through the state and pay an *effective dimension* v_t , between 1 and the number of states, in place of K . For the algorithm that is actually run we prove $\tilde{O}(\sqrt{V^-} + \sqrt{dT})$, where V^- is any budget on $\sum_t (v_t - 1)$ fixed in advance, so the action count leaves the leading term and the delay enters additively. The gain has a limit: even when handed the exact losses from d rounds ago, no algorithm escapes $\Omega(\sqrt{d\mathcal{E} \min\{1 + \log J, T/d\}})$, where J counts the drifting directions and $\mathcal{E}$ bounds how far the losses move while the learner waits. On generated data the state channel cuts regret by up to 79 per cent against action-level weighting, and on the funnel family by 32 to 68 per cent against the rate-optimal action-only baseline of Zimmert and Seldin (2020) under the tuning protocol we report. The guarantees assume the action-to-state map is known, and every experiment is synthetic.

1 Introduction

Operations commit before they learn. A markdown is set and sell-through is known a season later; staff are rostered and the shift's outcome lands after it; marketing spend is committed before attribution resolves; an item is recommended and the purchase it may cause is days away. In each, the decision is made now and the quantity it will be judged by arrives late, while something cheaper arrives quickly: a fitting-room visit, a shift's first hour, a click. Whether that early signal can cut the cost of learning is a question between two fields. What may be inferred from the signal is a matter of data; the wait before the outcome resolves is a constraint on the operation, which no estimator removes, and the decision commits while only the first of the two is settled.

The cost of learning under delay should be set by how many genuinely different regimes the actions induce rather than by how many actions there are: a catalogue of thousands of items that funnels users through half a dozen engagement patterns should be no harder to learn from than a catalogue of six.

The mechanism behind that claim is that the outcome depends on the action only through the intermediate state it produced, so one late outcome can be charged to every action that could have




Figure 1: **The claim and its boundary, on generated data.** **Left**, what each state predicts changes while the action-to-state map stays fixed, so losses from d rounds earlier grow inaccurate; E_2 of Section 2 tracks it. The marked points are the true loss at $t - d$ and the same value handed to the learner at t , which the rule spans. **Right**, more overlap allows more sharing, at matched difficulty, $K = 40$, $|\mathcal{S}| = 8$, $T = 10000$, $d = 50$, 20 paired seeds, bars one standard error of the paired difference (Appendix E, Experiment 1).

produced the state observed rather than to the one action played, pooling it across them. Bandits with intermediate observations formalize the setting: the learner picks one action and sees the consequence of that action alone. Esposito et al. (2023) show that the *state-to-loss* map decides the complexity, an unrestricted one giving only the rate available with no intermediate observations; we fix that map and ask what the *action-to-state* structure buys. The delayed-feedback literature has since developed along other axes: Bar-On and Mansour (2025) obtain instance-dependent bounds when the observations themselves evolve, and Levy et al. (2025) treat contextual actions under delay with general function approximation. In neither does an intermediate state carry one delayed outcome across several actions.

We answer that question with an *effective dimension* v_t , measuring how distinguishable the action-induced state distributions are under the current play; it lies between 1 and the number of states and can fall far below K . Theorem 1 pays v_t in place of K for the algorithm that is actually run, and across the families the measured gain tracks v_t rather than K . The claim has a boundary. Waiting carries a cost of its own, which sharing cannot remove. Writing $m_t = c_{t-d}$ for the action-loss vector from d rounds earlier and $E_2 = \sum_t \|c_t - m_t\|_\infty^2$ for how far the losses move while the learner waits, Theorem 2 lower-bounds the regret this forces even when m_t is supplied free (Figure 1). Adapting to overlap and drift at once remains open (Section 5). Four ingredients separate the neighbouring settings, and it is their combination that is new here.

	delayed	state seen	P known	instance-dependent
Zimmert and Seldin (2020)	✓	✗	✗	✗
Esposito et al. (2023)	✓	✓	✗	✗
Eldowa et al. (2024)	✗	✓	✓	✓
Bar-On and Mansour (2025)	✓	✗	✗	✓
this paper	✓	✓	✓	✓

Where this work sits.

Eldowa et al. (2024) is closest, and $v_t - 1$ is exactly the χ^2 mutual information they introduce, so that quantity is theirs; what is new is its role once the outcome is delayed. Other neighbouring settings are placed in Appendix C.

2 Problem formulation

The question is whether the intermediate state can lower the cost of learning below what K alone forces. The model that makes it precise is the following. There are T rounds, K actions, a finite state space $\mathcal{S}$, a known action-to-state matrix P , and a fixed known delay d ; all logarithms are natural and $\tilde{O}$ hides logarithmic factors. Every round follows the chain $A_t \xrightarrow{P} S_t \xrightarrow{\theta_t} X_t$: the learner draws A_t from a distribution x_t , observes $S_t \sim P(\cdot | A_t)$ at once, and receives the delayed outcome only at time $t + d$, encoded as a loss $X_t \in [0, 1]$ so that smaller is better, with $\mathbb{E}[X_t | \mathcal{F}_{t-1}, A_t, S_t] = \theta_t(S_t)$, where $\mathcal{F}_t$ is the environment’s history through round t , including outcomes generated but not yet delivered; outcomes are conditionally independent across rounds given the states. The conditional mean depends on A_t only through S_t . This *state-sufficiency* assumption licenses cross-action pooling: an outcome observed at state S_t may update every action that could

73 have produced that state. If an action affects the outcome directly, in a way not represented in S_t , the
 74 argument no longer applies. The learner may use all states observed so far, but among outcomes only
 75 those from rounds $r \leq t - d - 1$ have been delivered. STATE-EXP3 changes its play distribution
 76 only when an outcome arrives, so for this algorithm x_t is $\mathcal{F}_{t-d-1}$ -measurable.

77 The matrix $P \in [0, 1]^{K \times |\mathcal{S}|}$ is row-stochastic and fixed, and the *state-loss means* $\theta_t \in [0, 1]^{|\mathcal{S}|}$
 78 are *oblivious*: chosen in advance, not reacting to the learner. Action losses are $c_t = P\theta_t$ and
 79 performance is the static *pseudo-regret*, the gap to the best single action in hindsight, $\mathbb{E}[R_T] =$
 80 $\mathbb{E}[\sum_t c_t(A_t)] - \min_a \sum_t c_t(a)$. **Oracle for the lower bound.** For Theorem 2 only, the learner
 81 also receives the *stale action-loss vector* $m_t \triangleq c_{t-d}$ before choosing A_t , with $m_t = c_1$ for $t \leq d$.
 82 This oracle information is unavailable to STATE-EXP3, so a lower bound that holds even when it is
 83 supplied is stronger. We measure how misleading m_t is by $E_2 = \sum_t \|c_t - m_t\|_\infty^2 \leq T$, which stays
 84 small when errors are small even if frequent. The geometry of P governs sharing across actions and
 85 the evolution of θ_t governs staleness across time.

86 3 Pooling through the state

87 These two sources of difficulty separate cleanly. Pooling replaces the leading action-count depen-
 88 dence by an effective dimension, while no amount of sharing removes the cost of waiting.

89 An arriving outcome need not update the action played alone. Every action can be credited by how
 90 likely it was to have produced the observed state. Here P is known and m_t is not given. For each
 91 state $s \in \mathcal{S}$ let $q_t(s) = \sum_{a=1}^K x_t(a)P(s | a)$, and give each action $a \in \{1, \dots, K\}$ the estimate
 92 $\hat{c}_t(a) = P(S_t | a)X_t/q_t(S_t)$. It is correct on average, and its noise is governed by the *effective*
 93 *dimension* $v_t = \sum_s q_t(s)^{-1} \sum_a x_t(a)P(s | a)^2 \in [1, |\mathcal{S}|]$. STATE-EXP3 is EXP3 (Auer et al.,
 94 2002) at learning rate η on these estimates: it adds $\hat{c}_t$ to its running total once the outcome is usable,
 95 at round $t + d + 1$. A round costs $O(K)$ (Appendix A, with pseudocode and the χ^2 reading of v_t).

96 **Theorem 1.** *Let P be known and (θ_t) oblivious, and fix in advance any V^- with $\sum_t (v_t -$
 97 $1) \leq V^-$ always. Then STATE-EXP3 with any $\eta \leq 1/(e(d + 1))$ satisfies $\mathbb{E}[R_T] \leq$
 98 $\sqrt{2(3.3V^- + 2dT)} \log K + e(d + 1) \log K + d$.*

99 Thus the leading action-count dependence is replaced by V^- , which measures how differently the
 100 actions behave rather than how many there are, while the delay enters additively at $\tilde{O}(\sqrt{V^-} + \sqrt{dT})$,
 101 for the algorithm every experiment below runs. Writing $\bar{v}$ for the largest v_t over all ways of playing,
 102 $V^- = (\bar{v} - 1)T$ is admissible and depends on P alone. An interleaved variant that runs $d + 1$ copies,
 103 and a version that merges similar states, are deferred to Appendix B; neither adapts to E_2 .

104 The proof uses one key identity. The denominator of $\hat{c}_t$ is exactly the probability of the observed
 105 state under the current play, so $\langle x_t, \hat{c}_t \rangle = X_t \leq 1$ on every path. This limits how far the play
 106 distribution can move while an outcome is in flight and controls the extra term created by delay;
 107 subtracting $X_t \mathbf{1}$ in the analysis replaces v_t by $v_t - 1$ (Appendix F).

108 **The cost of waiting.** Waiting carries a cost that persists even when the stale losses are supplied
 109 free. Theorem 2 isolates that cost from time alone: the construction gives each of the J drifting
 110 directions a private state, which removes useful pooling and forces $J \leq |\mathcal{S}| - 1$ (Appendix D).

111 **Theorem 2.** *Fix T , a delay $1 \leq d \leq T/2$, a drift budget $\mathcal{E} \leq T/4$, and $1 \leq J \leq \min\{K - 1, |\mathcal{S}| - 1\}$.
 112 Every algorithm, even one handed the exact losses from d rounds ago, meets some instance fixed in
 113 advance with $E_2 \leq \mathcal{E}$ on which $\mathbb{E}[R_T] = \Omega(\sqrt{d\mathcal{E} \min\{1 + \log J, T/d\}})$. Without that help, and if
 114 $K \leq T$, the best regret any algorithm can guarantee is $\Omega(\sqrt{KT} + \sqrt{d\mathcal{E} \min\{1 + \log J, T/d\}})$, up
 115 to constants.*

116 The first term is the usual cost of exploring. The second is the cost of delay that pooling cannot
 117 remove: it grows with the wait, how far the losses move during it, and the number J of directions
 118 the best fixed action can use. They come from different hard instances, so adding them overstates
 119 the truth by at most a factor of two.




Figure 2: **Effective dimension, not action count.** **Left**, regret saved against action-level weighting as the effective dimension falls, with task difficulty held fixed so the sweep is not confounded by easier instances; $K = 40$, $|\mathcal{S}| = 8$, $T = 10000$, $d = 50$, 20 paired seeds, bars one standard error of the paired difference (Experiment 1). **Right**, the catalogue grows thirty-two-fold while $\bar{v}$ stays between 1.4 and 1.8, the separation Theorem 1 charges for (Experiment 2). $\hat{v}$ and $\bar{v}$ estimate $\sup_x v(x)$ by Monte Carlo.

4 Experiments

Theorem 1 predicts that the learning cost follows the number of distinct action behaviours rather than the number of actions, so the families below are built where the two come apart. Figure 2 separates them in both directions: the gain rises as the effective dimension falls at fixed task difficulty, and the effective dimension stays nearly constant as the catalogue grows. With $K = 200$ generated items and six levels of engagement the estimated dimension is 1.97, and STATE-EXP3 cuts regret against action-level weighting by 79.3, 63.6 and 38.9 per cent at $d = 10, 50, 200$, and against the rate-optimal action-only delayed-bandit baseline of Zimmert and Seldin (2020) tuned over a grid by 67.9, 51.4 and 31.5. That setting is hand-structured, not fitted to data, so it shows the mechanism, not a deployed system.

	d	action-level	tuned (Zimmert and Seldin, 2020)	STATE-EXP3	cut vs action	cut vs tuned
130	10	4577	2959	949	79.3%	67.9%
	50	4752	3557	1728	63.6%	51.4%
	200	5307	4735	3245	38.9%	31.5%

Funnel family, $K = 200$, $|\mathcal{S}| = 6$, $T = 20000$, 8 paired seeds; cuts are against action-level weighting and against the tuned baseline (Experiment 2).

These runs use the algorithm Theorem 1 covers. A best-state rule leads seven of nine settings, but where its assumption fails its regret grows linearly in T while ours stays sublinear (Appendix E).

5 Discussion

Two things stand between an action and its worth: how much its behaviour overlaps with other actions, and how long the outcome takes to arrive. The first is exploitable: when P is known, an estimate of $\bar{v}$ well below K indicates that delayed outcomes can be pooled across actions and that the state channel is worth using. The second cannot be pooled away: Theorem 2 shows that uncertainty from waiting persists even when the exact losses from d rounds ago are supplied for free. Thus, under state sufficiency, intermediate states can remove the leading action-count cost of a large catalogue, but not the temporal cost of stale information. Two limitations remain: $\bar{v}$ is estimated rather than computed exactly, and our upper bound captures overlap while the lower bound characterises drift on a separate axis. Adapting to both overlap and drift in one algorithm and one guarantee remains open.

References

N. Alon, N. Cesa-Bianchi, O. Dekel, and T. Koren. Online learning with feedback graphs: beyond bandits. *COLT*, PMLR 40, 2015.

- 149 P. Auer, N. Cesa-Bianchi, Y. Freund, and R. E. Schapire. The nonstochastic multiarmed bandit
150 problem. *SIAM Journal on Computing*, 32(1):48–77, 2002.
- 151 Y. Bar-On and Y. Mansour. Non-stochastic bandits with evolving observations. *ALT*, PMLR 272,
152 2025.
- 153 N. Cesa-Bianchi and G. Lugosi. *Prediction, Learning, and Games*. Cambridge University Press,
154 2006.
- 155 N. Cesa-Bianchi, C. Gentile, and Y. Mansour. Delay and cooperation in nonstochastic bandits. *JMLR*,
156 20, 2019.
- 157 K. Eldowa, N. Cesa-Bianchi, A. M. Metelli, and M. Restelli. Information capacity regret bounds for
158 bandits with mediator feedback. *JMLR*, 25(353), 2024.
- 159 E. Esposito, S. Masoudian, H. Qiu, D. van der Hoeven, N. Cesa-Bianchi, and Y. Seldin. Delayed
160 bandits: When do intermediate observations help? *ICML*, PMLR 202, 2023.
- 161 E. Esposito, F. Fusco, D. van der Hoeven, and N. Cesa-Bianchi. Learning on the edge: online
162 learning with stochastic feedback graphs. *NeurIPS*, 35, 2022.
- 163 G. Gabbianelli, G. Neu, and M. Papini. Online learning with off-policy feedback. *ALT*, PMLR 201,
164 2023.
- 165 O. Levy, L. Erez, A. Cohen, and Y. Mansour. Regret bounds for adversarial contextual bandits with
166 general function approximation and delayed feedback. *arXiv:2510.09127*, 2025.
- 167 O. Schlisselberg, T. Lancewicki, P. Auer, and Y. Mansour. Improved best-of-both-worlds regret for
168 bandits with delayed feedback. *arXiv:2505.24193*, 2025.
- 169 T. S. Thune, N. Cesa-Bianchi, and Y. Seldin. Nonstochastic multiarmed bandits with unrestricted
170 delays. *NeurIPS*, 32, 2019.
- 171 C. Vernade, A. György, and T. Mann. Non-stationary delayed bandits with intermediate observations.
172 *ICML*, PMLR 119, 2020.
- 173 S. Wang, A. Barik, and V. Y. F. Tan. Bandit convex optimization with gradient prediction adaptivity.
174 *arXiv:2605.22191*, 2026.
- 175 K. W. Zhang, T. Baldwin-McDonald, K. Ciosek, L. Maystre, and D. Russo. Impatient bandits: opti-
176 mizing for the long-term without delay. *JMLR*, to appear, 2026.
- 177 J. Zimmert and Y. Seldin. An optimal algorithm for adversarial bandits with arbitrary delays. *AIS-*
178 *TATS*, PMLR 108, 2020.

179 A The state-pooled estimator

180 **The algorithm.** Inputs are the action-to-state matrix P , the delay d , the number of copies m ,
181 and the rate $\eta > 0$. Setting $m = 1$ gives STATE-EXP3, the algorithm of the body and of every
182 experiment, covered by Theorem 1; setting $m = d + 1$ gives a distinct algorithm, INTERLEAVED-
183 STATE-EXP3, which runs $d + 1$ copies in rotation with copy $i \in \{0, \dots, m - 1\}$ taking every m th
184 round, and is covered by Theorem 3. Round r ’s outcome becomes usable at round $r + d + 1$, which
185 for $m = d + 1$ is the next round belonging to the copy that played round r .

- 186 1. Set $L_i \leftarrow (0, \dots, 0) \in \mathbb{R}^K$ for $i = 0, \dots, m - 1$, and let $\mathcal{P}$ be an empty table of rounds
187 awaiting their outcome.
- 188 2. For $t = 1, \dots, T$:
 - 189 (a) *Deliver.* If $r \triangleq t - d - 1 \geq 1$, read $(i_r, q_r(S_r), S_r, X_r)$ from $\mathcal{P}$, remove it, and update
190 only that copy, $L_{i_r}(a) \leftarrow L_{i_r}(a) + P(S_r \mid a)X_r/q_r(S_r)$ for every a . This is the
191 pooled step, and it touches every action rather than the one played.
 - 192 (b) *Select the copy.* $i \leftarrow t \bmod m$.
 - 193 (c) *Play.* $x_t(a) \propto \exp(-\eta L_i(a))$, draw $A_t \sim x_t$, and observe S_t at once.

194 (d) *Record*. Compute the scalar $q_t(S_t) = \sum_a x_t(a)P(S_t | a)$ and store $(i, q_t(S_t), S_t, \cdot)$
 195 in $\mathcal{P}$, to be completed when X_t arrives at time $t + d$.

196 Steps (a), (c) and (d) are each $O(K)$ and there is no optimization subproblem, so a round costs $O(K)$
 197 once P is stored. The state is held in the m vectors L_i , which is mK numbers, the matrix P , which
 198 is $K|\mathcal{S}|$ numbers, and four scalars for each of the at most d rounds in flight. Only the scalar $q_r(S_r)$
 199 is needed from round r , not the vector x_r .

200 **A worked example.** Take three actions and two states, with

$$P = \begin{pmatrix} 0.8 & 0.2 \\ 0.4 & 0.6 \\ 0 & 1 \end{pmatrix}, \quad x_t = (\tfrac{1}{2}, \tfrac{1}{4}, \tfrac{1}{4}),$$

201 so $q_t(s_1) = \frac{1}{2}(0.8) + \frac{1}{4}(0.4) + \frac{1}{4}(0) = 0.5$ and $q_t(s_2) = 0.5$. The third action cannot produce
 202 s_1 at all. Suppose the learner draws a_1 , observes s_1 , and receives $X_t = 1$ at time $t + d$. Then
 203 $\hat{c}_t(a) = P(s_1 | a)/0.5$, so $\hat{c}_t = (1.6, 0.8, 0)$ against the action-level $\tilde{c}_t = (2, 0, 0)$. One observation
 204 updates all three actions: a_2 is charged although it was never played, because it had a 0.4 chance of
 205 producing the state seen, and a_3 is charged nothing, because it could not have produced that state.
 206 None of this introduces bias: with $\theta_t = (1, 0)$ the true losses are $c_t = (0.8, 0.4, 0)$, and averaging $\hat{c}_t$
 207 over $S_t \sim q_t$ returns exactly that. The gain is in the second moment: $v_t = 1.44$ here, against $K = 3$
 208 for the action-level estimate, and with $\theta_t = (1, 1)$ both are attained exactly.

209 *Remark 1* (The χ^2 form). Under x_t the pair (A_t, S_t) has law $x_t(a)P(s | a)$ with marginals x_t and
 210 q_t , so

$$\begin{aligned} \chi^2(P_{A_t, S_t} \| P_{A_t} \otimes P_{S_t}) &= \sum_{a, s} \frac{(x_t(a)P(s | a) - x_t(a)q_t(s))^2}{x_t(a)q_t(s)} \\ &= \sum_{a, s} \frac{x_t(a)P(s | a)^2}{q_t(s)} - 2 + 1 = v_t - 1, \end{aligned}$$

211 the cross term summing to $\sum_{a, s} x_t(a)P(s | a) = 1$ and the last to $\sum_{a, s} x_t(a)q_t(s) = 1$. So
 212 $v_t = 1$ exactly when A_t and S_t are independent, which pins the rows of P only on the support
 213 of x_t and therefore means all rows agree whenever x_t has full support, as exponential weights
 214 does. If P is deterministic and every state is reached by some action in the support of x_t , then v_t
 215 equals the number of states that support reaches. In particular, under full-support play and an onto
 216 deterministic P , $v_t = |\mathcal{S}|$. This is an identity, not a bound; it is used for reading v_t , and once in the
 217 proof of Theorem 1.

218 **Proposition 1.** Fix t , let x_t be the play distribution, $q_t(s) = \sum_a x_t(a)P(s | a)$ the induced state
 219 distribution, and $\hat{c}_t(a) = P(S_t | a)X_t/q_t(S_t)$, which is well defined because only states with
 220 $q_t(s) > 0$ occur. For statements involving $1/x_t(a)$, restrict to actions with $x_t(a) > 0$; STATE-EXP3
 221 itself has full support on every action. Then $\hat{c}_t(a) \geq 0$, $\hat{c}_t(a) \leq 1/x_t(a)$, $\mathbb{E}[\hat{c}_t(a) | \mathcal{F}_{t-1}] = c_t(a)$,
 222 and

$$\sum_a x_t(a)\mathbb{E}[\hat{c}_t(a)^2 | \mathcal{F}_{t-1}] \leq v_t := \sum_{s: q_t(s) > 0} \frac{\sum_a x_t(a)P(s | a)^2}{q_t(s)}, \quad 1 \leq v_t \leq |\mathcal{S}|,$$

223 states of zero probability being omitted throughout, which is the convention $0/0 = 0$. The action-
 224 level $\hat{c}_t(a) = \mathbf{1}\{A_t = a\}X_t/x_t(a)$ has the same first three properties and weighted second moment
 225 at most K .

226 *Proof.* Nonnegativity is immediate from $X_t \geq 0$. Given $\mathcal{F}_{t-1}$ the state S_t has distribution q_t , and
 227 $\mathbb{E}[X_t | \mathcal{F}_{t-1}, S_t = s] = \theta_t(s)$ by the model, since that conditional mean does not depend on
 228 the action, so $\mathbb{E}[\hat{c}_t(a) | \mathcal{F}_{t-1}] = \sum_s q_t(s)P(s | a)\theta_t(s)/q_t(s) = c_t(a)$. For the range, $q_t(s) \geq$
 229 $x_t(a)P(s | a)$ for every a , so $P(s | a)/q_t(s) \leq 1/x_t(a)$ and $X_t \leq 1$. For the second moment,
 230 $X_t^2 \leq 1$ gives $\mathbb{E}[\hat{c}_t(a)^2 | \mathcal{F}_{t-1}] \leq \sum_s P(s | a)^2/q_t(s)$, and exchanging the sums gives v_t . Each
 231 summand of v_t is at most 1, since $P(s | a) \leq 1$ makes $\sum_a x_t(a)P(s | a)^2 \leq q_t(s)$, and at least
 232 $q_t(s)$, since Jensen applied to $a \mapsto P(s | a)$ under x_t gives $\sum_a x_t(a)P(s | a)^2 \geq q_t(s)^2$. Summing
 233 the two bounds gives $1 \leq v_t \leq |\mathcal{S}|$. The action-level claims are the same computation with the state
 234 replaced by the played action, whose weighted second moment is $\sum_a x_t(a)/x_t(a) = K$. $\square$

235 The two ends are attained. If every action induces the same state distribution then $v_t = \sum_s q_t(s) =$
 236 1, and if P is deterministic, each action reaching a single state, then every $P(s \mid a)$ is 0 or 1 and
 237 each summand is exactly 1, so v_t counts the states reached and equals $|\mathcal{S}|$ when every state has an
 238 action reaching it. Disjoint row supports give $v_t = K$ rather than $|\mathcal{S}|$, since the summands over
 239 one action's private states add to $\sum_s P(s \mid a) = 1$, so a sensor with many private states per action
 240 resolves few of them. So v_t reads the overlap between the rows of P rather than their length, and
 241 $\bar{v} = \sup_x v(x)$, the supremum over play distributions, depends on P alone. An outcome observed
 242 at a state updates the estimate for every action that can reach it, which is why the construction of
 243 Section 3 gives each drifting coordinate a private state, forcing $v_t = |\mathcal{S}|$ there. The immediately
 244 observed pairs (A_t, S_t) identify P separately, though Theorem 3 assumes it known.

245 B Proofs: pooling and grouping

246 Two results are stated here rather than in the body, since Theorem 1 covers the algorithm actually
 247 run. The first is the interleaved variant.

248 **Theorem 3.** *Let P be known and (θ_t) oblivious. Pick before the run any fixed number V with*
 249 *$\sum_t v_t \leq V$ always. Then INTERLEAVED-STATE-EXP3 at the matching learning rate satisfies*
 250 *$\mathbb{E}[R_T] \leq \sqrt{2(d+1)V \log K}$. Writing $\bar{v}$ for the largest v_t over all ways of playing, $V = \bar{v}T \leq |\mathcal{S}|T$*
 251 *is one such choice, and it depends on P alone.*

252 The second merges states that are not worth telling apart. Group them with a map g from $\mathcal{S}$ onto
 253 a smaller set $\mathcal{Z}$ and run the same algorithm on $g(S_t)$, where $P^g(z \mid a) = \sum_{s: g(s)=z} P(s \mid a)$.
 254 Grouping never raises the dimension, but the estimate now tracks a group average, so the widest
 255 spread of losses inside a group, $\delta_t(g) = \max_z \max_{s, s': g(s)=g(s')=z} |\theta_t(s) - \theta_t(s')|$, is the error
 256 introduced.

257 **Theorem 4.** *Under the hypotheses of Theorem 3, running on g with $\sum_t v_t(g) \leq V(g)$ gives*
 258 *$\mathbb{E}[R_T] \leq \sqrt{2(d+1)V(g) \log K} + 2 \sum_t \delta_t(g)$.*

259 Merging need not preserve state-sufficiency, so the grouped estimate is right not for c_t but for a stand-
 260 in loss differing from it by at most $\delta_t(g)$ on every action. Keeping states apart preserves differences
 261 that matter but learns more slowly; merging shares more but pays $2 \sum_t \delta_t(g)$. The best grouping is
 262 the true one: with 16 states from four hidden groups, sizes 1, 2, 4, 8, 16 give regret 1590, 831, 754,
 263 920, 1080, lowest at four, which no algorithm was told. The map g is fixed in advance; learning it
 264 from data remains open.

265 **Proof idea.** The pooled estimator is unbiased for every action, while Proposition 1 bounds its play-
 266 weighted second moment by v_t rather than K . With $m = d + 1$ interleaved copies, the outcome
 267 generated on a round belonging to one copy is available before that copy acts again, so each copy
 268 is an ordinary exponential-weights process with immediate feedback. Summing the $d + 1$ potential
 269 inequalities gives $(d + 1) \log K / \eta$, and the pooled second moments give $\frac{\eta}{2} \sum_t v_t$.

270 *Proof of Theorem 3.* Fix a copy i and let T_i be the number of rounds it owns, so $\sum_i T_i = T$. Round
 271 t 's outcome is usable from round $t + d + 1$, which with $m = d + 1$ is exactly $t + m$, the copy's next
 272 round, so within a copy the feedback is undelayed and L_i carries exactly the estimates of that copy's
 273 earlier rounds. Exponential weights on nonnegative losses give, for any action a^* and any $\eta > 0$,

$$\sum_{t \in i} (\langle x_t, \hat{c}_t \rangle - \hat{c}_t(a^*)) \leq \frac{\log K}{\eta} + \frac{\eta}{2} \sum_{t \in i} \sum_a x_t(a) \hat{c}_t(a)^2,$$

274 by the usual potential telescoping with $e^{-z} \leq 1 - z + z^2/2$, which needs $z \geq 0$ and therefore
 275 $\hat{c}_t \geq 0$, supplied by Proposition 1. Each x_t is $\mathcal{F}_{t-1}$ -measurable, because L_i then holds only rounds
 276 $r \leq t - m = t - d - 1$, so taking expectations turns the left side into $\sum_{t \in i} (\langle x_t, c_t \rangle - c_t(a^*))$ by
 277 unbiasedness and the quadratic term into at most $\mathbb{E}[\sum_{t \in i} v_t]$, again by Proposition 1. Summing over
 278 the $m = d + 1$ copies and using $\sum_i \min_a \sum_{t \in i} c_t(a) \leq \min_a \sum_t c_t(a)$, which is what lets the copies
 279 be compared against one common action, gives $\mathbb{E}[R_T] \leq (d + 1) \log K / \eta + \frac{\eta}{2} \mathbb{E}[\sum_t v_t]$. If $\sum_t v_t \leq$
 280 V surely then $\eta = \sqrt{2(d + 1) \log K / V}$ balances the two terms and gives $\sqrt{2(d + 1)V \log K}$, and
 281 $V = \bar{v}T$ is admissible by the definition of $\bar{v}$. $\square$

282 The delay enters as a factor because the copies do not share their estimates, each paying $\log K/\eta$
 283 over $T/(d+1)$ rounds. Sharing them gives the $m = 1$ algorithm of Theorem 1, whose additive
 284 delay term avoids this penalty.

285 *Proof of Theorem 4.* Running the algorithm on g is running it on the sensor $(\mathcal{Z}, P^g)$. The surrogate
 286 losses below depend on x_t and so are predictable rather than oblivious, but the proof of Theorem 3
 287 uses obliviousness nowhere, only that each c_t^g is $\mathcal{F}_{t-1}$ -measurable given the play, so it applies un-
 288 changed and bounds regret measured in the surrogate losses $c_t^g(a) = \sum_z P^g(z | a) \bar{\theta}_t(z)$, where
 289 $\bar{\theta}_t(z) = \sum_{s: g(s)=z} q_t(s) \theta_t(s) / q_t^g(z)$ is the group mean the grouped estimator is unbiased for, by the
 290 same computation as in Proposition 1 with S_t replaced by $g(S_t)$. Fix a and z . Both $\{q_t(s)/q_t^g(z)\}$
 291 and $\{P(s | a)/P^g(z | a)\}$ are probability vectors on $\{s : g(s) = z\}$, so their θ_t -averages lie in a
 292 common interval of length $\delta_t(g)$ and differ by at most $\delta_t(g)$. Weighting by $P^g(z | a)$ and summing
 293 over z gives $|c_t^g(a) - c_t(a)| \leq \delta_t(g)$ for every a , hence $\langle x_t, c_t \rangle - c_t(a^*) \leq \langle x_t, c_t^g \rangle - c_t^g(a^*) + 2\delta_t(g)$.
 294 Summing over t adds $2 \sum_t \delta_t(g)$. Merging leaves $V(g)$ no larger than a bound valid for the raw
 295 states: it suffices to merge two states, the general case following by induction on the merges. Writ-
 296 ing P_1, P_2 for the two columns, $B_j = \sum_a x_t(a) P_j(a)$ and $\lambda = B_1/(B_1 + B_2)$, Cauchy-Schwarz
 297 gives $(P_1 + P_2)^2 \leq \lambda^{-1} P_1^2 + (1 - \lambda)^{-1} P_2^2$ pointwise in a , so after dividing by the merged denom-
 298 inator the merged summand is at most the two it replaces. $\square$

299 C Relation to other inaccuracy measures

300 **Neighbouring settings.** Vernade et al. (2020) vary the action-to-state map while fixing the state-to-
 301 loss map, where we do the reverse. Wang et al. (2026) and Zhang et al. (2026) measure a prediction
 302 error as E_2 does, without the delay. Feedback graphs (Alon et al., 2015; Esposito et al., 2022;
 303 Gabbianelli et al., 2023) reveal other actions' losses outright, where here the state's loss must still
 304 be estimated from the one outcome that arrives.

305 Write $\Lambda_2 = \sum_t \min\{1, \|c_t - m_t\|_2^2\}$ for the inaccuracy measure of Bar-On and Mansour (2025).
 306 Unlike E_2 , which charges only the worst action in a round, Λ_2 grows with how many actions are
 307 predicted wrongly, and since $\|v\|_\infty \leq \|v\|_2 \leq \sqrt{J} \|v\|_\infty$ for a vector v with J nonzero coordinates,

$$E_2 \leq \Lambda_2 \leq J E_2,$$

308 with both ends attained: the left when a single action is mispredicted, the right when all J are
 309 mispredicted equally. So a bound in E_2 is never weaker and can be J times stronger, which is why
 310 Theorem 2 is stated in E_2 . The two measures agree exactly when the prediction error is concentrated
 311 on one action.

312 D Proofs: the lower bound

313 The measure E_2 is controlled by the state-loss variation W , and no converse bound holds in general.

314 **Lemma 1** (Delay window). $E_2 \leq d^2 W$ for $W = \sum_t \|\theta_t - \theta_{t-1}\|_\infty^2$. The ratio $E_2/(d^2 W)$ tends
 315 to one when θ drifts monotonically in one coordinate at a constant rate and $T/d \rightarrow \infty$, the first d
 316 rounds contributing $\sum_{j < d} j^2$ rather than d^2 each under the convention $m_t = c_1$.

317 In words, accumulated squared prediction error is at most the state-loss variation inflated by the
 318 square of the delay, and no better bound in W alone is available.

319 *Proof of Lemma 1.* Throughout write $\Delta_j = \theta_j - \theta_{j-1}$, set $\theta_0 = \theta_1$, so $\Delta_1 = 0$, and read c_r as c_1 for
 320 $r \leq 0$, matching the convention $m_t = c_1$ for $t \leq d$ of Section 2. Then $c_t - c_{t-d} = P \sum_{j=t-d+1}^t \Delta_j$.
 321 Because the rows of P are probability vectors, averaging cannot increase the supremum norm, so
 322 $\|c_t - c_{t-d}\|_\infty \leq \sum_j \|\Delta_j\|_\infty$. By Cauchy-Schwarz applied to the d summands, $\|c_t - c_{t-d}\|_\infty^2 \leq$
 323 $d \sum_{j=t-d+1}^t \|\Delta_j\|_\infty^2$. Summing over t and using that each increment belongs to at most d windows
 324 gives $E_2 \leq d^2 W$. The ratio $E_2/(d^2 W)$ approaches one when every increment has the same mag-
 325 nitude, sign and active coordinate, and some action reaches that coordinate deterministically, since
 326 then no cancellation occurs within a full window and P passes the increment through undiminished.
 327 Without the second condition the drift can lie in the kernel of P and the ratio can be zero. $\square$

328 The construction in full. Take states $s_0, \dots, s_J$, a stationary deterministic map sending a_j to s_j for
 329 $j \leq J$ and every other action to s_0 , and independent Rademacher signs $\sigma_b^{(j)}$ for $2 \leq b \leq B = \lfloor T/h \rfloor$.
 330 Set $\theta_t(s_0) = \frac{1}{2}$ throughout, $\theta_t \equiv \frac{1}{2}$ on block 1, and $\theta_t(s_j) = \frac{1}{2} - \varepsilon \sigma_b^{(j)}$ on each later block b . The
 331 neutral first block makes $c_1 = \frac{1}{2} \mathbf{1}$, so the convention $m_t = c_1$ for $t \leq d$ of Section 2 supplies a
 332 constant over exactly the rounds that precede any usable feedback, which is what the isolation step
 333 below needs at $b = 1$. Rounds after the last complete block carry no drift, costing a constant factor.

334 **Lemma 2** (Rademacher lower tail). *Let S_B be a sum of $B \geq 32$ independent Rademacher signs.*
 335 *For every $1 \leq u \leq \sqrt{B}/32$, $\Pr(S_B \geq u\sqrt{B}) \geq \frac{u}{4} e^{-64u^2}$.*

336 *Proof.* Write $S_B = 2X - B$ with $X \sim \text{Bin}(B, \frac{1}{2})$, so the event is $\{X \geq B/2 + u\sqrt{B}/2\}$; passing
 337 to X removes the lattice gap, since X takes every integer value in $[0, B]$ whereas S_B takes only
 338 those of the parity of B . Put $n = \lfloor B/2 \rfloor$ and $t = \lceil u\sqrt{B}/2 \rceil + 1$, so $X \geq n+t$ implies the event and
 339 $t \leq u\sqrt{B}$. Consecutive binomial probabilities have ratio $(B-n-j+1)/(n+j)$, whose distance
 340 from one is at most $8j/B$; for $j \leq B/16$ that is at most $\frac{1}{2}$, where $\log(1-x) \geq -2x$ holds, so
 341 summing over $i \leq j$ gives $\Pr(X = n+j) \geq \Pr(X = n) e^{-16j^2/B}$ with $\Pr(X = n) \geq 1/(2\sqrt{B})$.
 342 Summing over the t integers $j \in \{t, \dots, 2t-1\}$, all at most $B/16$ because $u \leq \sqrt{B}/32$, and using
 343 $t \geq u\sqrt{B}/2$, gives $\Pr(X \geq n+t) \geq t e^{-64u^2}/(2\sqrt{B}) \geq \frac{u}{4} e^{-64u^2}$. $\square$

344 Three regimes give the maximal inequality the lower bound needs. Each $S_B^{(j)}$ is sub-Gaussian
 345 with variance proxy B , so $\mathbb{E} \max_j (S_B^{(j)})^+ \leq \sqrt{2B \log J} + \sqrt{B}$ throughout. For the matching
 346 lower bound, first suppose $\log J \leq B/8$ and $\log J \geq 128$. Taking $u = \sqrt{\log J/128}$, which
 347 then lies in $[1, \sqrt{B}/32]$, makes $e^{-64u^2} = J^{-1/2}$, so $\Pr(S_B^{(j)} \geq u\sqrt{B}) \geq \frac{u}{4} J^{-1/2}$ and inde-
 348 pendence across j gives $\Pr(\max_j S_B^{(j)} < u\sqrt{B}) \leq e^{-u\sqrt{J}/4}$, at most $\frac{1}{2}$ beyond an absolute J .
 349 Hence $\mathbb{E} \max_j (S_B^{(j)})^+ \geq \frac{1}{2} u\sqrt{B} = \Omega(\sqrt{B \log J})$. Second, if $\log J > B/8$ and $B \geq 1024$, take
 350 $u = \sqrt{B}/32$, so $u\sqrt{B} = B/32$ and $e^{-64u^2} = e^{-B/16}$, and $J \geq e^{B/8}$ makes $J e^{-B/16}$ diverge,
 351 giving $\mathbb{E} \max_j (S_B^{(j)})^+ = \Omega(B)$, which is $\Omega(\sqrt{B \min\{1 + \log J, B\}})$ in that regime. Third, when
 352 $\log J < 128$ or $B < 1024$ the quantity $\sqrt{B \min\{1 + \log J, B\}}$ is $O(\sqrt{B})$, and a single walk already
 353 gives $\mathbb{E}(S_B^{(1)})^+ = \frac{1}{2} \mathbb{E}|S_B| \geq \frac{1}{2} (\mathbb{E} S_B^2)^{3/2} (\mathbb{E} S_B^4)^{-1/2} \geq \frac{1}{2} \sqrt{B/3}$ by Hölder, since $\mathbb{E} S_B^2 = B$ and
 354 $\mathbb{E} S_B^4 \leq 3B^2$. Combining the three regimes, $\mathbb{E} \max_j (S_B^{(j)})^+ = \Theta(\sqrt{B \min\{1 + \log J, B\}})$. This is
 355 the route by which maxima of independent random walks fix the minimax rate for prediction with
 356 expert advice (Cesa-Bianchi and Lugosi, 2006).

357 **Proof idea.** Divide time into blocks of length d and give J actions private states whose losses are
 358 perturbed by independent Rademacher signs within each block. Every observation available when
 359 an action is chosen comes from an earlier block, and the supplied $m_t = c_{t-d}$ depends only on
 360 earlier signs, so the learner cannot predict the signs governing its current block. Its expected loss is
 361 therefore $T/2$, while the best fixed action in hindsight benefits from the maximum of J independent
 362 random walks. Choosing the amplitude to make $E_2 \leq \mathcal{E}$ gives the stated scale.

363 *Proof of Theorem 2.* Take $h = d$ and $\varepsilon = \frac{1}{2} \sqrt{\mathcal{E}/T}$. Every outcome usable when choosing A_t comes
 364 from a round $r \leq t-d-1$, and $t \leq bh$ with $h \leq d$ forces $r < (b-1)h$, so A_t is independent of block
 365 b 's own signs; the same computation places the stale vector m_t in blocks strictly before b , so handing
 366 it to the learner changes nothing. Every algorithm's expected loss therefore equals $T/2$, since its
 367 action is independent of the current block's signs and all states have mean $\frac{1}{2}$ marginally. Its regret
 368 is therefore exactly the comparator's fluctuation advantage, $\varepsilon d \mathbb{E} [\max_{j \leq J} (\sum_{2 \leq b \leq B} \sigma_b^{(j)})^+]$ with
 369 $B = \lfloor T/d \rfloor$, over the $B-1$ randomized blocks. Per round $\|c_t - m_t\|_\infty \leq 2\varepsilon$, because the supremum
 370 norm charges only the largest coordinate and each coordinate moves by at most 2ε at a block bound-
 371 ary; the gap is ε across the first randomized block, where m_t is still the neutral c_1 , and lies in $\{0, 2\varepsilon\}$
 372 thereafter. Hence $E_2 \leq 4\varepsilon^2 T = \mathcal{E}$ holds deterministically. The budget must bind on the realized
 373 sequence, since a budget holding only in expectation would not constrain the instance selected by
 374 Yao's principle. By Lemma 2 and the discussion following it, the expected maximum of the positive
 375 parts of J independent Rademacher sums of length $B-1$ is $\Theta(\sqrt{(B-1) \min\{1 + \log J, B-1\}})$,

which covers both regimes and the case $J = 1$. The hypothesis $d \leq T/2$ gives $B \geq 2$ and hence $B - 1 \geq B/2$, so replacing $B - 1$ by B costs a factor at most $\sqrt{2}$ in each branch of the minimum, including the saturated branch where the minimum is $B - 1$ rather than $1 + \log J$. Substituting $B = \lfloor T/d \rfloor$ and ε therefore gives $\Theta(\sqrt{d\varepsilon} \min\{1 + \log J, T/d\})$. Rounds after the last complete block carry no drift and change the constant only. When the oracle is withheld, the $\sqrt{KT}$ term is Esposito et al. (2023, Thm. 5.1), whose construction is stationary and therefore lies in the present model class. $\square$

Remark 2 (Where the bound is matched). Suppose $X_t = c_t(A_t)$. On instances whose drift is carried by a single action only one coordinate of $c_t - m_t$ is nonzero, so $\|c_t - m_t\|_2 = \|c_t - m_t\|_\infty$; on the family of Section 3 that coordinate has magnitude at most $2\varepsilon \leq 1$, so the truncation at one in Λ_2 is inactive and $\Lambda_2 = E_2$. The bound of Bar-On and Mansour (2025) then reads $\tilde{O}(\sqrt{KT} + dE_2)$, whose $\sqrt{dE_2}$ contribution Theorem 2 matches up to the logarithmic factor at $J = 1$, so on that family the drift term of the published upper bound is unimprovable in order. Its $\sqrt{KT}$ term is not matched, the oracle supplying an entire stale loss vector, so the oracle-assisted problem need not retain bandit exploration complexity. Neither result provides an E_2 -adaptive upper bound under noisy outcomes.

E Numerical detail

Sixteen studies were run; the three that bear most directly on the main claim and its boundary are reported here as Experiments 1 to 3, each with an observation, an interpretation and a limitation, and each named with the script that produces it. The code for all sixteen is available. Experiment 1 draws P from a Dirichlet; Experiment 2 uses a hand-structured funnel map; and Experiment 3 uses the deterministic hard-instance map of Theorem 2. All policies run on the same environment parameters and common random-number streams, so a seed fixes everything except which estimator is formed, and differences are taken within a seed before averaging. Round r 's outcome is consumed before round $r + d + 1$, matching Section 2. Reported $\pm$ is one standard error of the paired difference. *Sample-path pseudo-regret* means $\sum_t c_t(A_t) - \min_a \sum_t c_t(a)$ on one realization of the learner and environment randomness, computed from the true losses rather than the sampled outcomes; averaging it over seeds estimates the expected pseudo-regret the theorems bound.

The first two experiments probe the state channel of the claim, the first under tight control and the second in an operational shape, and the last probes its temporal boundary. Experiment 1 holds K and the task difficulty fixed and varies the overlap, which separates Theorem 1's prediction from the instances simply becoming easier; Experiment 2 shows the regime in an environment built to the shape of a funnel and measures it against a rate-optimal baseline. Experiment 3 then tests Theorem 2 on the drift channel, where pooling cannot help. Theorems 3 and 4, on the interleaved algorithm and on coarsening, have no experiment here; both are stated and proved in Appendix B and left untested in this submission. Other studies whose conclusions were superseded by a later control, or that replicated one already reported, are left to the released code.




Figure 3: **Left**, Experiment 1. The same overlap sweep with and without holding the task difficulty fixed. Matching the uniform-play regret reverses the trend, supporting the interpretation that the decline in the unmatched sweep is driven by task-difficulty confounding rather than by the pooling mechanism itself. Bars are one standard error of the paired difference.

413 **Experiment 1 (matched_control.py). Overlap drives the gain, not easier tasks.** Sweeping
414 the overlap without further control finds the advantage falling as $\bar{v}$ falls, which read at face value
415 contradicts Theorem 1, whose budget V^- falls with $\bar{v}$. That sweep is confounded, since raising the
416 overlap also flattens $c_t = P\theta_t$ across actions, collapsing the comparator gap so that both regrets
417 vanish together. This experiment holds $K = 40$, $|S| = 8$, $T = 10000$, $d = 50$ and the difficulty
418 fixed, and varies only $\hat{v}$. Difficulty is the uniform-play regret $\sum_t \text{mean}_a c_t(a) - \min_a \sum_t c_t(a)$, a
419 property of the instance with no algorithm in it, and the amplitude of θ is bisected in every cell to hit
420 a common target. The map keeps one contrast direction alive at every overlap and θ is aligned with
421 it, which is what lets the gap survive as the rows agree.

422 **Observed.** Across 20 paired seeds the relative gain rises from 11.7 per cent at $\hat{v} = 3.380$ to 19.3
423 per cent at $\hat{v} = 1.120$, with absolute gains 85.5 ± 8.9 to 106.2 ± 7.9 and a correlation of -0.929
424 between $\hat{v}$ and the gain. The unmatched sweep falls from 34.5 to 3.1 per cent over a comparable
425 range (Figure 3), so the apparent contradiction is an artefact of the confound rather than of the
426 mechanism.

427 **Interpretation.** The two sweeps disagree in sign, and the matched-difficulty sweep agrees with the
428 overlap dependence Theorem 1 predicts for the single-copy algorithm this experiment runs. Overlap
429 is what pooling exploits, and the unmatched sweep measures overlap confounded with an instance
430 becoming trivial.

431 **Limitation.** At the two highest overlaps some seeds cannot reach the common target and the realized
432 difficulty is 5 to 13 per cent below it, so the matching is close rather than exact. Since a smaller
433 gap should if anything shrink the advantage, the measured rise is conservative in that direction. The
434 matched gains are also smaller in magnitude than the unmatched ones, which is what fixing the
435 difficulty costs.




Figure 4: Experiment 2, a recommendation funnel. **Left**, the effective dimension stays near two while the catalogue grows thirty-two-fold, so the gap between K and $\hat{v}$ comes from the catalogue having categories, not from a chosen K . **Right**, the two arms at $T = 20000$ and 8 paired seeds, labelled by whether $\bar{v}(d+1) \log K \leq K + d$ holds. Bars are one standard error. No real data is used anywhere in this experiment.

436 **Experiment 2 (funnel.py, sota_check.py). A recommendation funnel, and a rate-optimal**
 437 **baseline on it.** Unlike Experiment 1, this experiment does not draw P from a Dirichlet; the Dirichlet
 438 construction is the assumption under test rather than evidence for it. This one builds an environment
 439 to the shape of a recommendation funnel instead. Actions are $K = 200$ catalogue items, states are
 440 $|\mathcal{S}| = 6$ ordered engagement depths, and the outcome is a delayed conversion. Three structural facts
 441 are imposed because funnels have them, namely that items fall into a modest number of categories
 442 sharing an engagement profile, that item quality is heavy-tailed, and that the loss falls with engage-
 443 ment depth and drifts seasonally rather than adversarially. The generated catalogue puts 74 per cent
 444 of its mass on no engagement and 1.4 per cent on the deepest. The same environments, seeds and
 445 delay convention then carry a comparison against the hybrid Tsallis-plus-entropy FTRL of Zimmert
 446 and Seldin (2020), which attains the optimal rate for delayed feedback, its delay term being additive
 447 rather than multiplicative. The algorithms published since match that adversarial rate rather than
 448 improve on it, adding adaptivity to stochastic instances instead (Schlisselberg et al., 2025); all of
 449 them read the action alone.

450 Observed. The effective dimension is $\hat{v} = 1.97$ against $K = 200$, a hundredfold gap, and it stays
 451 between 1.41 and 1.81 as the catalogue grows from 25 to 800 items (Figure 4). Against action-level
 452 weighting the state-pooled variant saves 79.3, 63.6 and 38.9 per cent at $d = 10, 50$ and 200, the
 453 largest margins in this appendix.

	family, cell	STATE-EXP3	Zimmert and Seldin	same, grid-best	STATE-EXP3, grid-best
454	funnel, $d = 10$	949.0	2959.4	2413.8	27.0
	funnel, $d = 50$	1728.1	3557.0	2556.4	95.9
	funnel, $d = 200$	3244.8	4734.7	2807.7	199.7

455 Experiment 2. Regret on the funnel under three tunings; grid-best columns are chosen by reading these
 456 benchmarks.

457 Against the rate-optimal baseline, with STATE-EXP3 at the single tuning used everywhere else here
 458 and Zimmert and Seldin (2020) at the best of a six-point grid around its own value, STATE-EXP3
 459 wins all three funnel cells, cutting regret by 31.5 to 67.9 per cent. With both at their grid-best it wins
 460 all three again and by more. When only Zimmert and Seldin (2020) is grid-tuned, STATE-EXP3
 461 wins at $d = 10$ and $d = 50$ and loses at $d = 200$.

462 Interpretation. The separation between K and $\hat{v}$ arrives from catalogue structure rather than from
 463 a tuned parameter, and it does not close as the catalogue grows, which is the claim the Dirichlet
 464 families cannot make. Where the state-pooled variant is not held to the worse tuning of the two,
 465 the state channel helps against a rate-optimal method, which is expected, since that method cannot
 466 read the state at all; grid tuning helps it far more than it helps Zimmert and Seldin (2020), consistent
 467 with pooling cutting estimator variance and lower variance admitting a larger learning rate. The
 468 condition under which interleaving’s bound would improve on the action-level rate fails at $d = 50$




Figure 5: Experiment 3, the drift channel. **Left**, measured regret against the scale Theorem 2 predicts, over twelve cells in which d varies 20-fold, E_2 100-fold and J 16-fold. The dashed line has slope one. **Right**, the same points divided by that scale, grouped by which parameter the cell sweeps, with the shaded band spanning the observed range. STATE-EXP3 appears here only as one more algorithm on the hard family, where it has no advantage.

and $d = 200$, while the single-copy variant still saves most of the regret there, which is the regime Theorem 1 was proved for.

Limitation. *No real data is used.* There is no interaction log in this environment, nothing here is fitted to one, and this is not a semi-synthetic benchmark in the usual sense of that phrase. The structure is imposed by hand from qualitative properties, so it shows that the regime is describable and self-consistent, not that any deployed system occupies it. The funnel cell at $d = 200$ lost when Zimmert and Seldin (2020) alone is grid-tuned is a real loss and reported as one. No column is a symmetric protocol: the two grids are finite and unequal, and STATE-EXP3’s grid-best sits at its top multiplier in every cell, so its own minimum may lie outside the range searched. Every grid-best multiplier is chosen by reading these benchmarks, so only the first column is a procedure a practitioner could run. Eight paired seeds on the funnel and ten on the matched sweep. The arms run $m = 1$, covered by Theorem 1, and the environment has no adversarial component, so it exercises the pooling channel and not the drift channel.

A heuristic baseline. The best-state rule identifies the state with the lowest mean loss and plays the action most likely to reach it, ignoring every other state. Across nine settings, five families at $d = 10$ and 50 plus one drift cell, it leads seven and STATE-EXP3 two. Its two losses are the spread family, built so that the best action mixes over several good states while a decoy action puts more mass than any other on the single best one: the rule takes the decoy and records 1496 ± 329 and 1495 ± 327 at $d = 10$ and 50 against our 486 ± 11 and 814 ± 20 . Sweeping T from 2500 to 40000 on that family and fitting regret to a power of T gives an exponent of 1.00 for the rule, 0.58 for STATE-EXP3 and 0.67 for action-level weighting, so the rule’s regret grows linearly where both learners’ grow sublinearly. Its spread across seeds averages 4.6 times ours over the nine settings. The rule needs one state to be both best and reachable, and nothing in it detects when that fails, which is why the gain it shows on seven settings is not a gain a practitioner can rely on.

Experiment 3 (drift_scaling.py). The drift bound describes the right scaling. The construction of Theorem 2 at $K = 40$, $T = 8000$, $h = d$ and 12 paired seeds, sweeping d from 10 to 200, the amplitude ε from 0.02 to 0.20 so that E_2 runs from 12 to 1193, and the drifting coordinates J from 1 to 16. We run four policies: uniform, action-level EXP3, STATE-EXP3, and the greedy policy on the exact stale vector m_t , the last because Theorem 2 claims to survive it. E_2 is measured on the realized instance rather than predicted, so the normalization uses the quantity the theorem uses.

Observed. The predicted scale $\sqrt{d E_2 \min\{1 + \log J, T/d\}}$ ranges from 38 to 377 across the grid and the measured regret tracks it (Figure 5). Normalized, the four policies read 0.230 to 0.429, 0.228 to 0.424, 0.238 to 0.417 and 0.237 to 0.437, a spread below 1.9 against a tenfold range in the predictor. The oracle-assisted policy has the largest mean ratio, 0.345 against 0.311, 0.314 and 0.310 for the three that never see m_t .

Interpretation. The three arguments enter the predicted scale differently and the collapse holds in all three sweeps, so it is the form of the bound rather than one fitted constant that tracks the data.

506 Greedy play on the exact m_t sits at the top of the band rather than below it, which is what the bound
 507 surviving the oracle amounts to in these runs.

508 Limitation. Every cell has $T/d > 1 + \log J$, so only the $1 + \log J$ branch of the minimum is
 509 exercised and the T/d saturation branch, which needs d comparable to T , is not probed. A lower
 510 bound quantifies over all algorithms and four policies cannot establish that. What these runs check
 511 is the scaling of the construction.

512 **Code and environment.** The code producing every reported number and figure is at [https://](https://github.com/melikabaghi/state-exp3)
 513 github.com/melikabaghi/state-exp3. It pins the software environment used for the reported
 514 numbers, python 3.14.3 with numpy 2.4.2 and matplotlib 3.10.8, states the seed count for each
 515 study, lists one command per numbered study, and commits the stored outputs each figure reads so
 516 that every figure rebuilds without rerunning an experiment.

517 Taken together these experiments cover both channels. Experiments 1 and 2 probe the state channel
 518 and are consistent with the effective dimension, rather than $|\mathcal{S}|$ or K , governing what pooling buys;
 519 Experiment 1 had to control for task difficulty before an overlap sweep pointed that way. Among
 520 the state-channel experiments, the funnel is the only environment whose matrix is not drawn from a
 521 Dirichlet and it reports the largest margins here, on a structure imposed by hand rather than measured.
 522 It also places STATE-EXP3 against a rate-optimal delayed-bandit method: ahead in every cell when
 523 both are tuned alike, behind at $d = 200$ when that method alone is grid-tuned. Experiment 3 turns
 524 to the other channel and finds the drift scale of Theorem 2 tracking the data, including for a learner
 525 handed the stale losses, which is the cost no amount of pooling reaches.

526 F Proof of Theorem 1: additive delay without interleaving

527 Interleaving cannot give an additive bound. Copy i , which absorbs a share V_i of the budget $V =$
 528 $\sum_i V_i$, has regret $\sqrt{2V_i \log K}$, and Cauchy–Schwarz over the m copies gives $\sqrt{2mV \log K}$, tight
 529 when the V_i agree, so the $\sqrt{m}$ in Theorem 3 is not an artifact of the analysis. An additive bound has
 530 to come from a single copy.

531 Throughout, $L_t = \sum_{r \leq t-d-1} \hat{c}_r$ and $x_t \propto e^{-\eta L_t}$, so $L_{t+1} - L_t = \hat{c}_{t-d}$, with $\hat{c}_r = 0$ for $r < 1$.
 532 The filtration $\mathcal{F}_t$ is the environment’s, as defined in Section 2; under it x_t is $\mathcal{F}_{t-d-1}$ -measurable, so
 533 x_{r+d} is $\mathcal{F}_{r-1}$ -measurable.

534 **The drift lemma.** The natural route prices the change of measure from x_r to x_{r+d} through the
 535 normaliser Z_r , whose reciprocal has no uniformly bounded moment. That is avoidable. The pooled
 536 estimate satisfies

$$\langle x_t, \hat{c}_t \rangle = \sum_a x_t(a) \frac{P(S_t | a) X_t}{q_t(S_t)} = X_t \leq 1$$

537 identically, on every path, because the denominator is the state marginal of the same distribution that
 538 supplies the numerator.

539 **Lemma 3.** *If $\eta \leq 1/(e(d+1))$ then $x_{t+1}(a) \leq (1 + 1/d) x_t(a)$ for every t and every a , surely, and*
 540 *hence $x_{t+d}(a) \leq e x_t(a)$.*

541 *Proof.* Strong induction on t . For $t \leq d$ nothing is delivered and $x_{t+1} = x_t$. For $t > d$,
 542 assume the claim for every $s < t$ and compose it over the d steps from $t-d$ to t , giving
 543 $x_t(a) \leq (1 + 1/d)^d x_{t-d}(a) \leq e x_{t-d}(a)$. Contracting with the fixed nonnegative vector $\hat{c}_{t-d}$ and
 544 using the identity above, $\langle x_t, \hat{c}_{t-d} \rangle \leq e \langle x_{t-d}, \hat{c}_{t-d} \rangle = e X_{t-d} \leq e$. With $Z_t = \sum_b x_t(b) e^{-\eta \hat{c}_{t-d}(b)}$
 545 and $e^{-z} \geq 1 - z$, this gives $Z_t \geq 1 - \eta \langle x_t, \hat{c}_{t-d} \rangle \geq 1 - \eta e \geq d/(d+1)$, so $x_{t+1}(a) =$
 546 $x_t(a) e^{-\eta \hat{c}_{t-d}(a)} / Z_t \leq x_t(a) / Z_t \leq (1 + 1/d) x_t(a)$. $\square$

547 What the induction bounds is $\langle x_t, \hat{c}_{t-d} \rangle = X_{t-d} q_t(S_{t-d}) / q_{t-d}(S_{t-d})$, the state-marginal ratio that
 548 the $1/Z_r$ route could not reach. Individual $\hat{c}_r(a)$ remain unbounded; only their x_t -weighted mass is
 549 controlled. Cesa-Bianchi et al. (2019) prove the action-level version under $\eta \leq 1/(Ke(d+1))$ and
 550 Thune et al. (2019) remove the K , using that their estimate is carried by the single action played, so
 551 the weight cancels into a same-action ratio. A pooled estimate is spread over every action and does
 552 not collapse that way; the identity above replaces the cancellation.

553 **Centring.** Let $\zeta_t = \hat{c}_t - X_t \mathbf{1}$, with $\zeta_r = 0$ for $r < 1$. Exponential weights is unchanged by a
 554 shift that is constant across actions, so x_t is the same iterate and this is a change of analysis, not of
 555 algorithm (Eldowa et al., 2024). Keeping $\gamma(s) = \mathbb{E}[X_t^2 \mid S_t = s]$ inside the state sum,

$$\mathbb{E}[\langle x_t, \zeta_t^2 \rangle \mid \mathcal{F}_{t-1}] = \sum_s \gamma(s) \left[\frac{\sum_a x_t(a) P(s \mid a)^2}{q_t(s)} - q_t(s) \right] \leq v_t - 1,$$

556 each bracket being nonnegative by Jensen and $\gamma(s) \leq 1$, the final step being Remark 1. Bounding
 557 $\langle x_t, \hat{c}_t^2 \rangle$ by v_t first and then subtracting does not work: it leaves $v_t - \mathbb{E}[X_t^2]$.

558 *Proof of Theorem 1.* Since $\zeta_r(a) \geq -X_r \geq -1$ we have $\eta \zeta_r(a) \geq -\eta$, and Lagrange's re-
 559 mainder gives $e^{-z} \leq 1 - z + (e^\eta/2)z^2$ for $z \geq -\eta$. With $\log(1 + y) \leq y$, the potential
 560 $W_t = \sum_a \exp(-\eta \sum_{r \leq t-d-1} \zeta_r(a))$ obeys $\log(W_{t+1}/W_t) \leq -\eta \langle x_t, \zeta_{t-d} \rangle + (\eta^2 e^\eta/2) \langle x_t, \zeta_{t-d}^2 \rangle$.
 561 Telescoping from $W_1 = K$ and using $\log W_{T+1} \geq -\eta \sum_{r \leq T-d} \zeta_r(a^*)$ for the comparator a^* ,

$$\sum_{r=1}^{T-d} \langle x_{r+d}, \zeta_r \rangle \leq \sum_{r=1}^{T-d} \zeta_r(a^*) + \frac{\log K}{\eta} + \frac{\eta e^\eta}{2} \sum_{r=1}^{T-d} \langle x_{r+d}, \zeta_r^2 \rangle.$$

562 The shift cancels between the two sides. As x_r is $\mathcal{F}_{r-d-1}$ -measurable, $\mathbb{E}[\langle x_r, \hat{c}_r \rangle] = \mathbb{E}[\langle x_r, c_r \rangle]$
 563 and $\mathbb{E}[\hat{c}_r(a^*)] = c_r(a^*)$, and the d rounds undelivered by $T + 1$ contribute at most d , so

$$\mathbb{E}[R_T] \leq \frac{\log K}{\eta} + \frac{\eta e^\eta}{2} \sum_r \mathbb{E}[\langle x_{r+d}, \zeta_r^2 \rangle] + \sum_r \mathbb{E}[\langle x_r - x_{r+d}, \hat{c}_r \rangle] + d.$$

564 Lemma 3 gives $x_{r+d} \leq e x_r$ surely, and x_{r+d} is $\mathcal{F}_{r-1}$ -measurable, so the second sum is at most
 565 $e \sum_r \mathbb{E}[v_r - 1] \leq eV^-$.

566 For the third, write $x_{r+d}(a) = x_r(a) e^{-\eta G_r(a)} / Z'$ with $G_r = \sum_{u=r-d}^{r-1} \hat{c}_u \geq 0$ and normaliser
 567 $Z' \leq 1$, distinct from the one-step Z_t of the Lemma. Then $x_r(a) - x_{r+d}(a) \leq x_r(a)(1 -$
 568 $e^{-\eta G_r(a)}) \leq \eta x_r(a) G_r(a)$ for every a , trivially where the left side is negative. Since $\hat{c}_r \geq 0$ and
 569 x_r, G_r are $\mathcal{F}_{r-1}$ -measurable, the conditional expectation is at most $\eta \langle x_r, G_r \rangle$, and $\mathbb{E}[\langle x_r, G_r \rangle] =$
 570 $\sum_{u=r-d}^{r-1} \mathbb{E}[\langle x_r, \hat{c}_u \rangle] \leq d$: for each such u , x_r is $\mathcal{F}_{r-d-1}$ -measurable and $r - d - 1 \leq u - 1$, so
 571 it leaves the conditional expectation and $\mathbb{E}[\langle x_r, \hat{c}_u \rangle] = \langle x_r, c_u \rangle \leq 1$. This is measurability, not
 572 independence, and at $u = r - d$ the two σ -fields coincide, so the step rests on the timing convention
 573 of Section 2. The sum is at most ηdT .

574 Finally $\eta \leq 1/(2e)$ gives $e^{1+\eta} < 3.3$, so $\mathbb{E}[R_T] \leq \log K/\eta + \frac{\eta}{2}(3.3V^- + 2dT) + d$. For $f(\eta) =$
 575 $A/\eta + B\eta$ restricted to $\eta \leq \eta_{\max}$ one has $\min f \leq 2\sqrt{AB} + A/\eta_{\max}$, in the interior and in the
 576 binding case alike; with $A = \log K$, $B = (3.3V^- + 2dT)/2$ and $\eta_{\max} = 1/(e(d+1))$ this is the
 577 stated bound. $\square$

578 The cap on η binds only when $d \gtrsim T/(e^2 \log K)$, and not at the settings of Appendix E: the fun-
 579 nel's own rates are 0.0058, 0.0031 and 0.0016 at $d = 10, 50, 200$ against caps 0.0334, 0.0072 and
 580 0.0018, so Theorem 1 covers those runs as they were executed. Over three seeds and 20000 rounds
 581 the induction has room to spare, $\max_a x_{t+1}(a)/x_t(a)$ reaching 1.006, 1.003 and 1.002 against the
 582 permitted 1.100, 1.020 and 1.005.

583 The two bounds are numerically close on the funnel, where $\bar{v} \approx 2$: the additive form carries $3.3(\bar{v} -$
 584 $1) + 2d$ and the interleaved one $(d+1)\bar{v}$, and these cross near $\bar{v} = 2$. The additive form gains as $\bar{v}$
 585 rises toward $|S|$, and it describes the algorithm that is actually run, which is the reason to prefer it.