
Pooling and Drift in Delayed Bandits

Melika Baghi

Georgia Institute of Technology
mbaghi3@gatech.edu

Abstract

A system often has to act long before it learns whether the act worked: a recommender sees a click in seconds and a purchase in days. With K choices and a delay of d rounds, the best rate available is $\tilde{O}(\sqrt{(K+d)T})$ over T rounds (Esposito et al., 2023), so a longer menu is always more expensive to learn from. It need not be. If the outcome depends on the choice only through the state that choice produced, a click or a browse or a cart, then one late outcome informs every choice that could have produced the state observed, and the price is set by how many genuinely different states the choices produce rather than by how many choices there are. We measure that by an *effective dimension* v_t between 1 and the number of states, and prove $\tilde{O}(\sqrt{(d+1)V \log K})$ for a rotating algorithm and $\tilde{O}(\sqrt{V^-} + \sqrt{dT})$ for the single-copy one that is actually used, for any budget fixed in advance; merging similar states lowers the price again, at a bias we make explicit. The idea has a limit. Even handed the exact losses from d rounds ago, no algorithm escapes $\Omega(\sqrt{d\mathcal{E} \min\{1 + \log J, T/d\}})$, where J counts the drifting directions and $\mathcal{E}$ bounds how far the losses move while it waits. On generated data the state channel cuts regret by up to 79 per cent against choice-level weighting and by 32 to 68 per cent against the tuned minimax-optimal method of Zimmert and Seldin (2020).

1 Introduction

Delayed outcomes are common in sequential decision-making: a system must act now even though the outcome it cares about may arrive much later. In recommendation, an item is shown immediately, an engagement signal such as a click or browsing depth is observed soon after, and conversion or retention may only be known days or weeks later, so decisions are committed before their economically relevant consequences appear. The central question is when such a signal can reduce the cost of learning before the final outcome arrives.

The pattern is not special to recommendation: a markdown is set and sell-through is known a season later, staff are rostered and the service outcome lands after the shift, marketing spend commits before attribution resolves. A fast proxy arrives at once; the outcome the decision is judged on does not. Here a choice is an item that could be shown, a state is what the user does next, and the delayed outcome is whether they eventually buy. Different items may induce similar distributions over engagement states, so an outcome observed after one item can carry information about many others. The difficulty scales not with the number of choices K but with how many genuinely different operational regimes those choices induce, so thousands of items funnelling users through a few engagement patterns stay easy to learn from.

Bandits with intermediate observations formalize this: the learner repeatedly picks one action and sees the consequence of that action alone. Esposito et al. (2023) show that the *state-to-loss* map decides the complexity, an unrestricted one giving only the rate available with no intermediate observations; we fix that map and ask how much the *action-to-state* structure helps.

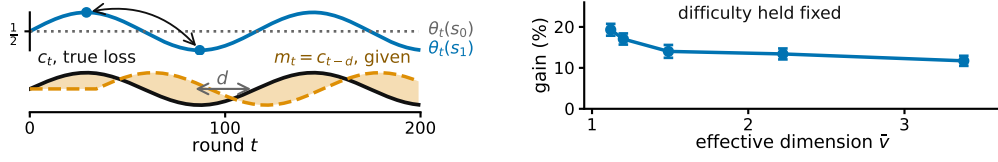


Figure 1: **An illustration of the two things that limit intermediate feedback, on generated data.** **Left**, what each state predicts about the outcome changes while the map from choices to states stays fixed, so the losses from d rounds earlier grow increasingly inaccurate; E_2 of Section 3 tracks that. **Right**, more overlap between choices allows more sharing, at matched difficulty, $K = 40$, $|\mathcal{S}| = 8$, $T = 10000$, $d = 50$, 20 paired seeds, bars one standard error of the paired difference (Appendix I, study 10).

Intermediate observations face two independent limitations: whether information can be shared across actions, and whether it stays useful across time.

Across actions: because similar actions reach similar states, we can *pool through the state*, charging a delayed outcome to every action that could have produced the state observed. The resulting statistical cost is governed by an *effective dimension* v_t , which measures how distinguishable the action-induced state distributions are under the learner’s current play. When many actions behave similarly through the state channel, v_t can be far smaller than K .

Across time: delay creates a second difficulty even when the learner is handed an exact description of the past. Let $m_t = c_{t-d}$ be the action-loss vector from d rounds earlier. Its usefulness depends on how far the losses have moved, which we measure by $E_2 = \sum_t \|c_t - m_t\|_\infty^2$ (Figure 1). Our lower bound asks how much regret this mismatch forces even when m_t is supplied for free.

In short, with P known the number of choices stops setting the price (Theorem 1), but no algorithm escapes the cost of waiting, even one handed the stale losses outright (Theorem 2). An unknown map can be estimated, with a weaker guarantee (Appendix E), and adapting to overlap and drift at once remains open (Section 6).

2 Related work

Our results connect two largely separate lines of work: delayed feedback, and information sharing across actions. *Delayed bandits*. Vernade et al. (2020) vary the action-to-state map while fixing the state-to-loss map; we do the reverse, letting the latter vary obliviously. Zimmert and Seldin (2020) give the optimal $\tilde{O}(\sqrt{KT} + \sqrt{D \log K})$ for total delay D , and Esposito et al. (2023) allow an unknown, possibly adversarial action-to-state map; fixing that channel is what makes its overlap geometry measurable through v_t . *Stale predictions*. Bar-On and Mansour (2025) derive upper bounds adapting to the inaccuracy of evolving observations; our lower bound instead asks what delay-dependent cost is unavoidable in stale-prediction error. Wang et al. (2026) bound squared gradient prediction error without delay, so the delay factor and its saturation are new here; Zhang et al. (2026) share the motivation in a stochastic Bayesian model where ours is adversarial. *Mediator feedback* is the closest structural connection: $v_t - 1$ is the χ^2 mutual information of Eldowa et al. (2024). Our contribution is not that quantity but its role when the outcome at the observed state is delayed, and the separation of that benefit from an independent temporal obstruction (Appendix F). Feedback graphs and off-policy coverage (Gabbianelli et al., 2023) are related forms of information sharing; Appendix A compares them.

3 Problem formulation

There are T rounds, K actions, a finite state space $\mathcal{S}$, a known action-to-state matrix P , and a fixed known delay d ; all logarithms are natural, $\tilde{O}$ hides logarithmic factors, and Appendix B tabulates every symbol. Every round follows the chain $A_t \xrightarrow{P} S_t \xrightarrow{\theta_t} X_t$: the learner draws A_t from a distribution x_t , observes $S_t \sim P(\cdot | A_t)$ at once, and receives the delayed outcome only at time $t+d$, encoded as a loss $X_t \in [0, 1]$ so that smaller is better, with $\mathbb{E}[X_t | \mathcal{F}_{t-1}, A_t, S_t] = \theta_t(S_t)$, where

$\mathcal{F}_{t-1}$ is everything seen before round t , and outcomes conditionally independent across rounds given the states. The conditional mean depends on A_t only through S_t . This *state-sufficiency* assumption licenses cross-action pooling: an outcome observed at state S_t may update every action that could have produced that state. If an action affects the outcome directly, in a way not represented in S_t , the argument no longer applies. Here x_t depends only on $\mathcal{F}_{t-1}$; X_r may first affect the action distribution at round $r + d + 1$, equivalently the round- t decision uses outcomes only from rounds $r \leq t - d - 1$.

That matrix $P \in [0, 1]^{K \times |\mathcal{S}|}$ is row-stochastic and fixed, and the *state-loss means* $\theta_t \in [0, 1]^{|\mathcal{S}|}$ are *oblivious*: chosen in advance, not reacting to the learner. Action losses are $c_t = P\theta_t$ and performance is the static *pseudo-regret*, the gap to the best single action in hindsight, $\mathbb{E}[R_T] = \mathbb{E}[\sum_t c_t(A_t)] - \min_a \sum_t c_t(a)$. **Oracle for the lower bound.** For Theorem 2 only, the learner also receives the *stale action-loss vector* $m_t \triangleq c_{t-d}$ before choosing A_t , with $m_t = c_1$ for $t \leq d$. This is not a function of the observables, so a lower bound surviving it is stronger, and STATE-EXP3 never uses it. We measure how misleading m_t is by $E_2 = \sum_t \|c_t - m_t\|_\infty^2 \leq T$, which stays small when errors are small even if frequent. The geometry of P thus governs sharing across actions and the evolution of θ_t governs staleness across time, the two axes of Sections 4 and 4.

4 What the theory guarantees

Two results follow. The first says the number of choices stops setting the price, for the algorithm one would actually run; the second says where that stops, since no sharing removes the cost of waiting. Two further results, a rotating variant and one for merged states, are stated and proved in Appendix D.

When an outcome arrives, do not update only the action played: update every action by how likely it was to produce the observed state. Here P is known and the learner does not receive m_t . For each state $s \in \mathcal{S}$ let $q_t(s) = \sum_{a=1}^K x_t(a)P(s | a)$ be the chance of seeing it this round, and give each choice $a \in \{1, \dots, K\}$ the estimate $\hat{c}_t(a) = P(S_t | a)X_t/q_t(S_t)$. It is correct on average, and how noisy it is depends on how much the actions overlap, through the *effective dimension* $v_t = \sum_s q_t(s)^{-1} \sum_a x_t(a)P(s | a)^2 \in [1, |\mathcal{S}|]$. STATE-EXP3 runs m copies of EXP3 (Auer et al., 2002) in rotation at learning rate η , copy $i \in \{0, \dots, m-1\}$ taking every m th round, and adds $\hat{c}_t$ to a copy's running total once the outcome is usable. A round costs $O(K)$ (Appendix C, with pseudocode and the χ^2 reading of v_t).

Theorem 1. *Let P be known and (θ_t) oblivious, and pick any $V^- \geq \sum_t (v_t - 1)$ before the run. Then STATE-EXP3 with $m = 1$ and any $\eta \leq 1/(e(d+1))$ satisfies $\mathbb{E}[R_T] \leq \sqrt{2(3.3V^- + 2dT) \log K} + e(d+1) \log K + d$.*

So K has been replaced by V^- , which counts how differently the choices behave rather than how many there are, and the delay enters additively: the bound is $\tilde{O}(\sqrt{V^-} + \sqrt{dT})$, and it covers the algorithm every experiment below runs, at rates they already satisfy. Writing $\bar{v}$ for the largest v_t over all ways of playing, $V^- = (\bar{v} - 1)T$ is one admissible choice and depends on P alone. Rotating $d + 1$ copies instead gives $\tilde{O}(\sqrt{(d+1)V \log K})$, and merging similar states lowers V again at a bias we make explicit (Appendix D). None of the three adapts to E_2 .

Where this stops. How much does waiting cost even when the stale losses are free? Theorem 2 shows a cost from time alone, not a joint overlap-drift bound: the construction gives each of the J drifting directions a private state, removing useful pooling and forcing $J \leq |\mathcal{S}| - 1$ (Appendix G).

Theorem 2. *Fix T , a delay $1 \leq d \leq T/2$, a drift budget $\mathcal{E} \leq T/4$, and $1 \leq J \leq \min\{K-1, |\mathcal{S}|-1\}$. Every algorithm, even one handed the exact losses from d rounds ago, meets some instance fixed in advance with $E_2 \leq \mathcal{E}$ on which $\mathbb{E}[R_T] = \Omega(\sqrt{d\mathcal{E} \min\{1 + \log J, T/d\}})$. Without that help, and if $K \leq T$, the best regret any algorithm can guarantee is $\Omega(\sqrt{KT} + \sqrt{d\mathcal{E} \min\{1 + \log J, T/d\}})$, up to constants.*

The first term is the usual cost of exploring. The second is the cost of delay: it grows with the wait, how far the losses move during it, and the number J of directions the best fixed action can use. They come from different hard instances, so adding them overstates the truth by at most a factor of two.

5 Experiments

The claim is about instances, so it must be checked where the number of choices and the number of distinct behaviours come apart. In a generated recommendation setting with $K = 200$ items and six levels of engagement the estimated dimension is 1.97, and the state-pooled variant cuts regret against action-level weighting by 79.3, 63.6 and 38.9 per cent at $d = 10, 50, 200$. It is hand-structured, not fitted to data, so it shows the mechanism, not a deployed system. An online plug-in $\hat{P}_t$ stays within 0.5 per cent of known P (study 2).

The advantage widens with the menu: at $|\mathcal{S}| = 4$ and 100 paired seeds the single copy beats action-level weighting on 73, 99 and 100 seeds at $K = 8, 40, 80$, the paired gain rising from 19.3 ± 3.4 to 169.3 ± 9.9 (study 1).

On the same family the state channel also beats the strongest published method for delayed bandits, the hybrid FTRL of Zimmert and Seldin (2020), tuned over a grid.

d	action-level	tuned (Zimmert and Seldin, 2020)	STATE-EXP3	cut	cut
10	4577	2959	949	79.3%	67.9%
50	4752	3557	1728	63.6%	51.4%
200	5307	4735	3245	38.9%	31.5%

Funnel family, $K = 200$, $|\mathcal{S}| = 6$, $T = 20000$, 8 paired seeds; the two cuts are against action-level weighting and against the tuned baseline (study 16).

These runs use the single-copy algorithm, the one Theorem 1 covers; Appendix I also reports the rotating variant. A best-state rule leads seven of nine settings, but where its assumption fails its regret grows linearly and its spread across seeds is four times ours. Against a minimax-optimal action-level delayed-bandit baseline the state channel still helps, cutting regret by 32 to 68 per cent on the funnel under the reported tuning protocol (Appendix I).

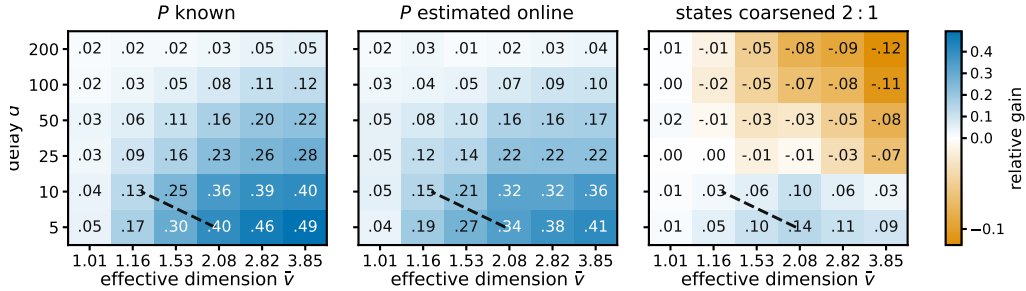


Figure 2: **Where pooling pays.** Relative gain $(R_{\text{action}} - R_{\text{state}})/R_{\text{action}}$ of pooling through the state, single copy, $K = 40$, $|\mathcal{S}| = 8$, $T = 5000$, 12 paired seeds; blue is a gain. $\hat{v}$ estimates $\sup_x v(x)$ for the raw map by Monte Carlo, which for the coarsened panel is not the $\bar{v}(P^g)$ that Theorem 6 charges. The dashed line compares two upper bounds for the rotating variant, so it is a reference and not a prediction for the variant plotted.

Figure 2 sweeps row overlap against delay and marks where pooling pays, in the three settings the studies treat separately: P known, P estimated, and states coarsened.

6 Discussion

Pool when many choices lead to similar states, coarsen when grouped states agree, and distrust old information when what the states predict moves fast. Missing is one algorithm adapting to both at once, at $\tilde{O}(\sqrt{KT} + \sqrt{d E_2 \log K})$; that remains unproved.

References

- N. Alon, N. Cesa-Bianchi, O. Dekel, and T. Koren. Online learning with feedback graphs: beyond bandits. *COLT*, PMLR 40, 2015.
- P. Auer, N. Cesa-Bianchi, Y. Freund, and R. E. Schapire. The nonstochastic multiarmed bandit problem. *SIAM Journal on Computing*, 32(1):48–77, 2002.
- Y. Bar-On and Y. Mansour. Non-stochastic bandits with evolving observations. *ALT*, PMLR 272, 2025.
- N. Cesa-Bianchi and G. Lugosi. *Prediction, Learning, and Games*. Cambridge University Press, 2006.
- N. Cesa-Bianchi, C. Gentile, and Y. Mansour. Delay and cooperation in nonstochastic bandits. *JMLR*, 20, 2019.
- K. Eldowa, N. Cesa-Bianchi, A. M. Metelli, and M. Restelli. Information capacity regret bounds for bandits with mediator feedback. *JMLR*, 25(353), 2024.
- E. Esposito, S. Masoudian, H. Qiu, D. van der Hoeven, N. Cesa-Bianchi, and Y. Seldin. Delayed bandits: When do intermediate observations help? *ICML*, PMLR 202, 2023.
- E. Esposito, F. Fusco, D. van der Hoeven, and N. Cesa-Bianchi. Learning on the edge: online learning with stochastic feedback graphs. *NeurIPS*, 35, 2022.
- G. E. Flaspohler, F. Orabona, J. Cohen, S. Mouatadid, M. Oprescu, P. Orenstein, and L. Mackey. Online learning with optimism and delay. *ICML*, PMLR 139, 2021.
- G. Guinet, S. Amin, and P. Jaillet. Effective dimension in bandit problems under censorship. *NeurIPS*, 35, 2022.
- G. Gabbianelli, G. Neu, and M. Papini. Online learning with off-policy feedback. *ALT*, PMLR 201, 2023.
- Y.-G. Hsieh, F. Iutzeler, J. Malick, and P. Mertikopoulos. Multi-agent online optimization with delays. *JMLR*, 23, 2022.
- P. Joulani, A. György, and C. Szepesvári. Online learning under delayed feedback. *ICML*, 2013.
- M. Mussi, S. Drago, M. Restelli, and A. M. Metelli. Factored-reward bandits with intermediate observations. *ICML*, PMLR 235, 2024.
- T. S. Thune, N. Cesa-Bianchi, and Y. Seldin. Nonstochastic multiarmed bandits with unrestricted delays. *NeurIPS*, 32, 2019.
- C. Vernade, A. György, and T. Mann. Non-stationary delayed bandits with intermediate observations. *ICML*, PMLR 119, 2020.
- S. Wang, A. Barik, and V. Y. F. Tan. Bandit convex optimization with gradient prediction adaptivity. *arXiv:2605.22191*, 2026.
- C.-Y. Wei, H. Luo, and A. Agarwal. Taking a hint: How to leverage loss predictors in contextual bandits? *COLT*, PMLR 125, 2020.
- K. W. Zhang, T. Baldwin-McDonald, K. Ciosek, L. Maystre, and D. Russo. Impatient bandits: optimizing for the long-term without delay. *JMLR*, to appear, 2026.
- M. Zhang. Toward optimal second-order path-length guarantee for adversarial multi-armed bandits. *arXiv:2608.15996*, 2026.
- J. Zimmert and Y. Seldin. An optimal algorithm for adversarial bandits with arbitrary delays. *AISTATS*, PMLR 108, 2020.

194 A Extended related work

195 Guinet et al. (2022) use the name *effective dimension* for an unrelated quantity, the number of direc-
196 tions a censored-feedback problem can resolve; it is not v_t and the observation model differs.

197 The minimax characterization of delayed adversarial bandits traces to Cesa-Bianchi et al. (2019),
198 and Joulani et al. (2013) give the additive form of the delay penalty in the stochastic case. On acting
199 before the loss is known, Flaspohler et al. (2021) study optimism under delay directly, Wei et al.
200 (2020) delimit when prediction error controls regret, Hsieh et al. (2022) show that an optimistic step
201 must run at a larger scale than the update, and Zhang (2026) gives second-order path-length bounds.

202 Structure that lets one observation inform several actions is the other half of the picture. Feedback
203 graphs (Alon et al., 2015) fix which actions are revealed together, and their stochastic version (Es-
204 posito et al., 2022) randomizes that set. Factored rewards with intermediate observations (Mussi
205 et al., 2024) decompose the reward instead. Off-policy coverage (Gabbianelli et al., 2023) is the
206 closest analogue of our effective dimension, since its mismatch coefficient measures the same kind
207 of overlap between the distribution that generated the data and the one being evaluated. The dis-
208 tinction from all of these is that a feedback graph randomizes *who* is observed, whereas our sensor
209 randomizes *what* is observed and the sharing is a consequence of P rather than a modeling primitive.

210 B Notation

Table 1: Symbols used in the paper. Every symbol is also defined at first use in the body, so this table is a reference rather than a dependency.

Symbol	Meaning
T, K, d	rounds, actions, fixed feedback delay
$\mathcal{S}, S_t$	finite state space, state observed at round t
$P \in [0, 1]^{K \times \mathcal{S} }$	row-stochastic action-to-state matrix, fixed known in Theorems 5 and 6, estimated in Appendix E
$\theta_t \in [0, 1]^{ \mathcal{S} }$	state-loss means at round t , oblivious
$c_t = P\theta_t \in [0, 1]^K$	induced loss over actions
$q_t(s) = \sum_a x_t(a)P(s a)$	state distribution induced by the play
$X_t \in [0, 1]$	outcome at time $t + d$, $\mathbb{E}[X_t \mathcal{F}_{t-1}, A_t, S_t] = \theta_t(S_t)$
x_t, A_t	action distribution at round t , action drawn from it
$m_t = c_{t-d}$	stale action-loss vector, given before the action, $m_t = c_1$ for $t \leq d$
E_2	$\sum_t \ c_t - m_t\ _\infty^2$, accumulated squared prediction error
Λ_2	$\sum_t \min\{1, \ c_t - m_t\ _2^2\}$, the measure of Bar-On and Mansour (2025)
W	$\sum_t \ \theta_t - \theta_{t-1}\ _\infty^2$, state-loss variation
$\hat{c}_t, \tilde{c}_t, \check{c}_t$	state-pooled, action-level, centered loss estimates
$\hat{g}_t$	the Safe-Pool blend $(1 - \lambda_t)\tilde{c}_t + \lambda_t\check{c}_t$, a distinct object
$\bar{c}_t, c_t^g, b_t$	conditional mean of $\hat{c}_t$, grouped surrogate loss, $\bar{c}_t - c_t$
$v_t, \bar{v}, V$	effective dimension $\sum_s q_t(s)^{-1} \sum_a x_t(a)P(s a)^2$, its supremum, its budget
$\hat{\bar{v}}$	Monte Carlo estimate of $\bar{v}$, used for measured values in Appendix I
$g, \mathcal{Z}, \delta_t(g)$	state grouping, its range, widest within-group loss spread
$\mathcal{E}, J$	budget constraining E_2 , number of independently drifting coordinates
$h, B, \varepsilon, \sigma_b^{(j)}$	block length, block count, amplitude, Rademacher signs
η, β, C	learning rate, Tsallis parameter, residual scale
α	smoothing pseudocount in $\hat{P}_t$
κ	sensor balance, $\bar{p}(s) \geq 1/(\kappa \mathcal{S})$ for $\bar{p} = K^{-1} \sum_a P(\cdot a)$
m, L_i, G_r	interleaving factor, estimate of copy i , estimates outstanding at r
$D, d_{\max}, \bar{\Delta}, \Gamma$	total delay, maximal delay, inaccuracy bound, the measure of Wang et al. (2026)

211 C The state-pooled estimator

212 **The algorithm.** Inputs are the action-to-state matrix P , the delay d , the number of copies m , and
213 the rate $\eta > 0$. The analyzed setting is $m = d + 1$ and the practical variant is $m = 1$. Round r 's
214 outcome becomes usable at round $r + d + 1$, which for $m = d + 1$ is the next round belonging to
215 the copy that played round r .

- 216 1. Set $L_i \leftarrow (0, \dots, 0) \in \mathbb{R}^K$ for $i = 0, \dots, m-1$, and let $\mathcal{P}$ be an empty table of rounds
 217 awaiting their outcome.
 218 2. For $t = 1, \dots, T$:
 219 (a) *Deliver*. If $r \triangleq t - d - 1 \geq 1$, read $(i_r, q_r(S_r), S_r, X_r)$ from $\mathcal{P}$, remove it, and update
 220 only that copy, $L_{i_r}(a) \leftarrow L_{i_r}(a) + P(S_r | a)X_r/q_r(S_r)$ for every a . This is the
 221 pooled step, and it touches every action rather than the one played.
 222 (b) *Select the copy*. $i \leftarrow t \bmod m$.
 223 (c) *Play*. $x_t(a) \propto \exp(-\eta L_i(a))$, draw $A_t \sim x_t$, and observe S_t at once.
 224 (d) *Record*. Compute the scalar $q_t(S_t) = \sum_a x_t(a)P(S_t | a)$ and store $(i, q_t(S_t), S_t, \cdot)$
 225 in $\mathcal{P}$, to be completed when X_t arrives at time $t + d$.

226 Steps (a), (c) and (d) are each $O(K)$ and there is no optimization subproblem, so a round costs $O(K)$
 227 once P is stored. The state is held in the m vectors L_i , which is mK numbers, the matrix P , which
 228 is $K|\mathcal{S}|$ numbers, and four scalars for each of the at most d rounds in flight. Only the scalar $q_r(S_r)$
 229 is needed from round r , not the vector x_r .

230 **A worked example.** Take three actions and two states, with

$$P = \begin{pmatrix} 0.8 & 0.2 \\ 0.4 & 0.6 \\ 0 & 1 \end{pmatrix}, \quad x_t = (\tfrac{1}{2}, \tfrac{1}{4}, \tfrac{1}{4}),$$

231 so $q_t(s_1) = \frac{1}{2}(0.8) + \frac{1}{4}(0.4) + \frac{1}{4}(0) = 0.5$ and $q_t(s_2) = 0.5$. The third action cannot produce
 232 s_1 at all. Suppose the learner draws a_1 , observes s_1 , and receives $X_t = 1$ at time $t + d$. Then
 233 $\hat{c}_t(a) = P(s_1 | a)/0.5$, so $\hat{c}_t = (1.6, 0.8, 0)$ against the action-level $\tilde{c}_t = (2, 0, 0)$. One observation
 234 updates all three actions: a_2 is charged although it was never played, because it had a 0.4 chance of
 235 producing the state seen, and a_3 is charged nothing, because it could not have produced that state.
 236 Nothing is biased by this. With $\theta_t = (1, 0)$ the true losses are $c_t = (0.8, 0.4, 0)$, and averaging $\hat{c}_t$
 237 over $S_t \sim q_t$ returns exactly that. What is bought is the second moment: $v_t = 1.44$ here, against
 238 $K = 3$ for the action-level estimate, and with $\theta_t = (1, 1)$ both are attained exactly.

239 *Remark 3* (The χ^2 form). Under x_t the pair (A_t, S_t) has law $x_t(a)P(s | a)$ with marginals x_t and
 240 q_t , so

$$\begin{aligned} \chi^2(P_{A_t, S_t} \| P_{A_t} \otimes P_{S_t}) &= \sum_{a, s} \frac{(x_t(a)P(s | a) - x_t(a)q_t(s))^2}{x_t(a)q_t(s)} \\ &= \sum_{a, s} \frac{x_t(a)P(s | a)^2}{q_t(s)} - 2 + 1 = v_t - 1, \end{aligned}$$

241 the cross term summing to $\sum_{a, s} x_t(a)P(s | a) = 1$ and the last to $\sum_{a, s} x_t(a)q_t(s) = 1$. So
 242 $v_t = 1$ exactly when A_t and S_t are independent, which pins the rows of P only on the support
 243 of x_t and therefore means all rows agree whenever x_t has full support, as exponential weights
 244 does. If P is deterministic and every state is reached by some action in the support of x_t , then v_t
 245 equals the number of states that support reaches. In particular, under full-support play and an onto
 246 deterministic P , $v_t = |\mathcal{S}|$. This is an identity, not a bound; it is used for reading v_t , and once in the
 247 proof of Theorem 1.

248 **Proposition 4.** Fix t , let x_t be the play distribution, $q_t(s) = \sum_a x_t(a)P(s | a)$ the induced state
 249 distribution, and $\hat{c}_t(a) = P(S_t | a)X_t/q_t(S_t)$, which is well defined because only states with
 250 $q_t(s) > 0$ occur. For statements involving $1/x_t(a)$, restrict to actions with $x_t(a) > 0$; STATE-EXP3
 251 itself has full support on every action. Then $\hat{c}_t(a) \geq 0$, $\hat{c}_t(a) \leq 1/x_t(a)$, $\mathbb{E}[\hat{c}_t(a) | \mathcal{F}_{t-1}] = c_t(a)$,
 252 and

$$\sum_a x_t(a)\mathbb{E}[\hat{c}_t(a)^2 | \mathcal{F}_{t-1}] \leq v_t := \sum_{s: q_t(s) > 0} \frac{\sum_a x_t(a)P(s | a)^2}{q_t(s)}, \quad 1 \leq v_t \leq |\mathcal{S}|,$$

253 states of zero probability being omitted throughout, which is the convention $0/0 = 0$. The action-
 254 level $\tilde{c}_t(a) = \mathbf{1}\{A_t = a\}X_t/x_t(a)$ has the same first three properties and weighted second moment
 255 at most K .

256 *Proof.* Nonnegativity is immediate from $X_t \geq 0$. Given $\mathcal{F}_{t-1}$ the state S_t has distribution q_t , and
 257 $\mathbb{E}[X_t \mid \mathcal{F}_{t-1}, S_t = s] = \theta_t(s)$ by the model, since that conditional mean does not depend on
 258 the action, so $\mathbb{E}[\hat{c}_t(a) \mid \mathcal{F}_{t-1}] = \sum_s q_t(s) P(s \mid a) \theta_t(s) / q_t(s) = c_t(a)$. For the range, $q_t(s) \geq$
 259 $x_t(a) P(s \mid a)$ for every a , so $P(s \mid a) / q_t(s) \leq 1 / x_t(a)$ and $X_t \leq 1$. For the second moment,
 260 $X_t^2 \leq 1$ gives $\mathbb{E}[\hat{c}_t(a)^2 \mid \mathcal{F}_{t-1}] \leq \sum_s P(s \mid a)^2 / q_t(s)$, and exchanging the sums gives v_t . Each
 261 summand of v_t is at most 1, since $P(s \mid a) \leq 1$ makes $\sum_a x_t(a) P(s \mid a)^2 \leq q_t(s)$, and at least
 262 $q_t(s)$, since Jensen applied to $a \mapsto P(s \mid a)$ under x_t gives $\sum_a x_t(a) P(s \mid a)^2 \geq q_t(s)^2$. Summing
 263 the two bounds gives $1 \leq v_t \leq |\mathcal{S}|$. The action-level claims are the same computation with the state
 264 replaced by the played action, whose weighted second moment is $\sum_a x_t(a) / x_t(a) = K$. $\square$

265 The two ends are attained. If every action induces the same state distribution then $v_t = \sum_s q_t(s) =$
 266 1, and if P is deterministic, each action reaching a single state, then every $P(s \mid a)$ is 0 or 1 and
 267 each summand is exactly 1, so v_t counts the states reached and equals $|\mathcal{S}|$ when every state has an
 268 action reaching it. Disjoint row supports give $v_t = K$ rather than $|\mathcal{S}|$, since the summands over
 269 one action's private states add to $\sum_s P(s \mid a) = 1$, so a sensor with many private states per action
 270 resolves few of them. So v_t reads the overlap between the rows of P rather than their length, and
 271 $\bar{v} = \sup_x v(x)$, the supremum over play distributions, depends on P alone. An outcome observed
 272 at a state updates the estimate for every action that can reach it, which is why the construction of
 273 Section 4 gives each drifting coordinate a private state, forcing $v_t = |\mathcal{S}|$ there. The immediately
 274 observed pairs (A_t, S_t) identify P separately, though Theorem 5 assumes it known.

275 D Proofs: pooling and grouping

276 Two results are stated here rather than in the body, since Theorem 1 covers the algorithm actually
 277 run. The first is the interleaved variant.

278 **Theorem 5.** *Let P be known and (θ_t) oblivious. Pick before the run any fixed number V with*
 279 *$\sum_t v_t \leq V$ always. Then STATE-EXP3 with $m = d + 1$ and the matching learning rate satisfies*
 280 *$\mathbb{E}[R_T] \leq \sqrt{2(d+1)V \log K}$. Writing $\bar{v}$ for the largest v_t over all ways of playing, $V = \bar{v}T \leq |\mathcal{S}|T$*
 281 *is one such choice, and it depends on P alone.*

282 The second merges states that are not worth telling apart. Group them with a map g from $\mathcal{S}$ onto
 283 a smaller set $\mathcal{Z}$ and run the same algorithm on $g(S_t)$, where $P^g(z \mid a) = \sum_{s: g(s)=z} P(s \mid a)$.
 284 Grouping never raises the dimension, but the estimate now tracks a group average, so the widest
 285 spread of losses inside a group, $\delta_t(g) = \max_z \max_{s, s': g(s)=g(s')=z} |\theta_t(s) - \theta_t(s')|$, is the error
 286 introduced.

287 **Theorem 6.** *Under the hypotheses of Theorem 5, running on g with $\sum_t v_t(g) \leq V(g)$ gives*
 288 *$\mathbb{E}[R_T] \leq \sqrt{2(d+1)V(g) \log K} + 2 \sum_t \delta_t(g)$.*

289 Merging need not preserve state-sufficiency, so the grouped estimate is right not for c_t but for a stand-
 290 in loss differing from it by at most $\delta_t(g)$ on every action. Keeping states apart preserves differences
 291 that matter but learns more slowly; merging shares more but pays $2 \sum_t \delta_t(g)$. The best grouping is
 292 the true one: with 16 states from four hidden groups, sizes 1, 2, 4, 8, 16 give regret 1590, 831, 754,
 293 920, 1080, lowest at four, which no algorithm was told. The map g is fixed in advance; learning it
 294 from data remains open.

295 **Proof idea.** The pooled estimator is unbiased for every action, while Proposition 4 bounds its play-
 296 weighted second moment by v_t rather than K . With $m = d + 1$ interleaved copies, the outcome
 297 generated on a round belonging to one copy is available before that copy acts again, so each copy
 298 is an ordinary exponential-weights process with immediate feedback. Summing the $d + 1$ potential
 299 inequalities gives $(d + 1) \log K / \eta$, and the pooled second moments give $\frac{\eta}{2} \sum_t v_t$.

300 *Proof of Theorem 5.* Fix a copy i and let T_i be the number of rounds it owns, so $\sum_i T_i = T$. Round
 301 t 's outcome is usable from round $t + d + 1$, which with $m = d + 1$ is exactly $t + m$, the copy's next
 302 round, so within a copy the feedback is undelayed and L_i carries exactly the estimates of that copy's
 303 earlier rounds. Exponential weights on nonnegative losses give, for any action a^* and any $\eta > 0$,

$$\sum_{t \in i} (\langle x_t, \hat{c}_t \rangle - \hat{c}_t(a^*)) \leq \frac{\log K}{\eta} + \frac{\eta}{2} \sum_{t \in i} \sum_a x_t(a) \hat{c}_t(a)^2,$$

by the usual potential telescoping with $e^{-z} \leq 1 - z + z^2/2$, which needs $z \geq 0$ and therefore $\hat{c}_t \geq 0$, supplied by Proposition 4. Each x_t is $\mathcal{F}_{t-1}$ -measurable, because L_i then holds only rounds $r \leq t - m = t - d - 1$, so taking expectations turns the left side into $\sum_{t \in i} (\langle x_t, c_t \rangle - c_t(a^*))$ by unbiasedness and the quadratic term into at most $\mathbb{E}[\sum_{t \in i} v_t]$, again by Proposition 4. Summing over the $m = d+1$ copies and using $\sum_i \min_a \sum_{t \in i} c_t(a) \leq \min_a \sum_t c_t(a)$, which is what lets the copies be compared against one common action, gives $\mathbb{E}[R_T] \leq (d+1) \log K/\eta + \frac{\eta}{2} \mathbb{E}[\sum_t v_t]$. If $\sum_t v_t \leq V$ surely then $\eta = \sqrt{2(d+1) \log K/V}$ balances the two terms and gives $\sqrt{2(d+1)V \log K}$, and $V = \bar{v}T$ is admissible by the definition of $\bar{v}$. $\square$

The delay enters as a factor because the copies do not share their estimates, each paying $\log K/\eta$ over $T/(d+1)$ rounds. Sharing them is exactly the $m = 1$ algorithm of the open problem of Section 6.

Lemma 7 (Coarsening). *For every play distribution and every g , $v_t(g) \leq v_t$.*

Proof. It suffices to merge two states, the general case following by induction on the merges, and a merged group of zero probability contributes nothing on either side by the convention above. Write P_1, P_2 for the two columns, $A_j = \sum_a x_t(a) P_j(a)^2$, $B_j = \sum_a x_t(a) P_j(a)$ and $\lambda = B_1/(B_1 + B_2)$. By Cauchy–Schwarz, $(P_1 + P_2)^2 \leq \lambda^{-1} P_1^2 + (1-\lambda)^{-1} P_2^2$ pointwise in a , so $\sum_a x_t(a) (P_1 + P_2)^2 \leq (B_1 + B_2)(A_1/B_1 + A_2/B_2)$, and dividing by the merged denominator $B_1 + B_2$ shows the merged summand is at most $A_1/B_1 + A_2/B_2$, the two it replaces. $\square$

Proof of Theorem 6. Running the algorithm on g is running it on the sensor (Z, P^g) . The surrogate losses below depend on x_t and so are predictable rather than oblivious, but the proof of Theorem 5 uses obliviousness nowhere, only that each c_t^g is $\mathcal{F}_{t-1}$ -measurable given the play, so it applies unchanged and bounds regret measured in the surrogate losses $c_t^g(a) = \sum_z P^g(z | a) \bar{\theta}_t(z)$, where $\bar{\theta}_t(z) = \sum_{s: g(s)=z} q_t(s) \theta_t(s) / q_t^g(z)$ is the group mean the grouped estimator is unbiased for, by the same computation as in Proposition 4 with S_t replaced by $g(S_t)$. Fix a and z . Both $\{q_t(s)/q_t^g(z)\}$ and $\{P(s | a)/P^g(z | a)\}$ are probability vectors on $\{s : g(s) = z\}$, so their θ_t -averages lie in a common interval of length $\delta_t(g)$ and differ by at most $\delta_t(g)$. Weighting by $P^g(z | a)$ and summing over z gives $|c_t^g(a) - c_t(a)| \leq \delta_t(g)$ for every a , hence $\langle x_t, c_t \rangle - c_t(a^*) \leq \langle x_t, c_t^g \rangle - c_t^g(a^*) + 2\delta_t(g)$. Summing over t adds $2 \sum_t \delta_t(g)$, and Lemma 7 makes $V(g)$ no larger than a bound valid for the raw states. $\square$

E An unknown action-to-state map

The plug-in estimator. The pairs (A_t, S_t) arrive without delay, so P is estimable from the learner’s own play while the outcomes are still in flight. Let $N_a(t)$ count the rounds before t that played a and $N_{a,s}(t)$ those that also produced s , and set $\hat{P}_t(s | a) = (N_{a,s}(t) + \alpha)/(N_a(t) + \alpha|S|)$ for a smoothing constant $\alpha \in (0, 1]$ fixed independently of T , which is row-stochastic with every entry positive, so $\hat{q}_t = x_t^\top \hat{P}_t$ is positive too, and both are $\mathcal{F}_{t-1}$ -measurable. Smoothing moves an entry from the empirical row by at most $\alpha(|S| - 1)/(N_a(t) + \alpha|S|)$, attained when every observation of a sits on one state, and moves the row in ℓ_1 by at most twice that. Both are of lower order than the sampling term at fixed α and must be carried otherwise, and Appendix I shows the choice matters in practice. The plug-in estimate is $\hat{c}_t(a) = \hat{P}_t(S_t | a) X_t / \hat{q}_t(S_t)$. Write $\varepsilon_t = \max_a \|\hat{P}_t(\cdot | a) - P(\cdot | a)\|_1$ for its accuracy and $\rho_t = \max_{s: q_t(s) > 0} |\hat{q}_t(s) - q_t(s)|/q_t(s)$ for the relative error it induces on the state distribution, states of zero probability under the play being omitted since the estimator never charges them.

Proposition 8. *On $\{\rho_t \leq \frac{1}{2}\}$ the plug-in estimate satisfies $\hat{c}_t \geq 0$ and $\sum_a x_t(a) \mathbb{E}[\hat{c}_t(a)^2 | \mathcal{F}_{t-1}] \leq 2\hat{v}_t \leq 2|S|$, and its conditional mean $\bar{c}_t$ satisfies $\sum_a x_t(a) (\bar{c}_t(a) - c_t(a)) = 0$ exactly, for every $\hat{P}_t$ and with no hypothesis at all, and $|\bar{c}_t(a) - c_t(a)| \leq 2\varepsilon_t(a) + 2\rho_t$ for every a , where $\varepsilon_t(a) = \|\hat{P}_t(\cdot | a) - P(\cdot | a)\|_1$.*

Proof. Write $\Delta_t = \hat{P}_t - P$ and $\delta_t = \hat{q}_t - q_t = \sum_a x_t(a) \Delta_t(\cdot | a)$. Substituting these into $\bar{c}_t(a) = \sum_s q_t(s) \hat{P}_t(s | a) \theta_t(s) / \hat{q}_t(s)$ and cancelling gives $\bar{c}_t(a) - c_t(a) = \sum_s \theta_t(s) [q_t(s) \Delta_t(s |$

352 $a) - P(s | a)\delta_t(s)]/\hat{q}_t(s)$, so with $\theta_t \in [0, 1]^{|S|}$,

$$|\bar{c}_t(a) - c_t(a)| \leq \sum_s \frac{q_t(s)|\Delta_t(s | a)| + P(s | a)|\delta_t(s)|}{\hat{q}_t(s)}.$$

353 The play-weighted identity needs no hypothesis. By definition $\sum_a x_t(a)\hat{P}_t(s | a) = \hat{q}_t(s)$, so
 354 $\sum_a x_t(a)\hat{c}_t(a) = X_t$ whatever $\hat{P}_t$ is, and taking conditional expectations gives $\sum_a x_t(a)\bar{c}_t(a) =$
 355 $\sum_s q_t(s)\theta_t(s) = \langle x_t, c_t \rangle$. On $\{\rho_t \leq \frac{1}{2}\}$ we have $q_t \leq 2\hat{q}_t$, and for the per-action bias the
 356 first sum of the display is at most $2\|\Delta_t(\cdot | a)\|_1 = 2\varepsilon_t(a)$ and the second at most $2\sum_s P(s |$
 357 $a)|\delta_t(s)|/q_t(s) \leq 2\rho_t$, because $P(\cdot | a)$ is a probability vector. The play-weighted identity
 358 does *not* survive reweighting: at $x_t = (\frac{1}{2}, \frac{1}{2})$, $P = I$, $\hat{P} = (\frac{.9}{.2} \frac{.1}{.8})$ and $\theta_t = (0, 1)$ the
 359 bias is $(\frac{1}{9}, -\frac{1}{9})$ with zero x_t -average, while weights $(1, 0)$ leave $\frac{1}{18}$. For the second moment,
 360 $\sum_a x_t(a)\mathbb{E}[\hat{c}_t(a)^2 | \mathcal{F}_{t-1}] \leq \sum_s \frac{q_t(s)}{\hat{q}_t(s)} \cdot \frac{\sum_a x_t(a)\hat{P}_t(s|a)^2}{\hat{q}_t(s)} \leq 2\hat{v}_t$, and $\hat{v}_t \leq |S|$ by Proposition 4
 361 applied to $\hat{P}_t$. $\square$

362 **The warm-start guarantee.** WARM-START STATE-EXP3 plays uniformly for n_0 rounds, forms
 363 $\hat{P} = \hat{P}_{n_0+1}$ from those rounds alone, discards the cumulative estimates together with every outcome
 364 of a warm-start round that has not yet arrived, and then runs STATE-EXP3 with $m = d + 1$, that
 365 frozen $\hat{P}$, and each play mixed with the uniform action at rate γ . Freezing makes $\hat{P}$ measurable with
 366 respect to the first phase, so ε and ρ are fixed before the second phase begins, and discarding is what
 367 keeps the first phase's inaccurate updates out of every later distribution.

368 **Theorem 9.** Suppose $\bar{p}(s) \geq 1/(\kappa|S|)$ for every s , fix $0 < \gamma \leq \frac{1}{2}$ and $1 \leq d \leq T/2$, and take
 369 $n_0 \geq c\gamma^{-2}K\kappa^2|S|^2 \log(K|S|T)$ for a universal c . Then WARM-START STATE-EXP3 satisfies, for
 370 every $\eta > 0$,

$$\mathbb{E}[R_T] \leq n_0 + 1 + \gamma T + \frac{(d+1)\log K}{\eta} + 2\eta|S|T + T(4\bar{\varepsilon} + 4\bar{\rho}),$$

371 where $\bar{\varepsilon} = \tilde{O}(\sqrt{K|S|/n_0})$ and $\bar{\rho} = \tilde{O}(\kappa|S|\sqrt{K/n_0}/\gamma) \leq \frac{1}{2}$ bound ε and ρ on an event of proba-
 372 bility at least $1 - 1/T$. Tuning η and then $\gamma \asymp (\kappa|S|)^{1/2}(K/n_0)^{1/4}$ and $n_0 \asymp T^{4/5}(\kappa|S|)^{2/5}K^{1/5}$
 373 gives $\mathbb{E}[R_T] = \tilde{O}(\sqrt{d|S|T \log K} + (\kappa|S|)^{2/5}K^{1/5}T^{4/5})$, provided $T = \tilde{\Omega}(K(\kappa|S|)^2)$. That con-
 374 dition constrains T against K , $|S|$ and κ only, so the hypothesis $d \leq T/2$ is what carries the delay.
 375 Together they give $n_0 \leq T/2 \leq T - d$, and the same condition supplies the lower bound above with
 376 some $\gamma \leq \frac{1}{2}$. Outside that regime the tuning is vacuous and $R_T \leq T$ is the statement.

377 *Proof.* The first phase has losses in $[0, 1]$, so it contributes at most n_0 . Let G be the event that
 378 $N_a(n_0 + 1) \geq n_0/(2K)$ for every a , that $\|\hat{P}(\cdot | a) - P(\cdot | a)\|_1 \leq \bar{\varepsilon}$ for every a , and that
 379 $\max_{s,a} |\hat{P}(s | a) - P(s | a)| \leq \bar{\rho}\gamma/(\kappa|S|)$. Uniform play gives $\mathbb{E}[N_a] = n_0/K$, so a multiplicative
 380 Chernoff bound, the ℓ_1 deviation bound for a multinomial, Hoeffding, and a union bound over the
 381 $K|S|$ entries put $\Pr(G^c) \leq 1/T$ at the stated n_0 , the pseudo-counts adding $2\alpha(|S| - 1)/(N_a + \alpha|S|)$
 382 to each ℓ_1 deviation and half that to each entry, of lower order at fixed α . Off G bound the second
 383 phase by T , contributing at most 1.

384 On G the quantity $\hat{P}$ is measurable with respect to the first phase, and mixing gives $q_t(s) \geq \gamma\bar{p}(s) \geq$
 385 $\gamma/(\kappa|S|)$ at every later round, so $\rho_t \leq \kappa|S|\max_{s,a} |\hat{P}(s | a) - P(s | a)|/\gamma \leq \bar{\rho} \leq \frac{1}{2}$ for every
 386 $t > n_0$ and every realization of the second phase. Proposition 8 therefore applies at every such
 387 round, and conditionally on the first phase the second is exactly the algorithm of Theorem 5 driven
 388 by a nonnegative estimate with weighted second moment at most $2\hat{v}_t \leq 2|S|$. Its proof gives
 389 $\mathbb{E}[\sum_{t>n_0} (\langle p_t, \bar{c}_t \rangle - \bar{c}_t(a^*))] \leq (d+1)\log K/\eta + 2\eta|S|T$, the extra factor of two coming from
 390 $p_t \leq x_t/(1-\gamma) \leq 2x_t$. Playing x_t rather than p_t costs at most γ per round since $c_t \in [0, 1]^K$, so
 391 it remains to pass from $\bar{c}_t$ to c_t under p_t . Proposition 8 gives $\langle x_t, b_t \rangle = 0$ for the per-action bias b_t ,
 392 and $x_t = (1-\gamma)p_t + \gamma u$ for the uniform u , so $\langle p_t, b_t \rangle = -\frac{\gamma}{1-\gamma}\langle u, b_t \rangle$, which for $\gamma \leq \frac{1}{2}$ is at most
 393 $\max_a |b_t(a)|$ in absolute value. Each round therefore adds $|\langle p_t, b_t \rangle| + |b_t(a^*)| \leq 2\max_a |b_t(a)| \leq$
 394 $4\bar{\varepsilon} + 4\bar{\rho}$. Collecting the four contributions gives the display. For the rate, balancing γT against $2T\bar{\rho}$
 395 gives the stated γ and a combined $\tilde{O}(T(\kappa|S|)^{1/2}(K/n_0)^{1/4})$, and balancing that against n_0 gives
 396 the stated n_0 , under which $T\bar{\varepsilon}$ and the constraint on n_0 are of lower order. $\square$

397 **Safe blending.** Blending the plug-in with the action-level estimate makes the guarantee safe. Let
 398 $\lambda_t \in [0, 1]^K$ be $\mathcal{F}_{t-1}$ -measurable and put $\check{g}_t(a) = (1 - \lambda_t(a))\bar{c}_t(a) + \lambda_t(a)\hat{c}_t(a)$.

399 **Proposition 10.** $\check{g}_t \geq 0$, its conditional mean satisfies $\mathbb{E}[\check{g}_t(a) \mid \mathcal{F}_{t-1}] - c_t(a) = \lambda_t(a)(\bar{c}_t(a) -$
 400 $c_t(a))$ for the plug-in's conditional mean $\bar{c}_t$ of Proposition 8, and

$$\sum_a x_t(a) \mathbb{E}[\check{g}_t(a)^2 \mid \mathcal{F}_{t-1}] \leq (K - \Pi_t) + \hat{V}_t,$$

401 for $\Pi_t = \sum_a \lambda_t(a)$ and $\hat{V}_t = \sum_a \lambda_t(a) x_t(a) \mathbb{E}[\hat{c}_t(a)^2 \mid \mathcal{F}_{t-1}]$, which satisfies $\hat{V}_t \leq 2\hat{v}_t$ on $\{\rho_t \leq$
 402 $\frac{1}{2}\}$ and not in general, since the step $q_t/\hat{q}_t \leq 2$ behind it is exactly that event. We write $V_t^{\text{ub}}(\tau)$
 403 below for a computable upper bound on $\hat{V}_t$, reserving $\hat{V}_t$ for the conditional expectation itself.

404 *Proof.* Both parts are nonnegative and $\lambda_t(a) \in [0, 1]$. The mean identity is linearity together
 405 with $\mathbb{E}[\bar{c}_t(a) \mid \mathcal{F}_{t-1}] = c_t(a)$. For the second moment, $u \mapsto u^2$ is convex, so $\check{g}_t(a)^2 \leq$
 406 $(1 - \lambda_t(a))\bar{c}_t(a)^2 + \lambda_t(a)\hat{c}_t(a)^2$, and $x_t(a)\mathbb{E}[\bar{c}_t(a)^2 \mid \mathcal{F}_{t-1}] \leq 1$ by Proposition 4. $\square$

407 So a weight of $\lambda_t(a)$ on the plug-in buys a variance saving of $\lambda_t(a)$ and pays a bias of $\lambda_t(a)$ times the
 408 plug-in's, and both endpoints are exact. Let $r_t(a)$ and $\hat{\rho}_t$ be any sequences computable from the im-
 409 mediate pairs that are *simultaneously valid*, meaning $\varepsilon_t(a) \leq r_t(a)$ for every a and $\rho_t \leq \hat{\rho}_t$ at every
 410 $t \leq T$, on an event of probability at least $1 - 1/T$. Write $V_t^{\text{ub}}(\tau) = 2 \sum_{a: r_t(a) \leq \tau} x_t(a) \sum_s \hat{P}_t(s \mid$
 411 $a)^2 / \hat{q}_t(s)$, which the learner can evaluate and which dominates $\hat{V}_t$ on $\{\rho_t \leq \frac{1}{2}\}$. SAFE-POOL
 412 STATE-EXP3 sets $\lambda_t \equiv 0$ when $\hat{\rho}_t > \frac{1}{2}$, and otherwise picks the best member of the candidate
 413 set $\{\perp\} \cup \{r_t(a) : a \leq K\}$, where $\perp$ stands for $\lambda_t \equiv 0$ and a numerical candidate τ stands for
 414 $\lambda_t(a) = \mathbf{1}\{r_t(a) \leq \tau\}$, best meaning smallest

$$\Psi_t(\tau) = \frac{\eta}{2} [(K - \Pi_t(\tau)) + V_t^{\text{ub}}(\tau)] + (4\tau + 4\hat{\rho}_t) \mathbf{1}\{\Pi_t(\tau) > 0\}, \quad \Psi_t(\perp) = \frac{\eta}{2} K.$$

415 Write Ψ_t^* for that smallest value. The candidate $\perp$ is listed apart from the numerical thresholds
 416 because a valid radius may equal zero, in which case $\tau = 0$ pools the actions attaining it rather than
 417 abstaining, so $\perp$ rather than $\tau = 0$ is what keeps $\frac{\eta}{2} K$ available on every round.

418 **Theorem 11.** For every $\eta > 0$, SAFE-POOL STATE-EXP3 with $m = d + 1$ satisfies

$$\mathbb{E}[R_T] \leq 1 + \frac{(d+1) \log K}{\eta} + \mathbb{E}\left[\sum_{t=1}^T \Psi_t^*\right],$$

419 so $\eta = \sqrt{2(d+1) \log K / (KT)}$ gives $\mathbb{E}[R_T] \leq 1 + \sqrt{2(d+1)KT \log K}$ on every instance, which
 420 is the guarantee of action-level exponential weights, with an improvement on each round whose
 421 Ψ_t^* is strictly below $\frac{\eta}{2} K$. The certificate is computed from the learner's own observations and so
 422 is random, which is why the display carries an expectation. The tuned bound does not, because
 423 $\Psi_t^* \leq \frac{\eta}{2} K$ holds on every path.

424 *Proof.* On the validity event, of probability at least $1 - 1/T$ and contributing at most 1 off it, pooling
 425 happens only when $\hat{\rho}_t \leq \frac{1}{2}$ and hence only when $\rho_t \leq \frac{1}{2}$, which is what makes V_t^{ub} dominate the
 426 conditional second moment. Proposition 10 then supplies a nonnegative estimate whose weighted
 427 second moment is at most $(K - \Pi_t) + \hat{V}_t$, so the argument of Appendix D applies to the sur-
 428rogate losses. Passing to c_t adds both the played bias $|\sum_a x_t(a) \lambda_t(a) b_t(a)|$ and the comparator
 429 bias $\lambda_t(a^*) |b_t(a^*)|$, for $b_t = \bar{c}_t - c_t$. The cancellation of Proposition 8 is lost under per-action
 430 weights, as the counterexample there shows, so each is bounded separately. On a pooled action
 431 $\varepsilon_t(a) \leq r_t(a) \leq \tau$ gives $|b_t(a)| \leq 2\tau + 2\hat{\rho}_t$, and since $\sum_a x_t(a) \lambda_t(a) \leq 1$ and $\lambda_t(a^*) \leq 1$ the two
 432 together are at most $4\tau + 4\hat{\rho}_t$, charged only when something is pooled. The second-moment bound
 433 and the two bias terms are all functions of the observations, so the per-round costs they assemble
 434 into are the random Ψ_t^* and the sum is taken in expectation. For the tuned bound, $\perp$ belongs to the
 435 candidate set and $\Psi_t(\perp) = \frac{\eta}{2} K$, so $\Psi_t^* \leq \frac{\eta}{2} K$ on every path whatever the radii are. $\square$

436 **Why certification stays conservative.** The guarantee is safe and the algorithm is inert. Pooling
 437 is certified when $4(\tau + \hat{\rho}_t) \leq \frac{\eta}{2} (\Pi_t - V_t^{\text{ub}})$. At the tuned η both sides fall as $T^{-1/2}$, the left through
 438 $\tau \asymp \sqrt{K \log(K|\mathcal{S}|T)/T}$ under uniform play and the right through $\eta \asymp \sqrt{\log K / (KT)}$. Their ratio

is $\Theta(\sqrt{\log(K|\mathcal{S}|T)/\log K})$, which exceeds one and does not vanish with the horizon, growing slowly in T through the logarithmic factors, so no horizon makes the certificate clear. Appendix I reports behavior consistent with this explanation, the algorithm reproducing action-level regret to the digit with a pooled fraction of zero.

The obstruction is not the shape of the bias bound. Regret charges $\langle x_t, b_t \rangle - b_t(a^*)$ for the per-action bias b_t , while Proposition 8 charges $\max_a |b_t(a)|$. Measured against the budget itself at $K = 40$, $|\mathcal{S}| = 8$, $T = 8000$ and $d = 50$, the first is 0.42 times it and the second 1.67 times it, so on that trajectory the realized bias lies below the best-case variance budget and a tighter bias bound would close the analytical gap. Two exact refinements fail to deliver one. Both $\sum_s \Delta_t(s | a)$ and $\sum_s \delta_t(s)$ vanish, so a constant may be subtracted and a maximum replaced by an oscillation, and Cauchy-Schwarz against q_t replaces the comparator term by $\sqrt{\nu_t(a^*)} \chi_t$ for its own effective dimension ν_t and a chi-square χ_t between $\hat{q}_t$ and q_t . Measured, these are 6.1 and 63 times the truth against 3.9 for the maximum they replace.

What the algorithm can verify from its own counts is 104 times the budget, and Bernstein radii do not improve on Hoeffding ones (Appendix I), so the barrier is certifying ρ_t , the relative error the estimate induces on the state distribution. The play-weighted identity of Proposition 8 says the whole charge is the comparator's, and it routes through ρ_t only because $\hat{q}_t = x_t^\top \hat{P}_t$ predicts the state distribution out of every row of P . Predicting is expensive for that reason, needing $t = \Omega(K|\mathcal{S}| \log(K|\mathcal{S}|T)/\rho^2)$ rounds, and the resulting certificate clears for no $K \leq T$, since it would need $\log K \geq 16c|\mathcal{S}|$ for $c = \log(2|\mathcal{S}|T)$, which $\log K \leq \log T$ cannot meet.

Counting the states costs less and still buys nothing. With x_t frozen over a window the empirical state frequency is a plain multinomial estimate of q_t , needing $n = \Omega(|\mathcal{S}| \log(|\mathcal{S}|T)/\rho^2)$ samples, a factor K fewer than predicting and free of P , and the certificate does clear once $K \log K \geq 64c|\mathcal{S}|$, about $K \geq 10^3$ at $|\mathcal{S}| = 8$. The window that buys it must satisfy $n \geq 32c|\mathcal{S}|T/(K \log K)$, a constant fraction of the horizon, and freezing the play over blocks of length B costs $\sqrt{2TBK \log K}$, which at that B is $T\sqrt{128c|\mathcal{S}|}$ and so linear in T . Both routes fail for one reason: the certificate must resolve q_t to relative accuracy $\Theta(\sqrt{K \log K/T})$, finer than either estimate delivers from the data the learner has.

In words, an unknown action-to-state matrix costs a bias in two accuracies, the ℓ_1 error of $\hat{P}$ and the relative error it induces on the state distribution, and only the second needs forced exploration, which is what sets the rate. **The bound does not certify an improvement from the channel it analyzes.** Its second term exceeds $\sqrt{KT}$ at every horizon with $T > K$, by a factor $(T/K)^{3/10}$, so a learner holding this guarantee alone should run action-level exponential weights instead. The experiments suggest the analysis is conservative rather than the algorithm weak, since Appendix I shows it tracking known P closely, and sharpening it is the natural next result.

F Feedback models, and what transfers

Write $\Lambda_2 = \sum_t \min\{1, \|c_t - m_t\|_2^2\}$ for the inaccuracy measure of Bar-On and Mansour (2025). Unlike E_2 , which charges only the worst action in a round, Λ_2 grows with how many actions are predicted wrongly, and since $\|v\|_\infty \leq \|v\|_2 \leq \sqrt{J}\|v\|_\infty$ for a vector v with J nonzero coordinates, $E_2 \leq \Lambda_2 \leq J E_2$ with both ends attained.

Bar-On and Mansour (2025) analyze a model in which the played coordinate of the loss vector is observed exactly. Section 3 instead observes a random outcome whose conditional mean is the loss. The following records precisely what survives the substitution.

Proposition 12. *Take the action-level $\tilde{c}_t(a) = \mathbf{1}\{A_t = a\}X_t/x_t(a)$ of Proposition 4 under the model of Section 3. Then for every action a , (i) $\mathbb{E}[\tilde{c}_t(a) | \mathcal{F}_{t-1}] = c_t(a)$, (ii) $0 \leq \tilde{c}_t(a) \leq 1/x_t(a)$, and (iii) $\mathbb{E}[\tilde{c}_t(a)^2 | \mathcal{F}_{t-1}] \leq 1/x_t(a)$. The centered estimator $\tilde{c}_t(a) = m_t(a) + \mathbf{1}\{A_t = a\}(X_t - m_t(a))/x_t(a)$ also satisfies (i), and its residual obeys $\mathbb{E}[(\tilde{c}_t(a) - m_t(a))^2 | \mathcal{F}_{t-1}] \leq 1/x_t(a)$.*

Proof. For (i), by the tower rule $\mathbb{E}[\mathbf{1}\{A_t = a\}X_t | \mathcal{F}_{t-1}] = x_t(a) \mathbb{E}[X_t | A_t = a] = x_t(a)c_t(a)$, and dividing by $x_t(a)$ gives the claim. Item (ii) holds because $X_t \in [0, 1]$. For (iii), $\mathbb{E}[\tilde{c}_t(a)^2 | \mathcal{F}_{t-1}] = \mathbb{E}[X_t^2 \mathbf{1}\{A_t = a\}]/x_t(a)^2 = \mathbb{E}[X_t^2 | A_t = a]/x_t(a) \leq 1/x_t(a)$ because $X_t \leq 1$. The statements for $\tilde{c}_t$ follow the same computation with X_t replaced by $X_t - m_t(a)$, using $|X_t - m_t(a)| \leq 1$. $\square$

491 In words, replacing an exactly observed loss by a bounded unbiased draw of it changes neither the
 492 mean of the estimator nor the moment bounds an analysis consumes.

493 Under exact feedback the corresponding estimator satisfies (i) to (iii) with identical bounds, since
 494 $c_t(a) \leq 1$. Consequently any step of an exact-feedback analysis that uses the observation only
 495 through unbiasedness and these two moment bounds is unchanged under noisy outcomes. What this
 496 does not cover is the comparison between the estimate built from the stale action-loss vector and the
 497 estimate built from the true loss. Under exact feedback that difference is determined by the played
 498 action. Under noisy outcomes it carries an additional mean-zero term of magnitude up to $1/x_t(A_t)$,
 499 and bounding its cumulative effect over a window of d rounds is what requires the stability property
 500 falsified in Appendix H. An upper bound for the noisy model is Theorem 5. What remains open is
 501 one adaptive to E_2 .

502 G Proofs: the lower bound

503 The measure E_2 is controlled by the state-loss variation W , and no converse bound holds in general.

504 **Lemma 13** (Delay window). $E_2 \leq d^2 W$ for $W = \sum_t \|\theta_t - \theta_{t-1}\|_\infty^2$. The ratio $E_2/(d^2 W)$ tends
 505 to one when θ drifts monotonically in one coordinate at a constant rate and $T/d \rightarrow \infty$, the first d
 506 rounds contributing $\sum_{j < d} j^2$ rather than d^2 each under the convention $m_t = c_1$.

507 In words, accumulated squared prediction error is at most the state-loss variation inflated by the
 508 square of the delay, and no better bound in W alone is available.

509 *Proof of Lemma 13.* Throughout write $\Delta_j = \theta_j - \theta_{j-1}$, set $\theta_0 = \theta_1$, so $\Delta_1 = 0$, and read c_r as c_1 for
 510 $r \leq 0$, matching the convention $m_t = c_1$ for $t \leq d$ of Section 3. Then $c_t - c_{t-d} = P \sum_{j=t-d+1}^t \Delta_j$.
 511 Because the rows of P are probability vectors, averaging cannot increase the supremum norm, so
 512 $\|c_t - c_{t-d}\|_\infty \leq \sum_j \|\Delta_j\|_\infty$. By Cauchy-Schwarz applied to the d summands, $\|c_t - c_{t-d}\|_\infty^2 \leq$
 513 $d \sum_{j=t-d+1}^t \|\Delta_j\|_\infty^2$. Summing over t and using that each increment belongs to at most d windows
 514 gives $E_2 \leq d^2 W$. The ratio $E_2/(d^2 W)$ approaches one when every increment has the same mag-
 515 nitude, sign and active coordinate, and some action reaches that coordinate deterministically, since
 516 then no cancellation occurs within a full window and P passes the increment through undiminished.
 517 Without the second condition the drift can lie in the kernel of P and the ratio can be zero. $\square$

518 **Lemma 14** (Information isolation). If $h \leq d$ then A_t is independent of the signs governing its own
 519 block, and this remains true when the learner is provided with the exact stale action-loss vector m_t .

520 *Proof of Lemma 14.* For t in block b , every realized outcome usable when choosing A_t comes from
 521 a round $r \leq t - d - 1$. Since $t \leq bh$ and $h \leq d$, we have $r \leq bh - d - 1 < (b-1)h$, so every usable
 522 realized outcome depends only on signs from blocks strictly before b , whose joint law depends only
 523 on $\sigma_2, \dots, \sigma_{b-1}$. Hence by conditional independence A_t is independent of block b 's signs. For the
 524 oracle, take any $h \leq d$. If $t > d$ then $t - d \leq bh - d \leq (b-1)h$ by the same computation,
 525 so $m_t = c_{t-d}$ is measurable with respect to the signs of blocks strictly before b . If $t \leq d$ then
 526 $m_t = c_1 = \frac{1}{2}\mathbf{1}$, which is deterministic because block 1 is neutral. In both cases m_t is independent
 527 of block b 's signs, so adjoining it to the learner's information leaves the conclusion intact. $\square$

528 The construction in full. Take states $s_0, \dots, s_q$, a stationary deterministic map sending a_j to s_j for
 529 $j \leq J$ and every other action to s_0 , and independent Rademacher signs $\sigma_b^{(j)}$ for $2 \leq b \leq B = \lfloor T/h \rfloor$.
 530 Set $\theta_t(s_0) = \frac{1}{2}$ throughout, $\theta_t \equiv \frac{1}{2}$ on block 1, and $\theta_t(s_j) = \frac{1}{2} - \varepsilon \sigma_b^{(j)}$ on each later block b . The
 531 neutral first block makes $c_1 = \frac{1}{2}\mathbf{1}$, so the convention $m_t = c_1$ for $t \leq d$ of Section 3 supplies a
 532 constant over exactly the rounds that precede any usable feedback, which is what Lemma 14 needs
 533 at $b = 1$. Rounds after the last complete block carry no drift, costing a constant factor.

534 **Lemma 15** (Rademacher lower tail). Let S_B be a sum of $B \geq 32$ independent Rademacher signs.
 535 For every $1 \leq u \leq \sqrt{B}/32$, $\Pr(S_B \geq u\sqrt{B}) \geq \frac{u}{4} e^{-64u^2}$.

536 *Proof.* Write $S_B = 2X - B$ with $X \sim \text{Bin}(B, \frac{1}{2})$, so that $\{S_B \geq u\sqrt{B}\} = \{X \geq B/2 + u\sqrt{B}/2\}$.
 537 Passing to X removes the lattice gap, since X takes every integer value in $[0, B]$ whereas S_B takes

only those of the parity of B . Put $n = \lfloor B/2 \rfloor$ and $t = \lceil u\sqrt{B}/2 \rceil + 1$, so $X \geq n + t$ implies the event, and $t \leq u\sqrt{B}$ because $u\sqrt{B} \geq \sqrt{32} > 4$.

For $j \geq 1$ the ratio of consecutive binomial probabilities is $\Pr(X = n + j)/\Pr(X = n + j - 1) = (B - n - j + 1)/(n + j)$. Using $(B - 1)/2 \leq n \leq B/2$, its distance from one is at most $(2j - 1)/(n + j) \leq (4j - 2)/(B - 1) \leq 8j/B$ for $B \geq 2$. Restricting to $j \leq B/16$ keeps $8j/B \leq \frac{1}{2}$, where $\log(1 - x) \geq -2x$ holds, so summing over $i \leq j$ gives $\Pr(X = n + j) \geq \Pr(X = n)e^{-16j(j+1)/(2B)} \geq \Pr(X = n)e^{-16j^2/B}$. The central term satisfies $\Pr(X = n) \geq 1/(2\sqrt{B})$ for every $B \geq 2$.

Sum over the t integers $j \in \{t, \dots, 2t - 1\}$, all of which are at most $2t$ and hence at most $B/16$ because $u \leq \sqrt{B}/32$ forces $2t \leq 2u\sqrt{B} \leq B/16$. Each term is at least $e^{-64t^2/B}/(2\sqrt{B})$, and $t \leq u\sqrt{B}$ makes the exponent at most $64u^2$, so $\Pr(X \geq n + t) \geq te^{-64u^2}/(2\sqrt{B}) \geq \frac{u}{4}e^{-64u^2}$ using $t \geq u\sqrt{B}/2$. $\square$

Three regimes give the maximal inequality the lower bound needs. Each $S_B^{(j)}$ is sub-Gaussian with variance proxy B , so $\mathbb{E} \max_j (S_B^{(j)})^+ \leq \sqrt{2B \log J} + \sqrt{B}$ throughout. For the matching lower bound, first suppose $\log J \leq B/8$ and $\log J \geq 128$. Taking $u = \sqrt{\log J}/128$, which then lies in $[1, \sqrt{B}/32]$, makes $e^{-64u^2} = J^{-1/2}$, so $\Pr(S_B^{(j)} \geq u\sqrt{B}) \geq \frac{u}{4}J^{-1/2}$ and independence across j gives $\Pr(\max_j S_B^{(j)} < u\sqrt{B}) \leq e^{-u\sqrt{J}/4}$, at most $\frac{1}{2}$ beyond an absolute J . Hence $\mathbb{E} \max_j (S_B^{(j)})^+ \geq \frac{1}{2}u\sqrt{B} = \Omega(\sqrt{B \log J})$. Second, if $\log J > B/8$ and $B \geq 1024$, take $u = \sqrt{B}/32$, so $u\sqrt{B} = B/32$ and $e^{-64u^2} = e^{-B/16}$, and $J \geq e^{B/8}$ makes $J e^{-B/16}$ diverge, giving $\mathbb{E} \max_j (S_B^{(j)})^+ = \Omega(B)$, which is $\Omega(\sqrt{B \min\{1 + \log J, B\}})$ in that regime. Third, when $\log J < 128$ or $B < 1024$ the quantity $\sqrt{B \min\{1 + \log J, B\}}$ is $O(\sqrt{B})$, and a single walk already gives $\mathbb{E}(S_B^{(1)})^+ = \frac{1}{2}\mathbb{E}|S_B| \geq \frac{1}{2}(\mathbb{E}S_B^2)^{3/2}(\mathbb{E}S_B^4)^{-1/2} \geq \frac{1}{2}\sqrt{B/3}$ by Hölder, since $\mathbb{E}S_B^2 = B$ and $\mathbb{E}S_B^4 \leq 3B^2$. Combining the three regimes, $\mathbb{E} \max_j (S_B^{(j)})^+ = \Theta(\sqrt{B \min\{1 + \log J, B\}})$. This is the route by which maxima of independent random walks fix the minimax rate for prediction with expert advice (Cesa-Bianchi and Lugosi, 2006).

Proof idea. Divide time into blocks of length d and give J actions private states whose losses are perturbed by independent Rademacher signs within each block. Every observation available when an action is chosen comes from an earlier block, and the supplied $m_t = c_{t-d}$ depends only on earlier signs, so the learner cannot predict the signs governing its current block. Its expected loss is therefore $T/2$, while the best fixed action in hindsight benefits from the maximum of J independent random walks. Choosing the amplitude to make $E_2 \leq \mathcal{E}$ gives the stated scale.

Proof of Theorem 2. Take $h = d$ and $\varepsilon = \frac{1}{2}\sqrt{\mathcal{E}/T}$. By Lemma 14 every algorithm's expected loss equals $T/2$, since its action is independent of the current block's signs and all states have mean $\frac{1}{2}$ marginally. Its regret is therefore exactly the comparator's fluctuation advantage, $\varepsilon d \mathbb{E}[\max_{j \leq J} (\sum_{2 \leq b \leq B} \sigma_b^{(j)})^+]$ with $B = \lfloor T/d \rfloor$, over the $B - 1$ randomized blocks. Per round $\|c_t - m_t\|_\infty \leq 2\varepsilon$, because the supremum norm charges only the largest coordinate and each coordinate moves by at most 2ε at a block boundary; the gap is ε across the first randomized block, where m_t is still the neutral c_1 , and lies in $\{0, 2\varepsilon\}$ thereafter. Hence $E_2 \leq 4\varepsilon^2 T = \mathcal{E}$ holds deterministically. The budget must bind on the realized sequence, since a budget holding only in expectation would not constrain the instance selected by Yao's principle. By Lemma 15 and the discussion following it, the expected maximum of the positive parts of J independent Rademacher sums of length $B - 1$ is $\Theta(\sqrt{(B - 1) \min\{1 + \log J, B - 1\}})$, which covers both regimes and the case $J = 1$. The hypothesis $d \leq T/2$ gives $B \geq 2$ and hence $B - 1 \geq B/2$, so replacing $B - 1$ by B costs a factor at most $\sqrt{2}$ in each branch of the minimum, including the saturated branch where the minimum is $B - 1$ rather than $1 + \log J$. Substituting $B = \lfloor T/d \rfloor$ and ε therefore gives $\Theta(\sqrt{d\mathcal{E} \min\{1 + \log J, T/d\}})$. Rounds after the last complete block carry no drift and change the constant only. When the oracle is withheld, the $\sqrt{KT}$ term is Esposito et al. (2023, Thm. 5.1), whose construction is stationary and therefore lies in the present model class. $\square$

586 **Proposition 16.** Suppose $X_t = c_t(A_t)$. On instances whose drift is carried by a single action, $\Lambda_2 =$
587 E_2 , so the bound of Bar-On and Mansour (2025) reads $\tilde{O}(\sqrt{KT} + d E_2)$ and its drift contribution
588 $\sqrt{d E_2}$ matches Theorem 2 up to logarithmic factors.

589 In words, single-action drift makes the Euclidean and supremum measures agree, so on that family
590 the drift term of the published upper bound is unimprovable in order. Its $\sqrt{KT}$ term is not matched
591 because the oracle supplies an entire stale loss vector, so the oracle-assisted problem need not retain
592 bandit exploration complexity. Neither result provides an E_2 -adaptive upper bound under noisy
593 outcomes.

594 *Proof of Proposition 16.* If only one coordinate of $c_t - m_t$ is nonzero then $\|c_t - m_t\|_2 = \|c_t - m_t\|_\infty$.
595 On the family of Section 4 that coordinate has magnitude at most $2\varepsilon \leq 1$, so the truncation at one in
596 Λ_2 is inactive and $\Lambda_2 = E_2$. The bound of Bar-On and Mansour (2025) then reads $\tilde{O}(\sqrt{KT} + d E_2)$,
597 whose $\sqrt{d E_2}$ contribution Theorem 2 matches up to the logarithmic factor at $J = 1$. Its $\sqrt{KT}$
598 contribution is not matched. $\square$

599 H The two-action counterexample

600 Take $K = 2$ with the hybrid regularizer of Zimmert and Seldin (2020), Tsallis parameter $\beta =$
601 0.01 , learning rate $\eta = 0.5$ and residual scale $C = 1$. The iterate solves the stationarity condition
602 $g(x_1) + L_1 = g(x_2) + L_2$ with $x_1 + x_2 = 1$, where

$$g(x) = -\frac{1}{2\beta}x^{-1/2} + \frac{1}{\eta}\log x$$

603 and L_a is the cumulative estimate for action a . Suppose the cumulative estimates place action 1 at
604 probability x_1 . A single maximally negative centered residual $-C/x_1$ is then delivered to action 1.
605 This is realizable in the model of Section 3, since $m_t(1) = 1$ and $X_t = 0$ give exactly $\check{c}_t(1) -$
606 $m_t(1) = -C/x_1$ with $C = 1$. Solving the stationarity condition again gives the following.

	x_1 before	x_1 after	ratio
	9.15×10^{-3}	1.46×10^{-2}	1.6
	9.83×10^{-4}	7.39×10^{-3}	7.5
607	2.19×10^{-4}	9.98×10^{-1}	4561
	5.80×10^{-5}	1.00	17226
	4.86×10^{-6}	1.00	205641

608 The transition sits at $4\beta^2 C^2 = 4 \times 10^{-4}$, which is where the Tsallis term $x^{-1/2}$ ceases to dominate
609 the entropy term $\log x$ in g , so the one-step linearization underlying a constant-factor bound fails.
610 The violated statement is the local-stability inequality asserting that $x_{t'}(a)/x_t(a)$ stays bounded by
611 a constant for all t' within d rounds of t . A single round refutes it, so no accumulation argument is
612 needed. This invalidates that lemma and the proof route resting on it, and not every argument that
613 might use the same regularizer. It does not refute the open problem of Section 6, which concerns
614 the regret rate and may be provable by other means. Devices that bound the residual magnitude
615 introduce either a first-order bias dominating the target term or a probability floor narrowing the
616 delay regime.

617 I Numerical detail

618 Sixteen studies were run. The six that bear on the claim are reported at length below, each with
619 an observation, an interpretation and a limitation; the other ten are summarised in one table at the
620 end. Studies one to thirteen draw P from a Dirichlet and fourteen to sixteen do not. STATE-EXP3
621 is beaten in study fifteen, and in four of six cells of study sixteen when the baseline alone is grid-
622 tuned. Study 10 is a control for study 6 and reverses its conclusion. All policies run on the same
623 environment parameters and common random-number streams, so a seed fixes everything except
624 which estimator is formed, and differences are taken within a seed before averaging. Round r 's
625 outcome is consumed before round $r + d + 1$, matching Section 3. Reported $\pm$ is one standard error
626 of the paired difference. Throughout this appendix, *sample-path pseudo-regret* means $\sum_t c_t(A_t) -$

627 $\min_a \sum_t c_t(a)$, evaluated on one realization of the learner and environment randomness. Averaging
628 it over seeds estimates the expected pseudo-regret that the theorems bound. We use that name rather
629 than “realized regret” to distinguish it from regret computed from the sampled outcomes X_t , which
630 is not what these studies report.

631 **Does the bound hold, and does the estimator help?** With $|\mathcal{S}| = 4$, $T = 10000$, $d = 50$, drift band
632 0.05 and 100 seeds:

	K	state, $m = d + 1$	state, $m = 1$	action, $m = 1$	paired gain	state better in
633	8	455.1	368.2	387.6	19.3 ± 3.4	73/100
	20	618.3	493.7	547.5	53.8 ± 5.6	87/100
	40	734.3	574.5	696.1	121.5 ± 9.4	99/100
	80	780.4	614.8	784.0	169.3 ± 9.9	100/100

634 Observed. Sample-path pseudo-regret for the analyzed variant stayed below $\sqrt{2(d+1)|\mathcal{S}|T \log K}$
635 in 100 of 100 seeds at every K , the bound running from 2913 to 4228 against realized 455 to 780.
636 The plain variant beat the action-level baseline at every K , by a margin growing from 19 to 169 as
637 K rose tenfold with $|\mathcal{S}|$ fixed, and it beat the analyzed variant throughout. Sweeping the delay at
638 $K = 40$ gave state-pooled 387, 569 and 731 against action-level 653, 699 and 773 at $d = 10, 50$
639 and 200.

640 Interpretation. The margin growing in K at fixed $|\mathcal{S}|$ is the prediction that separates the two estima-
641 tors, so the improvement is attributable to the state channel the effective dimension attributes to it.
642 Interleaving looks like a price of the analysis rather than of the estimator. The narrowing of the gap
643 as the delay grows is what an additive delay term predicts and a multiplicative one does not, which
644 is the diagnostic evidence behind the open problem of Section 6.

645 Limitation. This is one synthetic family with a fixed drift band, the plain variant carries no guarantee,
646 and a simulation cannot establish that a bound holds, only that sample-path pseudo-regret stayed
647 below it in the runs performed.

648 **Does the effective dimension, rather than $|\mathcal{S}|$, govern the gain?** Holding $|\mathcal{S}| = 8$, $K = 40$,
649 $T = 5000$, $d = 20$ and 24 seeds, and varying only the overlap between the rows of P , the measured
650 mean of v_t moves while $|\mathcal{S}|$ does not:

	row overlap	mean v_t	state-pooled	action-level	paired gain
651	0.00	2.719	392.7	599.3	206.6 ± 20.0
	0.25	2.013	361.1	528.1	167.0 ± 21.3
	0.50	1.532	306.5	406.5	100.0 ± 10.5
	0.75	1.168	199.9	229.2	29.3 ± 3.2
	0.95	1.009	49.2	50.8	1.6 ± 0.5

652 Observed. Both the regret and the advantage over action-level weighting fall monotonically with v_t ,
653 from a paired gain of 206.6 at $v_t = 2.719$ to 1.6 at $v_t = 1.009$. Nothing about the raw state count
654 changes across the sweep.

655 Interpretation. $|\mathcal{S}|$ is constant and cannot explain the pattern, while v_t moves with it, which is the
656 comparison Theorem 5 predicts.

657 Limitation. The two estimators do not become the same object at $v_t \approx 1$. Rather the rows of P agree,
658 so every action induces the same loss and both regrets collapse, making that endpoint uninformative
659 rather than favorable. One parameter of one map family was varied, so the sweep is consistent with
660 the effective dimension governing the gain rather than establishing it.

661 **Does the drift bound describe the right scaling?** The construction of Theorem 2 at $K = 40$,
662 $T = 8000$, $h = d$ and 12 paired seeds, sweeping d from 10 to 200, the amplitude ε from 0.02 to
663 0.20 so that E_2 runs from 12 to 1193, and the drifting coordinates J from 1 to 16. We run four
664 policies: uniform, action-level EXP3, STATE-EXP3, and the greedy policy on the exact stale vector
665 m_t , the last because Theorem 2 claims to survive it. E_2 is measured on the realized instance rather
666 than predicted, so the normalization uses the quantity the theorem uses.

667 Observed. The predicted scale $\sqrt{d E_2 \min\{1 + \log J, T/d\}}$ ranges from 38 to 377 across the grid
668 and the measured regret tracks it (Figure 3). Normalized, the four policies read 0.230 to 0.429,

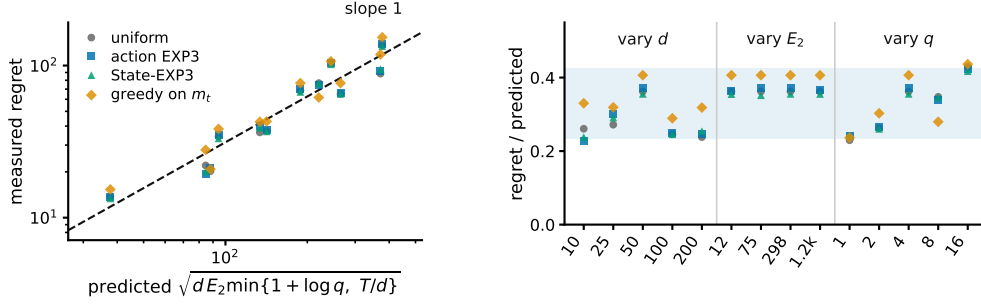


Figure 3: Study 8, the drift channel. **Left**, measured regret against the scale Theorem 2 predicts, over twelve cells in which d varies 20-fold, E_2 100-fold and J 16-fold. The dashed line has slope one. **Right**, the same points divided by that scale, grouped by which parameter the cell sweeps, with the shaded band spanning the observed range. STATE-EXP3 appears here only as one more algorithm on the hard family, where it has no advantage.

669 0.228 to 0.424, 0.238 to 0.417 and 0.237 to 0.437, a spread below 1.9 against a tenfold range in
670 the predictor. The oracle-assisted policy has the largest mean ratio, 0.345 against 0.311, 0.314 and
671 0.310 for the three that never see m_t .

672 Interpretation. The three arguments enter the predicted scale differently and the collapse holds in all
673 three sweeps, so it is the form of the bound rather than one fitted constant that tracks the data. That
674 greedy play on the exact m_t sits at the top of the band rather than below it is the concrete content of
675 the claim that the bound survives the oracle.

676 Limitation. Every cell has $T/d > 1 + \log J$, so only the $1 + \log J$ branch of the minimum is
677 exercised and the T/d saturation branch, which needs d comparable to T , is not probed. A lower
678 bound quantifies over all algorithms and four policies cannot establish that. What these runs check
679 is the scaling of the construction.

680 **The remaining studies in brief.** Ten further studies are summarised here rather than reported at
681 length, since none of them bears directly on the claim of Section 5. Each was run under the same
682 conventions, and the numbers below are what the runs produced.

Question	What the runs showed
What does an unknown map cost?	The online plug-in stays within 0.5 per cent of known P at every K and beats action-level weighting by 20 to 179 as K rises tenfold; forced exploration never helped, $\gamma = 0$ being best at every K .
Where does the plug-in fail?	On five deliberately hostile maps it still beat action-level weighting, two of them with κ above 100, and forced exploration cost regret in every cell. The Laplace default is not the best smoothing.
Does safe blending recover the pooling gain?	No. The rule reproduces action-level regret to the digit, so its guarantee is safe and the algorithm is inert.
Can the denominator be counted rather than predicted?	Counting is the noisier of the two on regret and on bias, and the predicted relative error is vacuous from $K = 40$ through $K = 4000$.
Is there a bias-variance trade in the representation?	Yes. Regret is lowest at four groups, the number of latent clusters, and stays there under heterogeneity 0.12.
Does the rotating bound describe the right dependence on T , d and $\hat{v}$?	Partly. The ratio $R/\sqrt{(d+1)\hat{v}T \log K}$ spans 35.9 times over the grid, 7.1 times once a near-degenerate column is dropped, so the bound tracks the shape and not the constant.
Is the plug-in fragile in the smoothing constant?	No cell of thirty loses to action-level weighting, and the gain roughly doubles as α falls from 1 to 0.03.
How far is the method from its certificate?	The plug-in stays within 1.6 per cent of known P , while the certificates themselves are far looser than the measured arms.
Can the coarsening be chosen from data?	No. The rule selects 15 groups at heterogeneity 0 and 16 at 0.12 against a true count of four. Learning g remains open.
Which assumption produces which benefit?	A best-state rule leads seven of the nine settings and STATE-EXP3 two, but where the rule's assumption fails its regret grows linearly and its spread across seeds is four times ours.

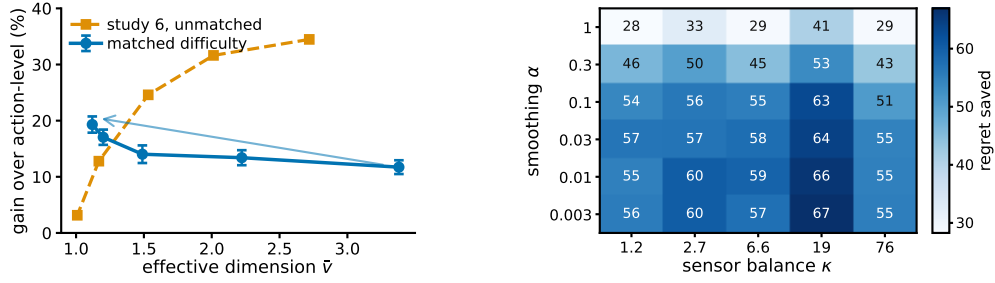


Figure 4: **Left**, study 10. The same overlap sweep with and without holding the task difficulty fixed. Matching the uniform-play regret reverses the trend, supporting the interpretation that the decline in Study 6 is driven by task-difficulty confounding rather than by the pooling mechanism itself. Bars are one standard error of the paired difference. **Right**, regret saved against action-level weighting over the smoothing pseudocount and the sensor balance, 10 paired seeds per cell, from the plug-in sensitivity study summarised in the closing table.

Does pooling help because the rows overlap, or because overlap made the task easy? Study 6 sweeps the overlap and finds the advantage falling as $\hat{v}$ falls, which read at face value contradicts Theorem 5, whose bound improves as $\bar{v}$ falls. That sweep is confounded, since raising the overlap also flattens $c_t = P\theta_t$ across actions, collapsing the comparator gap so that both regrets vanish together. This study holds $K = 40$, $|\mathcal{S}| = 8$, $T = 10000$, $d = 50$ and the difficulty fixed, and varies only $\hat{v}$. Difficulty is the uniform-play regret $\sum_t \text{mean}_a c_t(a) - \min_a \sum_t c_t(a)$, a property of the instance with no algorithm in it, and the amplitude of θ is bisected in every cell to hit a common target. The map keeps one contrast direction alive at every overlap and θ is aligned with it, which is what lets the gap survive as the rows agree.

Observed. Across 20 paired seeds the relative gain rises from 11.7 per cent at $\hat{v} = 3.380$ to 19.3 per cent at $\hat{v} = 1.120$, with absolute gains 85.5 ± 8.9 to 106.2 ± 7.9 and a correlation of -0.929 between $\hat{v}$ and the gain. Study 6’s unmatched sweep falls from 34.5 to 3.1 per cent over a comparable range (Figure 4).

Interpretation. The two sweeps disagree in sign, and the matched-difficulty sweep agrees with the overlap dependence suggested by Theorem 5, although this experiment uses the practical $m = 1$ variant rather than the theorem’s $m = d + 1$ algorithm. Overlap is what pooling exploits, and study 6 measures overlap confounded with an instance becoming trivial. This also explains study 9’s degenerate column without appealing to it.

Limitation. At the two highest overlaps some seeds cannot reach the common target and the realized difficulty is 5 to 13 per cent below it, so the matching is close rather than exact. Since a smaller gap should if anything shrink the advantage, the measured rise is conservative in that direction. The matched gains are also smaller in magnitude than the unmatched ones, which is what fixing the difficulty costs.

Does the matched-difficulty trend survive a second configuration? Study 10 is a single design point, so the same procedure was re-run at $K = 20$, $|\mathcal{S}| = 12$, $T = 8000$, $d = 20$ and 20 paired seeds, with the common target lowered to 500 so that it is reachable at every overlap. Nothing else changes.

overlap	$\hat{v}$	R_{unif}	state	action	paired gain
0.00	3.440	500.0	364.7	407.2	42.5 ± 8.2
0.35	2.225	500.0	346.8	401.3	54.6 ± 8.3
0.65	1.518	500.0	309.3	372.4	63.2 ± 6.4
0.85	1.240	493.6	271.8	346.0	74.3 ± 7.2
0.95	1.191	490.1	256.9	327.6	70.7 ± 7.8

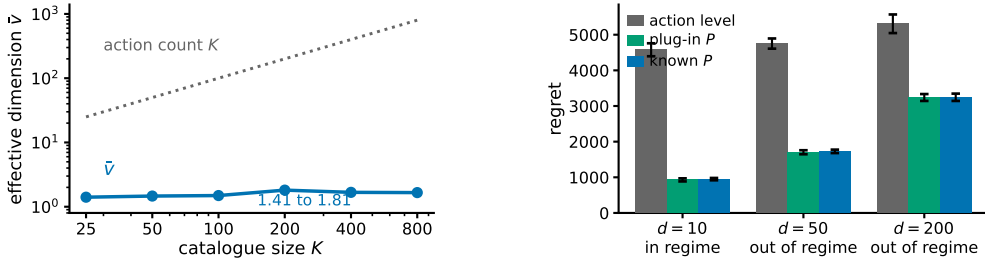


Figure 5: Study 14, a recommendation funnel. **Left**, the effective dimension stays near two while the catalogue grows thirty-two-fold, so the gap between K and $\hat{v}$ comes from the catalogue having categories, not from a chosen K . **Right**, the three arms at $T = 20000$ and 8 paired seeds, labelled by whether the regime condition of Section 4 holds. Bars are one standard error. No real data is used anywhere in this study.

Observed. The relative gain rises from 10.4 to 21.6 per cent as $\hat{v}$ falls from 3.440 to 1.191, and the correlation between $\hat{v}$ and the absolute gain is -0.972 against study 10’s -0.929 . Realized difficulty spans 490.1 to 500.0, a 1.020 times spread, so the matching is tighter here than in study 10.

Interpretation. The direction study 10 reports is not an artifact of its single configuration. Two independent design points, differing in K , $|\mathcal{S}|$, T and d , give the same sign and a comparable magnitude.

Limitation. Both configurations use the same construction and the same matching routine, so this tests configuration sensitivity and not construction sensitivity. Both run the practical $m = 1$ variant. An earlier attempt at this replication silently failed because `match_scale` binds its target as a default argument, so the bisection saturated and the difficulty was never matched; the run reported here passes the target explicitly.

Is “many actions resolved by few states” a natural regime? Every study so far draws P from a Dirichlet, which is the assumption under test rather than evidence for it. This one builds an environment to the shape of a recommendation funnel instead. Actions are $K = 200$ catalogue items, states are $|\mathcal{S}| = 6$ ordered engagement depths, and the outcome is a delayed conversion. Three structural facts are imposed because funnels have them, namely that items fall into a modest number of categories sharing an engagement profile, that item quality is heavy-tailed, and that the loss falls with engagement depth and drifts seasonally rather than adversarially. The generated catalogue puts 74 per cent of its mass on no engagement and 1.4 per cent on the deepest.

Observed. The effective dimension is $\hat{v} = 1.97$ against $K = 200$, a hundredfold gap, and it stays between 1.41 and 1.81 as the catalogue grows from 25 to 800 items (Figure 5). Against action-level weighting the state-pooled variant saves 79.3, 63.6 and 38.9 per cent at $d = 10, 50$ and 200, the largest margins anywhere in this appendix. The plug-in matches known P at every delay, the paired difference never reaching two standard errors. The regime condition $\bar{v}(d+1) \log K \leq K + d$ holds at $d = 10$ and fails at $d = 50$ and $d = 200$.

Interpretation. The separation between K and $\hat{v}$ arrives from catalogue structure rather than from a tuned parameter, and it does not close as the catalogue grows, which is the claim the Dirichlet families cannot make. That the condition Theorem 5 needs fails at the delays a funnel has, while the practical variant still saves most of the regret, is the same message as study 1 in the setting that motivates the paper. It is evidence for the open problem of Section 6 rather than for the theorem.

Limitation. *No real data is used.* There is no interaction log in this environment, nothing here is fitted to one, and this is not a semi-synthetic benchmark in the usual sense of that phrase. The structure is imposed by hand from qualitative properties, so it shows that the regime is describable and self-consistent, not that any deployed system occupies it. The arms run $m = 1$, which carries no guarantee, and the environment has no adversarial component, so it exercises the pooling channel and not the drift channel.

748 **Does pooling still help against the strongest published delayed-bandit algorithm?** Every com-
749 parison so far is against action-level weighting, EXP3-IX, naive sharing, or the best-state rule.
750 None is the strongest published method for delayed feedback: that is the hybrid Tsallis-plus-entropy
751 FTRL of Zimmert and Seldin (2020), whose delay term is additive rather than multiplicative. We
752 run it on both headline families, with the same environments, seeds and delay convention, under
753 three tunings.

family, cell	STATE-EXP3	Zimmert and Seldin	same, grid-best	STATE-EXP3, grid-best
funnel, $d = 10$	949.0	2959.4	2413.8	27.0
funnel, $d = 50$	1728.1	3557.0	2556.4	95.9
754 funnel, $d = 200$	3244.8	4734.7	2807.7	199.7
overlap 0.00	661.5	731.8	610.4	202.4
overlap 0.65	584.1	673.1	544.0	51.2
overlap 0.95	461.4	540.2	419.4	34.3

755 Observed. With STATE-EXP3 at the single tuning used everywhere else here and Zimmert and
756 Seldin (2020) at the best of a six-point grid around its own value, STATE-EXP3 wins every cell,
757 cutting regret by 31.5 to 67.9 per cent on the funnel and 9.6 to 14.6 per cent on the matched sweep.
758 With both at their grid-best, over seven multipliers for Zimmert and Seldin (2020) and six for STATE-
759 EXP3, STATE-EXP3 wins every cell again and by more, 66.8 to 98.9 per cent. With Zimmert and
760 Seldin (2020) at its grid-best and STATE-EXP3 left at its paper tuning, the third column against the
761 first, Zimmert and Seldin (2020) wins the funnel at $d = 200$ and all three matched cells, four of six.

762 Interpretation. No column is a symmetric protocol. The two grids are finite and unequal, and
763 STATE-EXP3's grid-best sits at its top multiplier in every cell, so its own minimum may lie outside
764 the range searched; the truncation runs against STATE-EXP3. Where it is not held to the worse
765 tuning of the two, the state channel helps against the strongest published method, which is expected,
766 since that method is action-level by construction and cannot read the state at all. Grid tuning helps
767 STATE-EXP3 far more than it helps Zimmert and Seldin (2020), which is consistent with pooling
768 cutting estimator variance and lower variance admitting a larger learning rate.

769 Limitation. The third column against the first is a real loss and is reported as one. These are our
770 own synthetic families, and a method built for adversarial worst cases is not shown at its best on
771 benign instances. Every grid-best multiplier is chosen by reading these benchmarks, so none of the
772 last three columns is a procedure a practitioner could run; only the first is. Eight paired seeds on the
773 funnel and ten on the matched sweep.

774 Taken together the sixteen studies cover both channels. Studies one, six, seven, nine, ten, fourteen
775 and fifteen probe the state channel. They are consistent with the effective dimension, rather than
776 $|\mathcal{S}|$ or K , governing what pooling buys, with the qualification that study ten had to control for task
777 difficulty before study six's sweep pointed the way Theorem 5 does, and study nine shows where
778 the theorem stops describing the dependence and attributes that to the interleaving. Studies two to
779 five, eleven and twelve probe what happens when the action-to-state matrix must be estimated, and
780 locate the difficulty in certifying the estimate rather than in forming it, study eleven adding that the
781 smoothing constant matters more than the coverage imbalance the analysis carries and study twelve
782 putting the measured arms and their certificates on one axis. Study thirteen closes the representation
783 thread by showing that the bound of Theorem 6 is too conservative to serve as its own selection
784 rule. Study fourteen is the only one whose action-to-state matrix is not drawn from a Dirichlet, and
785 it reports the largest margins here, though on a structure imposed by hand rather than measured.
786 Study fifteen widens the baseline set and finds STATE-EXP3 beaten in seven of nine settings by a
787 heuristic that assumes the best action is the one reaching the single best state, which holds on several
788 of these families and fails with linear regret when it does not. Study sixteen places STATE-EXP3
789 against the strongest published delayed-bandit method. It is ahead in every cell when both are tuned
790 alike, and behind in four of the six when that method alone is grid-tuned. Study eight is the drift
791 channel, and its collapse across d , E_2 and J is the strongest evidence here that Theorem 2 has the
792 right form. No study runs a method that uses both channels, which is the gap Section 6 describes.

793 J A route that does not work

794 Theorem 1 is proved in Appendix K by bounding a state-marginal ratio. The natural first attempt is
795 different and fails, which is worth recording so the detour is not repeated.

796 Run STATE-EXP3 with $m = 1$ and let $\tilde{x}_r$ be the iterate that has also absorbed the outstanding
797 estimates $G_r = \sum_{s=r-d}^{r-1} \hat{c}_s \geq 0$. The delay term $\mathbb{E}[\langle x_r - \tilde{x}_r, \hat{c}_r \rangle]$ is at most $\eta \sum_{s=r-d}^{r-1} \mathbb{E}[\langle x_r, \hat{c}_s \rangle] \leq$
798 ηd for any unbiased estimator, which is already the additive form. The quadratic term is where it
799 stops: the telescoping scores $\sum_a x_{r+d}(a) \hat{c}_r(a)^2$ while Proposition 4 bounds the sum against x_r ,
800 and the substitution costs $|\mathcal{S}|/Z_r$ with $Z_r = \mathbb{E}_{a \sim x_r}[e^{-\eta G_r(a)}]$, hence an exponential moment of
801 $\langle x_r, G_r \rangle$.

802 That moment has no bound uniform over state distributions. Here $\langle x_r, \hat{c}_s \rangle = q_r(S_s)X_s/q_s(S_s)$,
803 governed by the reciprocal probability of the state actually observed, whose mean is $|\mathcal{S}|$ but
804 whose exponential moment is not uniformly finite: already at $|\mathcal{S}| = 2$ with $q = (\varepsilon, 1 - \varepsilon)$,
805 $\mathbb{E}[e^{\eta/q(S)}] \geq \varepsilon e^{\eta/\varepsilon} \rightarrow \infty$ while the mean stays at 2. Flooring q restores the moment and costs
806 the product back, since $\eta \langle x_r, G_r \rangle \leq 1$ then needs $\gamma \asymp \eta d \kappa |\mathcal{S}|$ and the mixing cost γT reproduces
807 $\tilde{O}(\sqrt{d \kappa} |\mathcal{S}| T \log K)$. Empirically the obstruction never bites, $1/Z_r$ staying below 1.33 at the tuned
808 η over $d \in \{10, 50, 200\}$, which is what suggested it was an artifact of the device rather than of the
809 problem.

810 K Proof of Theorem 1: additive delay without interleaving

811 Interleaving cannot give an additive bound. Copy i , which absorbs a share V_i of the budget $V =$
812 $\sum_i V_i$, has regret $\sqrt{2V_i \log K}$, and Cauchy–Schwarz over the m copies gives $\sqrt{2mV \log K}$, tight
813 when the V_i agree, so the $\sqrt{m}$ in Theorem 5 is not an artifact of the analysis. An additive bound has
814 to come from a single copy.

815 Throughout, $L_t = \sum_{r \leq t-d-1} \hat{c}_r$ and $x_t \propto e^{-\eta L_t}$, so $L_{t+1} - L_t = \hat{c}_{t-d}$, with $\hat{c}_r = 0$ for $r < 1$. The
816 filtration $\mathcal{F}_t$ is the environment’s, including outcomes not yet delivered, as in Proposition 4; under it
817 x_t is $\mathcal{F}_{t-d-1}$ -measurable, so x_{r+d} is $\mathcal{F}_{r-1}$ -measurable.

818 **The drift lemma.** Appendix J priced the change of measure from x_r to x_{r+d} through $1/Z_r$. That
819 is avoidable. The pooled estimate satisfies

$$\langle x_t, \hat{c}_t \rangle = \sum_a x_t(a) \frac{P(S_t | a) X_t}{q_t(S_t)} = X_t \leq 1$$

820 identically, on every path, because the denominator is the state marginal of the same distribution that
821 supplies the numerator.

822 **Lemma 17.** *If $\eta \leq 1/(e(d+1))$ then $x_{t+1}(a) \leq (1 + 1/d)x_t(a)$ for every t and every a , surely,*
823 *and hence $x_{t+d}(a) \leq e x_t(a)$.*

824 *Proof.* Strong induction on t . For $t \leq d$ nothing is delivered and $x_{t+1} = x_t$. For $t > d$,
825 assume the claim for every $s < t$ and compose it over the d steps from $t-d$ to t , giving
826 $x_t(a) \leq (1 + 1/d)^d x_{t-d}(a) \leq e x_{t-d}(a)$. Contracting with the fixed nonnegative vector $\hat{c}_{t-d}$ and
827 using the identity above, $\langle x_t, \hat{c}_{t-d} \rangle \leq e \langle x_{t-d}, \hat{c}_{t-d} \rangle = e X_{t-d} \leq e$. With $Z_t = \sum_b x_t(b) e^{-\eta \hat{c}_{t-d}(b)}$
828 and $e^{-z} \geq 1 - z$, this gives $Z_t \geq 1 - \eta \langle x_t, \hat{c}_{t-d} \rangle \geq 1 - \eta e \geq d/(d+1)$, so $x_{t+1}(a) =$
829 $x_t(a) e^{-\eta \hat{c}_{t-d}(a)} / Z_t \leq x_t(a) / Z_t \leq (1 + 1/d)x_t(a)$. $\square$

830 What the induction bounds is $\langle x_t, \hat{c}_{t-d} \rangle = X_{t-d} q_t(S_{t-d}) / q_{t-d}(S_{t-d})$, the state-marginal ratio
831 Appendix J could not reach. Individual $\hat{c}_r(a)$ remain unbounded; only their x_t -weighted mass is
832 controlled. Cesa-Bianchi et al. (2019) prove the action-level version under $\eta \leq 1/(Ke(d+1))$ and
833 Thune et al. (2019) remove the K , using that their estimate is carried by the single action played, so
834 the weight cancels into a same-action ratio. A pooled estimate is spread over every action and does
835 not collapse that way; the identity above replaces the cancellation.

Centring. Let $\zeta_t = \hat{c}_t - X_t \mathbf{1}$, with $\zeta_r = 0$ for $r < 1$. Exponential weights is unchanged by a shift that is constant across actions, so x_t is the same iterate and this is a change of analysis, not of algorithm (Eldowa et al., 2024). Keeping $\gamma(s) = \mathbb{E}[X_t^2 \mid S_t = s]$ inside the state sum,

$$\mathbb{E}[\langle x_t, \zeta_t^2 \rangle \mid \mathcal{F}_{t-1}] = \sum_s \gamma(s) \left[\frac{\sum_a x_t(a) P(s \mid a)^2}{q_t(s)} - q_t(s) \right] \leq v_t - 1,$$

each bracket being nonnegative by Jensen and $\gamma(s) \leq 1$, the final step being Remark 3. Bounding $\langle x_t, \hat{c}_t^2 \rangle$ by v_t first and then subtracting does not work: it leaves $v_t - \mathbb{E}[X_t^2]$.

Proof of Theorem 1. Since $\zeta_r(a) \geq -X_r \geq -1$ we have $\eta \zeta_r(a) \geq -\eta$, and Lagrange’s remainder gives $e^{-z} \leq 1 - z + (e^\eta/2)z^2$ for $z \geq -\eta$. With $\log(1 + y) \leq y$, the potential $W_t = \sum_a \exp(-\eta \sum_{r \leq t-d-1} \zeta_r(a))$ obeys $\log(W_{t+1}/W_t) \leq -\eta \langle x_t, \zeta_{t-d} \rangle + (\eta^2 e^\eta/2) \langle x_t, \zeta_{t-d}^2 \rangle$. Telescoping from $W_1 = K$ and using $\log W_{T+1} \geq -\eta \sum_{r \leq T-d} \zeta_r(a^*)$ for the comparator a^* ,

$$\sum_{r=1}^{T-d} \langle x_{r+d}, \zeta_r \rangle \leq \sum_{r=1}^{T-d} \zeta_r(a^*) + \frac{\log K}{\eta} + \frac{\eta e^\eta}{2} \sum_{r=1}^{T-d} \langle x_{r+d}, \zeta_r^2 \rangle.$$

The shift cancels between the two sides. As x_r is $\mathcal{F}_{r-d-1}$ -measurable, $\mathbb{E}[\langle x_r, \hat{c}_r \rangle] = \mathbb{E}[\langle x_r, c_r \rangle]$ and $\mathbb{E}[\hat{c}_r(a^*)] = c_r(a^*)$, and the d rounds undelivered by $T + 1$ contribute at most d , so

$$\mathbb{E}[R_T] \leq \frac{\log K}{\eta} + \frac{\eta e^\eta}{2} \sum_r \mathbb{E}[\langle x_{r+d}, \zeta_r^2 \rangle] + \sum_r \mathbb{E}[\langle x_r - x_{r+d}, \hat{c}_r \rangle] + d.$$

Lemma 17 gives $x_{r+d} \leq e x_r$ surely, and x_{r+d} is $\mathcal{F}_{r-1}$ -measurable, so the second sum is at most $e \sum_r \mathbb{E}[v_r - 1] \leq e V^-$.

For the third, write $x_{r+d}(a) = x_r(a) e^{-\eta G_r(a)} / Z'$ with $G_r = \sum_{u=r-d}^{r-1} \hat{c}_u \geq 0$ and normaliser $Z' \leq 1$, distinct from the one-step Z_t of the Lemma. Then $x_r(a) - x_{r+d}(a) \leq x_r(a)(1 - e^{-\eta G_r(a)}) \leq \eta x_r(a) G_r(a)$ for every a , trivially where the left side is negative. Since $\hat{c}_r \geq 0$ and x_r, G_r are $\mathcal{F}_{r-1}$ -measurable, the conditional expectation is at most $\eta \langle x_r, G_r \rangle$, and $\mathbb{E}[\langle x_r, G_r \rangle] = \sum_{u=r-d}^{r-1} \mathbb{E}[\langle x_r, \hat{c}_u \rangle] \leq d$: for each such u , x_r is $\mathcal{F}_{r-d-1}$ -measurable and $r - d - 1 \leq u - 1$, so it leaves the conditional expectation and $\mathbb{E}[\langle x_r, \hat{c}_u \rangle] = \langle x_r, c_u \rangle \leq 1$. This is measurability, not independence, and at $u = r - d$ the two σ -fields coincide, so the step rests on the timing convention of Section 3. The sum is at most $\eta d T$.

Finally $\eta \leq 1/(2e)$ gives $e^{1+\eta} < 3.3$, so $\mathbb{E}[R_T] \leq \log K/\eta + \frac{\eta}{2}(3.3V^- + 2dT) + d$. For $f(\eta) = A/\eta + B\eta$ restricted to $\eta \leq \eta_{\max}$ one has $\min f \leq 2\sqrt{AB} + A/\eta_{\max}$, in the interior and in the binding case alike; with $A = \log K$, $B = (3.3V^- + 2dT)/2$ and $\eta_{\max} = 1/(e(d+1))$ this is the stated bound. $\square$

The cap on η binds only when $d \gtrsim T/(e^2 \log K)$, and not at the settings of Appendix I: the funnel study’s own rates are 0.0058, 0.0031 and 0.0016 at $d = 10, 50, 200$ against caps 0.0334, 0.0072 and 0.0018, so Theorem 1 covers those runs as they were executed. Over three seeds and 20000 rounds the induction has room to spare, $\max_a x_{t+1}(a)/x_t(a)$ reaching 1.006, 1.003 and 1.002 against the permitted 1.100, 1.020 and 1.005.

The two bounds are numerically close on the funnel, where $\bar{v} \approx 2$: the additive form carries $3.3(\bar{v} - 1) + 2d$ and the interleaved one $(d+1)\bar{v}$, and these cross near $\bar{v} = 2$. The additive form gains as $\bar{v}$ rises toward $|S|$, and it describes the algorithm that is actually run, which is the reason to prefer it.

L Stationary sensors

Observation 18. The multiplicative term $\sqrt{\min\{|S|, d\}T}$ is not forced by the known construction, and is absent at one endpoint of the present model class.

The $\Omega(\sqrt{\min\{|S|, d\}T})$ construction of Esposito et al. (2023, Thm. 5.2) re-splits states across actions adversarially per block, so its action-to-state map is time-varying and lies outside the model class of Section 3. At $E_2 = 0$ with stationary P the problem is a stochastic delayed bandit, where

875 delay enters additively (Joulani et al., 2013). One route to it is therefore unavailable and the term is
876 absent at one endpoint, but neither a matching upper bound for all $E_2 > 0$ nor a lower bound ruling
877 it out is established here.