
Pooling and Drift in Delayed Bandits

Melika Baghi

Georgia Institute of Technology
mbaghi3@gatech.edu

Abstract

1 An action’s consequences often arrive in two stages: an intermediate signal at
2 once, and the outcome that matters only after a delay, as a recommender sees
3 a click in seconds but a purchase in days. Over T rounds with K actions
4 and delay d , the rate available when the state-to-outcome map is unrestricted is
5 $\tilde{O}(\sqrt{(K+d)T})$ (Esposito et al., 2023), with $\tilde{O}$ hiding logarithmic factors, so
6 costs grow with the catalogue even when many items behave alike. When can
7 the intermediate signal remove that dependence? If the outcome mean depends
8 on the action only through the state it produced, one delayed outcome informs
9 every action that could have produced it, and the cost is set not by how many
10 actions there are but by how differently they behave, an *effective dimension* v_t
11 between 1 and the number of states $|S|$. Running $d+1$ copies of exponential
12 weights in rotation attains $\tilde{O}(\sqrt{(d+1)V \log K})$ for any $V \geq \sum_t v_t$ fixed in ad-
13 vance; a single copy makes the delay additive instead. Grouping similar states
14 lowers the cost further, at a bias we make explicit. Sharing cannot remove the
15 cost of delay itself: a learner handed the exact losses from d rounds ago still
16 suffers $\Omega(\sqrt{d\mathcal{E} \min\{1 + \log J, T/d\}})$, where J counts the drifting directions and
17 $\mathcal{E}$ bounds how far the losses drift while it waits. Intermediate feedback there-
18 fore helps only when the signal is both shared across actions and still current.
19 Synthetic experiments on the single-copy variant cut regret by up to 79 per cent
20 against action-level weighting, and by 32 to 68 per cent against the tuned minimax-
21 optimal method of Zimmert and Seldin (2020).

1 Introduction

23 Delayed outcomes are common in sequential decision-making: a system must act now even though
24 the outcome it cares about may arrive much later. In recommendation, an item is shown immediately,
25 an engagement signal such as a click or browsing depth is observed soon after, and conversion or
26 retention may only be known days or weeks later, so decisions are committed before their economi-
27 cally relevant consequences appear. The central question is when such a signal can reduce the cost
28 of learning before the final outcome arrives.

29 Here actions are catalogue items, states are intermediate engagement outcomes, and the delayed
30 outcome is conversion. Different items may induce similar distributions over engagement states,
31 so an outcome observed after one item can carry information about many others. The difficulty
32 scales not with the catalogue size K but with how differently the actions distribute users across the
33 intermediate states, so thousands of items funnelling users through a few engagement patterns stay
34 easy to learn from.

35 Bandits with intermediate observations formalize this: the learner repeatedly picks one action and
36 sees the consequence of that action alone. Esposito et al. (2023) show that the *state-to-loss* map
37 decides the complexity, an unrestricted one giving only the rate available with no intermediate ob-
38 servations; we fix that map and ask how much the *action-to-state* structure helps.




Figure 1: **Both panels synthetic.** **Left, the state-to-outcome relationship changes while the action-to-state channel stays fixed**, so the loss vector from d rounds earlier grows increasingly inaccurate. The drift measure E_2 of Section 3 tracks it. **Right, more overlap allows more sharing across actions**, at matched task difficulty, $K = 40$, $|\mathcal{S}| = 8$, $T = 10000$, $d = 50$, 20 paired seeds, bars one standard error of the paired difference. **The right panel runs the practical single-copy variant, which Theorem 1 does not cover and Theorem 2 does** (Appendix I, study 10).

39 Intermediate observations face two independent limitations: whether information can be shared
40 *across actions*, and whether it stays useful *across time*.

41 Across actions: because similar actions reach similar states, we can *pool through the state*, charg-
42 ing a delayed outcome to every action that could have produced the state observed. The resulting
43 statistical cost is governed by an *effective dimension* v_t , which measures how distinguishable the
44 action-induced state distributions are under the learner’s current play. When many actions behave
45 similarly through the state channel, v_t can be far smaller than K .

46 Across time: delay creates a second difficulty even when the learner is handed an exact description
47 of the past. Let $m_t = c_{t-d}$ be the action-loss vector from d rounds earlier. Its usefulness depends on
48 how far the losses have moved, which we measure by $E_2 = \sum_t \|c_t - m_t\|_\infty^2$ (Figure 1). Our lower
49 bound asks how much regret this mismatch forces even when m_t is supplied for free.

50 In short: with P known, running $d+1$ copies lets the states cut the cost of learning (Theorem 1), a
51 single copy makes the delay additive (Theorem 2), and a grouping fixed in advance trades dimension
52 against bias (Theorem 3); but no algorithm escapes the cost of delay, even one handed the stale losses
53 outright (Theorem 4). The first three address the difficulty across actions and the last the one across
54 time. An unknown action-to-state matrix can be estimated, with a weaker guarantee (Appendix E),
55 while jointly adapting to overlap and temporal drift remains open (Conjecture 5).

56 2 Related work

57 Our results connect two largely separate lines of work: delayed feedback, and information sharing
58 across actions. *Delayed bandits.* Vernade et al. (2020) vary the action-to-state map while fixing
59 the state-to-loss map; we do the reverse, letting the latter vary obliviously. Zimmert and Seldin
60 (2020) give the optimal $\tilde{O}(\sqrt{KT} + \sqrt{D \log K})$ for total delay D , and Esposito et al. (2023) allow
61 an unknown and possibly adversarial action-to-state map; we fix that channel instead, which is what
62 makes its overlap geometry measurable through v_t . *Stale predictions.* Bar-On and Mansour (2025)
63 derive upper bounds adapting to the inaccuracy of evolving observations; our lower bound instead
64 asks what delay-dependent cost is unavoidable in stale-prediction error. Wang et al. (2026) bound
65 squared gradient prediction error without delay, so the delay factor and its saturation are new here;
66 Zhang et al. (2026) share the motivation in a stochastic Bayesian model where ours is adversarial.
67 *Mediator feedback* is the closest structural connection: $v_t - 1$ is the χ^2 mutual information of Eldowa
68 et al. (2024). Our contribution is not that quantity but its role when the outcome at the observed state
69 is delayed, and the separation of that benefit from an independent temporal obstruction (Appendix F).
70 Feedback graphs and off-policy coverage (Gabbianelli et al., 2023) are related forms of information
71 sharing; Appendix A compares them.

72 3 Problem formulation

73 There are T rounds, K actions, a finite state space $\mathcal{S}$, a known action-to-state matrix P , and a
74 fixed known delay d ; all logarithms are natural and Appendix B tabulates every symbol. Every
75 round follows the chain $A_t \xrightarrow{P} S_t \xrightarrow{\theta_t} X_t$: the learner draws A_t from a distribution x_t , observes

76 $S_t \sim P(\cdot \mid A_t)$ at once, and receives the delayed outcome only at time $t + d$, encoded as a loss
77 $X_t \in [0, 1]$ so that smaller is better, with $\mathbb{E}[X_t \mid \mathcal{F}_{t-1}, A_t, S_t] = \theta_t(S_t)$, where $\mathcal{F}_{t-1}$ is everything
78 seen before round t , and outcomes conditionally independent across rounds given the states. The
79 conditional mean depends on A_t only through S_t . This *state-sufficiency* assumption licenses cross-
80 action pooling: an outcome observed at state S_t may update every action that could have produced
81 that state. If an action affects the outcome directly, in a way not represented in S_t , the argument no
82 longer applies. Here x_t depends only on $\mathcal{F}_{t-1}$; X_r may first affect the action distribution at round
83 $r + d + 1$, equivalently the round- t decision uses outcomes only from rounds $r \leq t - d - 1$.

84 That matrix $P \in [0, 1]^{K \times |S|}$ is row-stochastic and fixed, and the *state-loss means* $\theta_t \in [0, 1]^{|S|}$
85 are *oblivious*: chosen in advance, not reacting to the learner. Action losses are $c_t = P\theta_t$ and
86 performance is the static *pseudo-regret*, the gap to the best single action in hindsight, $\mathbb{E}[R_T] =$
87 $\mathbb{E}[\sum_t c_t(A_t)] - \min_a \sum_t c_t(a)$. **Oracle for the lower bound.** For Theorem 4 only, the learner also
88 receives the *stale action-loss vector* $m_t \triangleq c_{t-d}$ before choosing A_t , with $m_t = c_1$ for $t \leq d$. This is
89 not a function of the observables, so a lower bound surviving it is stronger, and STATE-EXP3 never
90 uses it. We measure how misleading m_t is by $E_2 = \sum_t \|c_t - m_t\|_\infty^2 \leq T$, which stays small when
91 errors are small even if frequent. The geometry of P thus governs sharing across actions and the
92 evolution of θ_t governs staleness across time, the two axes of Sections 4 and 6.

93 4 Pooling through the state

94 When an outcome arrives, do not update only the action played: update every action by how likely
95 it was to produce the observed state. Here P is known and the learner does not receive m_t . Let
96 $q_t(s) = \sum_a x_t(a)P(s \mid a)$ be the chance of seeing state s this round, and give each action $\hat{c}_t(a) =$
97 $P(S_t \mid a)X_t/q_t(S_t)$. It is correct on average, and how noisy it is depends on how much the actions
98 overlap, through the *effective dimension* $v_t = \sum_s q_t(s)^{-1} \sum_a x_t(a)P(s \mid a)^2 \in [1, |S|]$. STATE-
99 EXP3 runs m copies of EXP3 (Auer et al., 2002) in rotation at learning rate η , copy i taking every
100 m th round, and adds $\hat{c}_t$ to a copy's running total once the outcome is usable. A round costs $O(K)$
101 (Appendix C, with pseudocode and the χ^2 reading of v_t).

102 **Theorem 1.** *Let P be known and (θ_t) oblivious. Pick before the run any fixed number V with*
103 *$\sum_t v_t \leq V$ always. Then STATE-EXP3 with $m = d + 1$ and the matching learning rate satisfies*

$$\mathbb{E}[R_T] \leq \sqrt{2(d+1)V \log K}.$$

104 Writing $\bar{v}$ for the largest v_t over all ways of playing, $V = \bar{v}T \leq |S|T$ is one such choice, and it
105 depends on P alone.

106 So K has been replaced by V , which counts how different the actions really are. The proof uses two
107 facts, that the shared estimate is correct on average and no noisier than v_t , and that with $d + 1$ copies
108 in rotation each copy has its own previous outcome before it acts again. With $V = \bar{v}T$ this beats the
109 general rate $\tilde{O}(\sqrt{(K+d)T})$ (Esposito et al., 2023) whenever $\bar{v}(d+1) \log K \lesssim K + d$. Three gaps
110 remain: rotating copies multiplies the delay cost by $d + 1$; the bound uses v_t where $v_t - 1$ would be
111 correct, so it stays at $\sqrt{2(d+1)T \log K}$ even when every action behaves alike and the true regret is
112 zero; and it ignores E_2 . The first is an artifact of the rotation, and dropping it costs nothing.

113 **Theorem 2.** *Under the hypotheses of Theorem 1, pick any $V^- \geq \sum_t (v_t - 1)$ before the run. Then*
114 *STATE-EXP3 with $m = 1$ and any $\eta \leq 1/(e(d+1))$ satisfies $\mathbb{E}[R_T] \leq \sqrt{2(3.3V^- + 2dT) \log K} +$*
115 *$e(d+1) \log K + d$.*

116 That is $\tilde{O}(\sqrt{V^-} + \sqrt{dT})$: the delay is additive, and the algorithm the experiments run is covered, at
117 rates they already satisfy; the proof adapts a deterministic drift induction (Cesa-Bianchi et al., 2019;
118 Thune et al., 2019). Only the first of the three gaps above closes (Appendix K).

119 5 The representation trade-off

120 Some states may not be worth telling apart. Group them with a map g from S onto a smaller set
121 $\mathcal{Z}$ and run the same algorithm on $g(S_t)$, where $P^g(z \mid a) = \sum_{s: g(s)=z} P(s \mid a)$. Grouping never
122 raises the dimension, but the estimate now tracks a group average, so the widest spread of losses
123 inside a group, $\delta_t(g) = \max_z \max_{s, s': g(s)=g(s')=z} |\theta_t(s) - \theta_t(s')|$, is the error introduced.

Theorem 3. Under the hypotheses of Theorem 1, running on g with $\sum_t v_t(g) \leq V(g)$ gives $\mathbb{E}[R_T] \leq \sqrt{2(d+1)V(g)\log K} + 2\sum_t \delta_t(g)$.

Merging states need not preserve state-sufficiency, so the grouped estimate is right not for c_t but for a stand-in loss differing from it by at most $\delta_t(g)$ on every action (Appendix D). Keeping states apart preserves differences that matter but learns more slowly; merging shares more but pays $2\sum_t \delta_t(g)$. The best grouping is the true one: with 16 states built from four hidden groups, groupings of size 1, 2, 4, 8, 16 give regret 1590, 831, 754, 920, 1080, lowest at four, which no algorithm was told. The map g is fixed in advance; learning it from data remains open.

6 The cost of drift

How much does drift cost even when the stale action losses are free? Theorem 4 shows a cost from time alone, not a joint overlap-drift bound: the construction gives each of the J drifting directions a private state, removing useful pooling and forcing $J \leq |\mathcal{S}| - 1$ (Appendix G).

Theorem 4. Fix T , a delay $1 \leq d \leq T/2$, a drift budget $\mathcal{E} \leq T/4$, and $1 \leq J \leq \min\{K-1, |\mathcal{S}|-1\}$. Every algorithm, even one handed the exact losses from d rounds ago, meets some instance fixed in advance with $E_2 \leq \mathcal{E}$ on which $\mathbb{E}[R_T] = \Omega(\sqrt{d\mathcal{E}\min\{1+\log J, T/d\}})$. Without that help, and if $K \leq T$, the best regret any algorithm can guarantee is $\Omega(\sqrt{KT} + \sqrt{d\mathcal{E}\min\{1+\log J, T/d\}})$, up to constants.

The first term, $\sqrt{KT}$, is the usual cost of exploring. The second is the cost of delay: it grows with the wait, how far the losses move during it, and the number J of directions the best fixed action can use. The two come from different hard instances, so adding them overstates the truth by at most a factor of two.

7 What the experiments show

Pooling gives large gains when many actions share a small number of state patterns. In a synthetic recommendation funnel with $K = 200$ catalogue items and six engagement states the estimated effective dimension is only 1.97, and the state-pooled variant cuts regret against action-level weighting by 79.3, 63.6 and 38.9 per cent at $d = 10, 50, 200$. The funnel is hand-structured, not fitted to data, so it illustrates the mechanism, not a deployed system. On a Dirichlet family the gain grows with the catalogue, and an online plug-in $\hat{P}_t$ stays within 0.5 per cent of known P (study 1).

K	rotation	single copy	action-level	paired gain	wins
8	455.1	368.2	387.6	19.3 ± 3.4	73/100
40	734.3	574.5	696.1	121.5 ± 9.4	99/100
80	780.4	614.8	784.0	169.3 ± 9.9	100/100

$|\mathcal{S}| = 4$, $T = 10000$, $d = 50, 100$ paired seeds. Columns one and two are STATE-EXP3 with $d+1$ copies and with one; three is action-level. $\pm$ is one standard error of the paired difference.

These experiments study the $m = 1$ variant, covered by Theorem 2 but not by Theorem 1, whose proof needs $m = d + 1$; Appendix I also evaluates the interleaved variant. Broader comparisons reveal trade-offs. A best-state rule leads seven of nine settings, but where its assumption fails its regret grows linearly and its spread across seeds is four times ours. Against a minimax-optimal action-level delayed-bandit baseline the state channel still helps, cutting regret by 32 to 68 per cent on the funnel under the reported tuning protocol (Appendix I).

8 Discussion

The results suggest a diagnostic: pooling helps most when many actions lead to similar states, coarsening strengthens it when grouped states have similar outcome means, and fast drift limits how far delayed information can be trusted.

Conjecture 5 (open). Some algorithm attains $\mathbb{E}[R_T] = \tilde{O}(\sqrt{KT} + \sqrt{dE_2\log K})$, adapting to both axes at once.

References

- N. Alon, N. Cesa-Bianchi, O. Dekel, and T. Koren. Online learning with feedback graphs: beyond bandits. *COLT*, PMLR 40, 2015.
- P. Auer, N. Cesa-Bianchi, Y. Freund, and R. E. Schapire. The nonstochastic multiarmed bandit problem. *SIAM Journal on Computing*, 32(1):48–77, 2002.
- Y. Bar-On and Y. Mansour. Non-stochastic bandits with evolving observations. *ALT*, PMLR 272, 2025.
- N. Cesa-Bianchi and G. Lugosi. *Prediction, Learning, and Games*. Cambridge University Press, 2006.
- N. Cesa-Bianchi, C. Gentile, and Y. Mansour. Delay and cooperation in nonstochastic bandits. *JMLR*, 20, 2019.
- K. Eldowa, N. Cesa-Bianchi, A. M. Metelli, and M. Restelli. Information capacity regret bounds for bandits with mediator feedback. *JMLR*, 25(353), 2024.
- E. Esposito, S. Masoudian, H. Qiu, D. van der Hoeven, N. Cesa-Bianchi, and Y. Seldin. Delayed bandits: When do intermediate observations help? *ICML*, PMLR 202, 2023.
- E. Esposito, F. Fusco, D. van der Hoeven, and N. Cesa-Bianchi. Learning on the edge: online learning with stochastic feedback graphs. *NeurIPS*, 35, 2022.
- G. E. Flaspohler, F. Orabona, J. Cohen, S. Mouatadid, M. Oprescu, P. Orenstein, and L. Mackey. Online learning with optimism and delay. *ICML*, PMLR 139, 2021.
- G. Guinet, S. Amin, and P. Jaillet. Effective dimension in bandit problems under censorship. *NeurIPS*, 35, 2022.
- G. Gabbianelli, G. Neu, and M. Papini. Online learning with off-policy feedback. *ALT*, PMLR 201, 2023.
- Y.-G. Hsieh, F. Iutzeler, J. Malick, and P. Mertikopoulos. Multi-agent online optimization with delays. *JMLR*, 23, 2022.
- P. Joulani, A. György, and C. Szepesvári. Online learning under delayed feedback. *ICML*, 2013.
- M. Mussi, S. Drago, M. Restelli, and A. M. Metelli. Factored-reward bandits with intermediate observations. *ICML*, PMLR 235, 2024.
- G. Neu. Explore no more: improved high-probability regret bounds for non-stochastic bandits. *NeurIPS*, 28, 2015.
- T. S. Thune, N. Cesa-Bianchi, and Y. Seldin. Nonstochastic multiarmed bandits with unrestricted delays. *NeurIPS*, 32, 2019.
- C. Vernade, A. György, and T. Mann. Non-stationary delayed bandits with intermediate observations. *ICML*, PMLR 119, 2020.
- S. Wang, A. Barik, and V. Y. F. Tan. Bandit convex optimization with gradient prediction adaptivity. *arXiv:2605.22191*, 2026.
- C.-Y. Wei, H. Luo, and A. Agarwal. Taking a hint: How to leverage loss predictors in contextual bandits? *COLT*, PMLR 125, 2020.
- K. W. Zhang, T. Baldwin-McDonald, K. Ciosek, L. Maystre, and D. Russo. Impatient bandits: optimizing for the long-term without delay. *JMLR*, to appear, 2026.
- M. Zhang. Toward optimal second-order path-length guarantee for adversarial multi-armed bandits. *arXiv:2608.15996*, 2026.
- J. Zimmert and Y. Seldin. An optimal algorithm for adversarial bandits with arbitrary delays. *AISTATS*, PMLR 108, 2020.

A Extended related work

Guinet et al. (2022) use the name *effective dimension* for an unrelated quantity, the number of directions a censored-feedback problem can resolve; it is not v_t and the observation model differs.

The minimax characterization of delayed adversarial bandits traces to Cesa-Bianchi et al. (2019), and Joulani et al. (2013) give the additive form of the delay penalty in the stochastic case. On acting before the loss is known, Flaspohler et al. (2021) study optimism under delay directly, Wei et al. (2020) delimit when prediction error controls regret, Hsieh et al. (2022) show that an optimistic step must run at a larger scale than the update, and Zhang (2026) gives second-order path-length bounds.

Structure that lets one observation inform several actions is the other half of the picture. Feedback graphs (Alon et al., 2015) fix which actions are revealed together, and their stochastic version (Esposito et al., 2022) randomizes that set. Factored rewards with intermediate observations (Mussi et al., 2024) decompose the reward instead. Off-policy coverage (Gabbianelli et al., 2023) is the closest analogue of our effective dimension, since its mismatch coefficient measures the same kind of overlap between the distribution that generated the data and the one being evaluated. The distinction from all of these is that a feedback graph randomizes *who* is observed, whereas our sensor randomizes *what* is observed and the sharing is a consequence of P rather than a modeling primitive.

B Notation

Table 1: Symbols used in the paper. Every symbol is also defined at first use in the body, so this table is a reference rather than a dependency.

Symbol	Meaning
T, K, d	rounds, actions, fixed feedback delay
$\mathcal{S}, S_t$	finite state space, state observed at round t
$P \in [0, 1]^{K \times \mathcal{S} }$	row-stochastic action-to-state matrix, fixed known in Theorems 1 and 3, estimated in Appendix E
$\theta_t \in [0, 1]^{ \mathcal{S} }$	state-loss means at round t , oblivious
$c_t = P\theta_t \in [0, 1]^K$	induced loss over actions
$q_t(s) = \sum_a x_t(a)P(s a)$	state distribution induced by the play
$X_t \in [0, 1]$	outcome at time $t + d$, $\mathbb{E}[X_t \mathcal{F}_{t-1}, A_t, S_t] = \theta_t(S_t)$
x_t, A_t	action distribution at round t , action drawn from it
$m_t = c_{t-d}$	stale action-loss vector, given before the action, $m_t = c_1$ for $t \leq d$
E_2	$\sum_t \ c_t - m_t\ _\infty^2$, accumulated squared prediction error
Λ_2	$\sum_t \min\{1, \ c_t - m_t\ _2^2\}$, the measure of Bar-On and Mansour (2025)
W	$\sum_t \ \theta_t - \theta_{t-1}\ _\infty^2$, state-loss variation
$\hat{c}_t, \tilde{c}_t, \check{c}_t$	state-pooled, action-level, centered loss estimates
$\hat{g}_t$	the Safe-Pool blend $(1 - \lambda_t)\tilde{c}_t + \lambda_t\check{c}_t$, a distinct object
$\bar{c}_t, c_t^g, b_t$	conditional mean of $\hat{c}_t$, grouped surrogate loss, $\bar{c}_t - c_t$
$v_t, \bar{v}, V$	effective dimension $\sum_s q_t(s)^{-1} \sum_a x_t(a)P(s a)^2$, its supremum, its budget
$\hat{\bar{v}}$	Monte Carlo estimate of $\bar{v}$, used for measured values in Appendix I
$g, \mathcal{Z}, \delta_t(g)$	state grouping, its range, widest within-group loss spread
$\mathcal{E}, J$	budget constraining E_2 , number of independently drifting coordinates
$h, B, \varepsilon, \sigma_b^{(j)}$	block length, block count, amplitude, Rademacher signs
η, β, C	learning rate, Tsallis parameter, residual scale
α	smoothing pseudocount in $\hat{P}_t$
κ	sensor balance, $\bar{p}(s) \geq 1/(\kappa \mathcal{S})$ for $\bar{p} = K^{-1} \sum_a P(\cdot a)$
m, L_i, G_r	interleaving factor, estimate of copy i , estimates outstanding at r
$D, d_{\max}, \bar{\Delta}, \Gamma$	total delay, maximal delay, inaccuracy bound, the measure of Wang et al. (2026)

C The state-pooled estimator

The algorithm. Inputs are the action-to-state matrix P , the delay d , the number of copies m , and the rate $\eta > 0$. The analyzed setting is $m = d + 1$ and the practical variant is $m = 1$. Round r 's outcome becomes usable at round $r + d + 1$, which for $m = d + 1$ is the next round belonging to the copy that played round r .

- 233 1. Set $L_i \leftarrow (0, \dots, 0) \in \mathbb{R}^K$ for $i = 0, \dots, m-1$, and let $\mathcal{P}$ be an empty table of rounds
 234 awaiting their outcome.
 235 2. For $t = 1, \dots, T$:
 236 (a) *Deliver*. If $r \triangleq t - d - 1 \geq 1$, read $(i_r, q_r(S_r), S_r, X_r)$ from $\mathcal{P}$, remove it, and update
 237 only that copy, $L_{i_r}(a) \leftarrow L_{i_r}(a) + P(S_r | a)X_r/q_r(S_r)$ for every a . This is the
 238 pooled step, and it touches every action rather than the one played.
 239 (b) *Select the copy*. $i \leftarrow t \bmod m$.
 240 (c) *Play*. $x_t(a) \propto \exp(-\eta L_i(a))$, draw $A_t \sim x_t$, and observe S_t at once.
 241 (d) *Record*. Compute the scalar $q_t(S_t) = \sum_a x_t(a)P(S_t | a)$ and store $(i, q_t(S_t), S_t, \cdot)$
 242 in $\mathcal{P}$, to be completed when X_t arrives at time $t + d$.

243 Steps (a), (c) and (d) are each $O(K)$ and there is no optimization subproblem, so a round costs $O(K)$
 244 once P is stored. The state is held in the m vectors L_i , which is mK numbers, the matrix P , which
 245 is $K|\mathcal{S}|$ numbers, and four scalars for each of the at most d rounds in flight. Only the scalar $q_r(S_r)$
 246 is needed from round r , not the vector x_r .

247 *Remark 6* (The χ^2 form). Under x_t the pair (A_t, S_t) has law $x_t(a)P(s | a)$ with marginals x_t and
 248 q_t , so

$$\begin{aligned} \chi^2(P_{A_t, S_t} \parallel P_{A_t} \otimes P_{S_t}) &= \sum_{a, s} \frac{(x_t(a)P(s | a) - x_t(a)q_t(s))^2}{x_t(a)q_t(s)} \\ &= \sum_{a, s} \frac{x_t(a)P(s | a)^2}{q_t(s)} - 2 + 1 = v_t - 1, \end{aligned}$$

249 the cross term summing to $\sum_{a, s} x_t(a)P(s | a) = 1$ and the last to $\sum_{a, s} x_t(a)q_t(s) = 1$. So
 250 $v_t = 1$ exactly when A_t and S_t are independent, which pins the rows of P only on the support
 251 of x_t and therefore means all rows agree whenever x_t has full support, as exponential weights
 252 does. If P is deterministic and every state is reached by some action in the support of x_t , then v_t
 253 equals the number of states that support reaches. In particular, under full-support play and an onto
 254 deterministic P , $v_t = |\mathcal{S}|$. This is an identity, not a bound; it is used for reading v_t , and once in the
 255 proof of Theorem 2.

256 **Proposition 7.** Fix t , let x_t be the play distribution, $q_t(s) = \sum_a x_t(a)P(s | a)$ the induced state
 257 distribution, and $\hat{c}_t(a) = P(S_t | a)X_t/q_t(S_t)$, which is well defined because only states with
 258 $q_t(s) > 0$ occur. For statements involving $1/x_t(a)$, restrict to actions with $x_t(a) > 0$; STATE-EXP3
 259 itself has full support on every action. Then $\hat{c}_t(a) \geq 0$, $\hat{c}_t(a) \leq 1/x_t(a)$, $\mathbb{E}[\hat{c}_t(a) | \mathcal{F}_{t-1}] = c_t(a)$,
 260 and

$$\sum_a x_t(a)\mathbb{E}[\hat{c}_t(a)^2 | \mathcal{F}_{t-1}] \leq v_t := \sum_{s: q_t(s) > 0} \frac{\sum_a x_t(a)P(s | a)^2}{q_t(s)}, \quad 1 \leq v_t \leq |\mathcal{S}|,$$

261 states of zero probability being omitted throughout, which is the convention $0/0 = 0$. The action-
 262 level $\hat{c}_t(a) = \mathbf{1}\{A_t = a\}X_t/x_t(a)$ has the same first three properties and weighted second moment
 263 at most K .

264 *Proof.* Nonnegativity is immediate from $X_t \geq 0$. Given $\mathcal{F}_{t-1}$ the state S_t has distribution q_t , and
 265 $\mathbb{E}[X_t | \mathcal{F}_{t-1}, S_t = s] = \theta_t(s)$ by the model, since that conditional mean does not depend on
 266 the action, so $\mathbb{E}[\hat{c}_t(a) | \mathcal{F}_{t-1}] = \sum_s q_t(s)P(s | a)\theta_t(s)/q_t(s) = c_t(a)$. For the range, $q_t(s) \geq$
 267 $x_t(a)P(s | a)$ for every a , so $P(s | a)/q_t(s) \leq 1/x_t(a)$ and $X_t \leq 1$. For the second moment,
 268 $X_t^2 \leq 1$ gives $\mathbb{E}[\hat{c}_t(a)^2 | \mathcal{F}_{t-1}] \leq \sum_s P(s | a)^2/q_t(s)$, and exchanging the sums gives v_t . Each
 269 summand of v_t is at most 1, since $P(s | a) \leq 1$ makes $\sum_a x_t(a)P(s | a)^2 \leq q_t(s)$, and at least
 270 $q_t(s)$, since Jensen applied to $a \mapsto P(s | a)$ under x_t gives $\sum_a x_t(a)P(s | a)^2 \geq q_t(s)^2$. Summing
 271 the two bounds gives $1 \leq v_t \leq |\mathcal{S}|$. The action-level claims are the same computation with the state
 272 replaced by the played action, whose weighted second moment is $\sum_a x_t(a)/x_t(a) = K$. $\square$

273 The two ends are attained. If every action induces the same state distribution then $v_t = \sum_s q_t(s) =$
 274 1, and if P is deterministic, each action reaching a single state, then every $P(s | a)$ is 0 or 1 and
 275 each summand is exactly 1, so v_t counts the states reached and equals $|\mathcal{S}|$ when every state has an
 276 action reaching it. Disjoint row supports give $v_t = K$ rather than $|\mathcal{S}|$, since the summands over

one action's private states add to $\sum_s P(s | a) = 1$, so a sensor with many private states per action resolves few of them. So v_t reads the overlap between the rows of P rather than their length, and $\bar{v} = \sup_x v(x)$, the supremum over play distributions, depends on P alone. An outcome observed at a state updates the estimate for every action that can reach it, which is why the construction of Section 6 gives each drifting coordinate a private state, forcing $v_t = |S|$ there. The immediately observed pairs (A_t, S_t) identify P separately, though Theorem 1 assumes it known.

D Proofs of Theorems 1 and 3

Proof idea. The pooled estimator is unbiased for every action, while Proposition 7 bounds its play-weighted second moment by v_t rather than K . With $m = d + 1$ interleaved copies, the outcome generated on a round belonging to one copy is available before that copy acts again, so each copy is an ordinary exponential-weights process with immediate feedback. Summing the $d + 1$ potential inequalities gives $(d + 1) \log K / \eta$, and the pooled second moments give $\frac{\eta}{2} \sum_t v_t$.

Proof of Theorem 1. Fix a copy i and let T_i be the number of rounds it owns, so $\sum_i T_i = T$. Round t 's outcome is usable from round $t + d + 1$, which with $m = d + 1$ is exactly $t + m$, the copy's next round, so within a copy the feedback is undelayed and L_i carries exactly the estimates of that copy's earlier rounds. Exponential weights on nonnegative losses give, for any action a^* and any $\eta > 0$,

$$\sum_{t \in i} (\langle x_t, \hat{c}_t \rangle - \hat{c}_t(a^*)) \leq \frac{\log K}{\eta} + \frac{\eta}{2} \sum_{t \in i} \sum_a x_t(a) \hat{c}_t(a)^2,$$

by the usual potential telescoping with $e^{-z} \leq 1 - z + z^2/2$, which needs $z \geq 0$ and therefore $\hat{c}_t \geq 0$, supplied by Proposition 7. Each x_t is $\mathcal{F}_{t-1}$ -measurable, because L_i then holds only rounds $r \leq t - m = t - d - 1$, so taking expectations turns the left side into $\sum_{t \in i} (\langle x_t, c_t \rangle - c_t(a^*))$ by unbiasedness and the quadratic term into at most $\mathbb{E}[\sum_{t \in i} v_t]$, again by Proposition 7. Summing over the $m = d + 1$ copies and using $\sum_i \min_a \sum_{t \in i} c_t(a) \leq \min_a \sum_t c_t(a)$, which is what lets the copies be compared against one common action, gives $\mathbb{E}[R_T] \leq (d + 1) \log K / \eta + \frac{\eta}{2} \mathbb{E}[\sum_t v_t]$. If $\sum_t v_t \leq V$ surely then $\eta = \sqrt{2(d + 1) \log K / V}$ balances the two terms and gives $\sqrt{2(d + 1) V \log K}$, and $V = \bar{v}T$ is admissible by the definition of $\bar{v}$. $\square$

The delay enters as a factor because the copies do not share their estimates, each paying $\log K / \eta$ over $T / (d + 1)$ rounds. Sharing them is exactly the $m = 1$ algorithm of Conjecture 5, whose analysis is open.

Lemma 8 (Coarsening). *For every play distribution and every g , $v_t(g) \leq v_t$.*

Proof. It suffices to merge two states, the general case following by induction on the merges, and a merged group of zero probability contributes nothing on either side by the convention above. Write P_1, P_2 for the two columns, $A_j = \sum_a x_t(a) P_j(a)^2$, $B_j = \sum_a x_t(a) P_j(a)$ and $\lambda = B_1 / (B_1 + B_2)$. By Cauchy-Schwarz, $(P_1 + P_2)^2 \leq \lambda^{-1} P_1^2 + (1 - \lambda)^{-1} P_2^2$ pointwise in a , so $\sum_a x_t(a) (P_1 + P_2)^2 \leq (B_1 + B_2)(A_1 / B_1 + A_2 / B_2)$, and dividing by the merged denominator $B_1 + B_2$ shows the merged summand is at most $A_1 / B_1 + A_2 / B_2$, the two it replaces. $\square$

Proof of Theorem 3. Running the algorithm on g is running it on the sensor $(\mathcal{Z}, P^g)$. The surrogate losses below depend on x_t and so are predictable rather than oblivious, but the proof of Theorem 1 uses obliviousness nowhere, only that each c_t^g is $\mathcal{F}_{t-1}$ -measurable given the play, so it applies unchanged and bounds regret measured in the surrogate losses $c_t^g(a) = \sum_z P^g(z | a) \bar{\theta}_t(z)$, where $\bar{\theta}_t(z) = \sum_{s: g(s)=z} q_t(s) \theta_t(s) / q_t^g(z)$ is the group mean the grouped estimator is unbiased for, by the same computation as in Proposition 7 with S_t replaced by $g(S_t)$. Fix a and z . Both $\{q_t(s) / q_t^g(z)\}$ and $\{P(s | a) / P^g(z | a)\}$ are probability vectors on $\{s : g(s) = z\}$, so their θ_t -averages lie in a common interval of length $\delta_t(g)$ and differ by at most $\delta_t(g)$. Weighting by $P^g(z | a)$ and summing over z gives $|c_t^g(a) - c_t(a)| \leq \delta_t(g)$ for every a , hence $\langle x_t, c_t \rangle - c_t(a^*) \leq \langle x_t, c_t^g \rangle - c_t^g(a^*) + 2\delta_t(g)$. Summing over t adds $2 \sum_t \delta_t(g)$, and Lemma 8 makes $V(g)$ no larger than a bound valid for the raw states. $\square$

322 E An unknown action-to-state map

323 **The plug-in estimator.** The pairs (A_t, S_t) arrive without delay, so P is estimable from the
 324 learner's own play while the outcomes are still in flight. Let $N_a(t)$ count the rounds before t that
 325 played a and $N_{a,s}(t)$ those that also produced s , and set $\hat{P}_t(s | a) = (N_{a,s}(t) + \alpha) / (N_a(t) + \alpha|\mathcal{S}|)$
 326 for a smoothing constant $\alpha \in (0, 1]$ fixed independently of T , which is row-stochastic with every en-
 327 try positive, so $\hat{q}_t = x_t^\top \hat{P}_t$ is positive too, and both are $\mathcal{F}_{t-1}$ -measurable. Smoothing moves an entry
 328 from the empirical row by at most $\alpha(|\mathcal{S}| - 1) / (N_a(t) + \alpha|\mathcal{S}|)$, attained when every observation of a
 329 sits on one state, and moves the row in ℓ_1 by at most twice that. Both are of lower order than the sam-
 330 pling term at fixed α and must be carried otherwise, and Appendix I shows the choice matters in prac-
 331 tice. The plug-in estimate is $\hat{c}_t(a) = \hat{P}_t(S_t | a) X_t / \hat{q}_t(S_t)$. Write $\varepsilon_t = \max_a \|\hat{P}_t(\cdot | a) - P(\cdot | a)\|_1$
 332 for its accuracy and $\rho_t = \max_{s: q_t(s) > 0} |\hat{q}_t(s) - q_t(s)| / q_t(s)$ for the relative error it induces on the
 333 state distribution, states of zero probability under the play being omitted since the estimator never
 334 charges them.

335 **Proposition 9.** *On $\{\rho_t \leq \frac{1}{2}\}$ the plug-in estimate satisfies $\hat{c}_t \geq 0$ and $\sum_a x_t(a) \mathbb{E}[\hat{c}_t(a)^2 | \mathcal{F}_{t-1}] \leq$
 336 $2\hat{v}_t \leq 2|\mathcal{S}|$, and its conditional mean $\bar{c}_t$ satisfies $\sum_a x_t(a)(\bar{c}_t(a) - c_t(a)) = 0$ exactly, for every $\hat{P}_t$
 337 and with no hypothesis at all, and $|\bar{c}_t(a) - c_t(a)| \leq 2\varepsilon_t(a) + 2\rho_t$ for every a , where $\varepsilon_t(a) = \|\hat{P}_t(\cdot | a) -$
 338 $P(\cdot | a)\|_1$.*

339 *Proof.* Write $\Delta_t = \hat{P}_t - P$ and $\delta_t = \hat{q}_t - q_t = \sum_a x_t(a) \Delta_t(\cdot | a)$. Substituting these into
 340 $\bar{c}_t(a) = \sum_s q_t(s) \hat{P}_t(s | a) \theta_t(s) / \hat{q}_t(s)$ and cancelling gives $\bar{c}_t(a) - c_t(a) = \sum_s \theta_t(s) [q_t(s) \Delta_t(s | a) -$
 341 $a) - P(s | a) \delta_t(s)] / \hat{q}_t(s)$, so with $\theta_t \in [0, 1]^{|\mathcal{S}|}$,

$$|\bar{c}_t(a) - c_t(a)| \leq \sum_s \frac{q_t(s) |\Delta_t(s | a)| + P(s | a) |\delta_t(s)|}{\hat{q}_t(s)}.$$

342 The play-weighted identity needs no hypothesis. By definition $\sum_a x_t(a) \hat{P}_t(s | a) = \hat{q}_t(s)$, so
 343 $\sum_a x_t(a) \hat{c}_t(a) = X_t$ whatever $\hat{P}_t$ is, and taking conditional expectations gives $\sum_a x_t(a) \bar{c}_t(a) =$
 344 $\sum_s q_t(s) \theta_t(s) = \langle x_t, c_t \rangle$. On $\{\rho_t \leq \frac{1}{2}\}$ we have $q_t \leq 2\hat{q}_t$, and for the per-action bias the
 345 first sum of the display is at most $2\|\Delta_t(\cdot | a)\|_1 = 2\varepsilon_t(a)$ and the second at most $2\sum_s P(s | a) |\delta_t(s)| / q_t(s) \leq 2\rho_t$, because $P(\cdot | a)$ is a probability vector. The play-weighted identity
 346 does *not* survive reweighting: at $x_t = (\frac{1}{2}, \frac{1}{2})$, $P = I$, $\hat{P} = \begin{pmatrix} .9 & .1 \\ .2 & .8 \end{pmatrix}$ and $\theta_t = (0, 1)$ the
 347 bias is $(\frac{1}{9}, -\frac{1}{9})$ with zero x_t -average, while weights $(1, 0)$ leave $\frac{1}{18}$. For the second moment,
 348 $\sum_a x_t(a) \mathbb{E}[\hat{c}_t(a)^2 | \mathcal{F}_{t-1}] \leq \sum_s \frac{q_t(s)}{\hat{q}_t(s)} \cdot \frac{\sum_a x_t(a) \hat{P}_t(s | a)^2}{\hat{q}_t(s)} \leq 2\hat{v}_t$, and $\hat{v}_t \leq |\mathcal{S}|$ by Proposition 7
 349 applied to $\hat{P}_t$. $\square$

351 **The warm-start guarantee.** WARM-START STATE-EXP3 plays uniformly for n_0 rounds, forms
 352 $\hat{P} = \hat{P}_{n_0+1}$ from those rounds alone, discards the cumulative estimates together with every outcome
 353 of a warm-start round that has not yet arrived, and then runs STATE-EXP3 with $m = d + 1$, that
 354 frozen $\hat{P}$, and each play mixed with the uniform action at rate γ . Freezing makes $\hat{P}$ measurable with
 355 respect to the first phase, so ε and ρ are fixed before the second phase begins, and discarding is what
 356 keeps the first phase's inaccurate updates out of every later distribution.

357 **Theorem 10.** *Suppose $\bar{p}(s) \geq 1/(\kappa|\mathcal{S}|)$ for every s , fix $0 < \gamma \leq \frac{1}{2}$ and $1 \leq d \leq T/2$, and take
 358 $n_0 \geq c\gamma^{-2}K\kappa^2|\mathcal{S}|^2 \log(K|\mathcal{S}|T)$ for a universal c . Then WARM-START STATE-EXP3 satisfies, for
 359 every $\eta > 0$,*

$$\mathbb{E}[R_T] \leq n_0 + 1 + \gamma T + \frac{(d+1) \log K}{\eta} + 2\eta|\mathcal{S}|T + T(4\bar{\varepsilon} + 4\bar{\rho}),$$

360 where $\bar{\varepsilon} = \tilde{O}(\sqrt{K|\mathcal{S}|/n_0})$ and $\bar{\rho} = \tilde{O}(\kappa|\mathcal{S}| \sqrt{K/n_0}/\gamma) \leq \frac{1}{2}$ bound ε and ρ on an event of proba-
 361 bility at least $1 - 1/T$. Tuning η and then $\gamma \asymp (\kappa|\mathcal{S}|)^{1/2} (K/n_0)^{1/4}$ and $n_0 \asymp T^{4/5} (\kappa|\mathcal{S}|)^{2/5} K^{1/5}$
 362 gives $\mathbb{E}[R_T] = \tilde{O}(\sqrt{d|\mathcal{S}|T \log K} + (\kappa|\mathcal{S}|)^{2/5} K^{1/5} T^{4/5})$, provided $T = \tilde{\Omega}(K(\kappa|\mathcal{S}|)^2)$. That con-
 363 dition constrains T against K , $|\mathcal{S}|$ and κ only, so the hypothesis $d \leq T/2$ is what carries the delay.
 364 Together they give $n_0 \leq T/2 \leq T - d$, and the same condition supplies the lower bound above with
 365 some $\gamma \leq \frac{1}{2}$. Outside that regime the tuning is vacuous and $R_T \leq T$ is the statement.

366 *Proof.* The first phase has losses in $[0, 1]$, so it contributes at most n_0 . Let G be the event that
 367 $N_a(n_0 + 1) \geq n_0/(2K)$ for every a , that $\|\hat{P}(\cdot | a) - P(\cdot | a)\|_1 \leq \bar{\varepsilon}$ for every a , and that
 368 $\max_{s,a} |\hat{P}(s | a) - P(s | a)| \leq \bar{\rho}\gamma/(\kappa|\mathcal{S}|)$. Uniform play gives $\mathbb{E}[N_a] = n_0/K$, so a multiplicative
 369 Chernoff bound, the ℓ_1 deviation bound for a multinomial, Hoeffding, and a union bound over the
 370 $K|\mathcal{S}|$ entries put $\Pr(G^c) \leq 1/T$ at the stated n_0 , the pseudo-counts adding $2\alpha(|\mathcal{S}| - 1)/(N_a + \alpha|\mathcal{S}|)$
 371 to each ℓ_1 deviation and half that to each entry, of lower order at fixed α . Off G bound the second
 372 phase by T , contributing at most 1.

373 On G the quantity $\hat{P}$ is measurable with respect to the first phase, and mixing gives $q_t(s) \geq \gamma\bar{p}(s) \geq$
 374 $\gamma/(\kappa|\mathcal{S}|)$ at every later round, so $\rho_t \leq \kappa|\mathcal{S}| \max_{s,a} |\hat{P}(s | a) - P(s | a)|/\gamma \leq \bar{\rho} \leq \frac{1}{2}$ for every
 375 $t > n_0$ and every realization of the second phase. Proposition 9 therefore applies at every such
 376 round, and conditionally on the first phase the second is exactly the algorithm of Theorem 1 driven
 377 by a nonnegative estimate with weighted second moment at most $2\hat{v}_t \leq 2|\mathcal{S}|$. Its proof gives
 378 $\mathbb{E}[\sum_{t>n_0} (\langle p_t, \bar{c}_t \rangle - \bar{c}_t(a^*))] \leq (d+1) \log K/\eta + 2\eta|\mathcal{S}|T$, the extra factor of two coming from
 379 $p_t \leq x_t/(1-\gamma) \leq 2x_t$. Playing x_t rather than p_t costs at most γ per round since $c_t \in [0, 1]^K$, so
 380 it remains to pass from $\bar{c}_t$ to c_t under p_t . Proposition 9 gives $\langle x_t, b_t \rangle = 0$ for the per-action bias b_t ,
 381 and $x_t = (1-\gamma)p_t + \gamma u$ for the uniform u , so $\langle p_t, b_t \rangle = -\frac{\gamma}{1-\gamma} \langle u, b_t \rangle$, which for $\gamma \leq \frac{1}{2}$ is at most
 382 $\max_a |b_t(a)|$ in absolute value. Each round therefore adds $|\langle p_t, b_t \rangle| + |b_t(a^*)| \leq 2 \max_a |b_t(a)| \leq$
 383 $4\bar{\varepsilon} + 4\bar{\rho}$. Collecting the four contributions gives the display. For the rate, balancing γ/T against $2T\bar{\rho}$
 384 gives the stated γ and a combined $\tilde{O}(T(\kappa|\mathcal{S}|)^{1/2}(K/n_0)^{1/4})$, and balancing that against n_0 gives
 385 the stated n_0 , under which $T\bar{\varepsilon}$ and the constraint on n_0 are of lower order. $\square$

386 **Safe blending.** Blending the plug-in with the action-level estimate makes the guarantee safe. Let
 387 $\lambda_t \in [0, 1]^K$ be $\mathcal{F}_{t-1}$ -measurable and put $\check{g}_t(a) = (1 - \lambda_t(a))\bar{c}_t(a) + \lambda_t(a)\hat{c}_t(a)$.

388 **Proposition 11.** $\check{g}_t \geq 0$, its conditional mean satisfies $\mathbb{E}[\check{g}_t(a) | \mathcal{F}_{t-1}] - c_t(a) = \lambda_t(a)(\bar{c}_t(a) -$
 389 $c_t(a))$ for the plug-in's conditional mean $\bar{c}_t$ of Proposition 9, and

$$\sum_a x_t(a) \mathbb{E}[\check{g}_t(a)^2 | \mathcal{F}_{t-1}] \leq (K - \Pi_t) + \hat{V}_t,$$

390 for $\Pi_t = \sum_a \lambda_t(a)$ and $\hat{V}_t = \sum_a \lambda_t(a)x_t(a)\mathbb{E}[\hat{c}_t(a)^2 | \mathcal{F}_{t-1}]$, which satisfies $\hat{V}_t \leq 2\hat{v}_t$ on $\{\rho_t \leq$
 391 $\frac{1}{2}\}$ and not in general, since the step $q_t/\hat{q}_t \leq 2$ behind it is exactly that event. We write $V_t^{\text{ub}}(\tau)$
 392 below for a computable upper bound on $\hat{V}_t$, reserving $\hat{V}_t$ for the conditional expectation itself.

393 *Proof.* Both parts are nonnegative and $\lambda_t(a) \in [0, 1]$. The mean identity is linearity together
 394 with $\mathbb{E}[\bar{c}_t(a) | \mathcal{F}_{t-1}] = c_t(a)$. For the second moment, $u \mapsto u^2$ is convex, so $\check{g}_t(a)^2 \leq$
 395 $(1 - \lambda_t(a))\bar{c}_t(a)^2 + \lambda_t(a)\hat{c}_t(a)^2$, and $x_t(a)\mathbb{E}[\bar{c}_t(a)^2 | \mathcal{F}_{t-1}] \leq 1$ by Proposition 7. $\square$

396 So a weight of $\lambda_t(a)$ on the plug-in buys a variance saving of $\lambda_t(a)$ and pays a bias of $\lambda_t(a)$ times the
 397 plug-in's, and both endpoints are exact. Let $r_t(a)$ and $\hat{\rho}_t$ be any sequences computable from the im-
 398 mediate pairs that are *simultaneously valid*, meaning $\varepsilon_t(a) \leq r_t(a)$ for every a and $\rho_t \leq \hat{\rho}_t$ at every
 399 $t \leq T$, on an event of probability at least $1 - 1/T$. Write $V_t^{\text{ub}}(\tau) = 2 \sum_{a:r_t(a) \leq \tau} x_t(a) \sum_s \hat{P}_t(s |$
 400 $a)^2/\hat{q}_t(s)$, which the learner can evaluate and which dominates $\hat{V}_t$ on $\{\rho_t \leq \frac{1}{2}\}$. SAFE-POOL
 401 STATE-EXP3 sets $\lambda_t \equiv 0$ when $\hat{\rho}_t > \frac{1}{2}$, and otherwise picks the best member of the candidate
 402 set $\{\perp\} \cup \{r_t(a) : a \leq K\}$, where $\perp$ stands for $\lambda_t \equiv 0$ and a numerical candidate τ stands for
 403 $\lambda_t(a) = \mathbf{1}\{r_t(a) \leq \tau\}$, best meaning smallest

$$\Psi_t(\tau) = \frac{\eta}{2} [(K - \Pi_t(\tau)) + V_t^{\text{ub}}(\tau)] + (4\tau + 4\hat{\rho}_t) \mathbf{1}\{\Pi_t(\tau) > 0\}, \quad \Psi_t(\perp) = \frac{\eta}{2} K.$$

404 Write Ψ_t^* for that smallest value. The candidate $\perp$ is listed apart from the numerical thresholds
 405 because a valid radius may equal zero, in which case $\tau = 0$ pools the actions attaining it rather than
 406 abstaining, so $\perp$ rather than $\tau = 0$ is what keeps $\frac{\eta}{2}K$ available on every round.

407 **Theorem 12.** For every $\eta > 0$, SAFE-POOL STATE-EXP3 with $m = d + 1$ satisfies

$$\mathbb{E}[R_T] \leq 1 + \frac{(d+1) \log K}{\eta} + \mathbb{E}\left[\sum_{t=1}^T \Psi_t^*\right],$$

so $\eta = \sqrt{2(d+1)\log K/(KT)}$ gives $\mathbb{E}[R_T] \leq 1 + \sqrt{2(d+1)KT\log K}$ on every instance, which is the guarantee of action-level exponential weights, with an improvement on each round whose Ψ_t^* is strictly below $\frac{\eta}{2}K$. The certificate is computed from the learner's own observations and so is random, which is why the display carries an expectation. The tuned bound does not, because $\Psi_t^* \leq \frac{\eta}{2}K$ holds on every path.

Proof. On the validity event, of probability at least $1 - 1/T$ and contributing at most 1 off it, pooling happens only when $\hat{\rho}_t \leq \frac{1}{2}$ and hence only when $\rho_t \leq \frac{1}{2}$, which is what makes V_t^{ub} dominate the conditional second moment. Proposition 11 then supplies a nonnegative estimate whose weighted second moment is at most $(K - \Pi_t) + \hat{V}_t$, so the argument of Appendix D applies to the surrogate losses. Passing to c_t adds both the played bias $|\sum_a x_t(a)\lambda_t(a)b_t(a)|$ and the comparator bias $\lambda_t(a^*)|b_t(a^*)|$, for $b_t = \bar{c}_t - c_t$. The cancellation of Proposition 9 is lost under per-action weights, as the counterexample there shows, so each is bounded separately. On a pooled action $\varepsilon_t(a) \leq r_t(a) \leq \tau$ gives $|b_t(a)| \leq 2\tau + 2\hat{\rho}_t$, and since $\sum_a x_t(a)\lambda_t(a) \leq 1$ and $\lambda_t(a^*) \leq 1$ the two together are at most $4\tau + 4\hat{\rho}_t$, charged only when something is pooled. The second-moment bound and the two bias terms are all functions of the observations, so the per-round costs they assemble into are the random Ψ_t^* and the sum is taken in expectation. For the tuned bound, $\perp$ belongs to the candidate set and $\Psi_t(\perp) = \frac{\eta}{2}K$, so $\Psi_t^* \leq \frac{\eta}{2}K$ on every path whatever the radii are. $\square$

Why certification stays conservative. The guarantee is safe and the algorithm is inert. Pooling is certified when $4(\tau + \hat{\rho}_t) \leq \frac{\eta}{2}(\Pi_t - V_t^{\text{ub}})$. At the tuned η both sides fall as $T^{-1/2}$, the left through $\tau \asymp \sqrt{K\log(K|\mathcal{S}|T)/T}$ under uniform play and the right through $\eta \asymp \sqrt{\log K/(KT)}$. Their ratio is $\Theta(\sqrt{\log(K|\mathcal{S}|T)/\log K})$, which exceeds one and does not vanish with the horizon, growing slowly in T through the logarithmic factors, so no horizon makes the certificate clear. Appendix I reports behavior consistent with this explanation, the algorithm reproducing action-level regret to the digit with a pooled fraction of zero.

The obstruction is not the shape of the bias bound. Regret charges $\langle x_t, b_t \rangle - b_t(a^*)$ for the per-action bias b_t , while Proposition 9 charges $\max_a |b_t(a)|$. Measured against the budget itself at $K = 40$, $|\mathcal{S}| = 8$, $T = 8000$ and $d = 50$, the first is 0.42 times it and the second 1.67 times it, so on that trajectory the realized bias lies below the best-case variance budget and a tighter bias bound would close the analytical gap. Two exact refinements fail to deliver one. Both $\sum_s \Delta_t(s | a)$ and $\sum_s \delta_t(s)$ vanish, so a constant may be subtracted and a maximum replaced by an oscillation, and Cauchy-Schwarz against q_t replaces the comparator term by $\sqrt{\nu_t(a^*)}\chi_t$ for its own effective dimension ν_t and a chi-square χ_t between $\hat{q}_t$ and q_t . Measured, these are 6.1 and 63 times the truth against 3.9 for the maximum they replace.

What the algorithm can verify from its own counts is 104 times the budget, and Bernstein radii do not improve on Hoeffding ones (Appendix I), so the barrier is certifying ρ_t , the relative error the estimate induces on the state distribution. The play-weighted identity of Proposition 9 says the whole charge is the comparator's, and it routes through ρ_t only because $\hat{q}_t = x_t^\top \hat{P}_t$ predicts the state distribution out of every row of P . Predicting is expensive for that reason, needing $t = \Omega(K|\mathcal{S}|\log(K|\mathcal{S}|T)/\rho^2)$ rounds, and the resulting certificate clears for no $K \leq T$, since it would need $\log K \geq 16c|\mathcal{S}|$ for $c = \log(2|\mathcal{S}|T)$, which $\log K \leq \log T$ cannot meet.

Counting the states costs less and still buys nothing. With x_t frozen over a window the empirical state frequency is a plain multinomial estimate of q_t , needing $n = \Omega(|\mathcal{S}|\log(|\mathcal{S}|T)/\rho^2)$ samples, a factor K fewer than predicting and free of P , and the certificate does clear once $K\log K \geq 64c|\mathcal{S}|$, about $K \geq 10^3$ at $|\mathcal{S}| = 8$. The window that buys it must satisfy $n \geq 32c|\mathcal{S}|T/(K\log K)$, a constant fraction of the horizon, and freezing the play over blocks of length B costs $\sqrt{2TBK\log K}$, which at that B is $T\sqrt{128c|\mathcal{S}|}$ and so linear in T . Both routes fail for one reason: the certificate must resolve q_t to relative accuracy $\Theta(\sqrt{K\log K/T})$, finer than either estimate delivers from the data the learner has.

In words, an unknown action-to-state matrix costs a bias in two accuracies, the ℓ_1 error of $\hat{P}$ and the relative error it induces on the state distribution, and only the second needs forced exploration, which is what sets the rate. **The bound does not certify an improvement from the channel it analyzes.** Its second term exceeds $\sqrt{KT}$ at every horizon with $T > K$, by a factor $(T/K)^{3/10}$, so a learner holding this guarantee alone should run action-level exponential weights instead. The

experiments suggest the analysis is conservative rather than the algorithm weak, since Appendix I shows it tracking known P closely, and sharpening it is the natural next result.

F Feedback models, and what transfers

Write $\Lambda_2 = \sum_t \min\{1, \|c_t - m_t\|_2^2\}$ for the inaccuracy measure of Bar-On and Mansour (2025). Unlike E_2 , which charges only the worst action in a round, Λ_2 grows with how many actions are predicted wrongly, and since $\|v\|_\infty \leq \|v\|_2 \leq \sqrt{J}\|v\|_\infty$ for a vector v with J nonzero coordinates, $E_2 \leq \Lambda_2 \leq J E_2$ with both ends attained.

Bar-On and Mansour (2025) analyze a model in which the played coordinate of the loss vector is observed exactly. Section 3 instead observes a random outcome whose conditional mean is the loss. The following records precisely what survives the substitution.

Proposition 13. *Take the action-level $\tilde{c}_t(a) = \mathbf{1}\{A_t = a\}X_t/x_t(a)$ of Proposition 7 under the model of Section 3. Then for every action a , (i) $\mathbb{E}[\tilde{c}_t(a) \mid \mathcal{F}_{t-1}] = c_t(a)$, (ii) $0 \leq \tilde{c}_t(a) \leq 1/x_t(a)$, and (iii) $\mathbb{E}[\tilde{c}_t(a)^2 \mid \mathcal{F}_{t-1}] \leq 1/x_t(a)$. The centered estimator $\tilde{c}_t(a) = m_t(a) + \mathbf{1}\{A_t = a\}(X_t - m_t(a))/x_t(a)$ also satisfies (i), and its residual obeys $\mathbb{E}[(\tilde{c}_t(a) - m_t(a))^2 \mid \mathcal{F}_{t-1}] \leq 1/x_t(a)$.*

Proof. For (i), by the tower rule $\mathbb{E}[\mathbf{1}\{A_t = a\}X_t \mid \mathcal{F}_{t-1}] = x_t(a) \mathbb{E}[X_t \mid A_t = a] = x_t(a)c_t(a)$, and dividing by $x_t(a)$ gives the claim. Item (ii) holds because $X_t \in [0, 1]$. For (iii), $\mathbb{E}[\tilde{c}_t(a)^2 \mid \mathcal{F}_{t-1}] = \mathbb{E}[X_t^2 \mathbf{1}\{A_t = a\}]/x_t(a)^2 = \mathbb{E}[X_t^2 \mid A_t = a]/x_t(a) \leq 1/x_t(a)$ because $X_t \leq 1$. The statements for $\tilde{c}_t$ follow the same computation with X_t replaced by $X_t - m_t(a)$, using $|X_t - m_t(a)| \leq 1$. $\square$

In words, replacing an exactly observed loss by a bounded unbiased draw of it changes neither the mean of the estimator nor the moment bounds an analysis consumes.

Under exact feedback the corresponding estimator satisfies (i) to (iii) with identical bounds, since $c_t(a) \leq 1$. Consequently any step of an exact-feedback analysis that uses the observation only through unbiasedness and these two moment bounds is unchanged under noisy outcomes. What this does not cover is the comparison between the estimate built from the stale action-loss vector and the estimate built from the true loss. Under exact feedback that difference is determined by the played action. Under noisy outcomes it carries an additional mean-zero term of magnitude up to $1/x_t(A_t)$, and bounding its cumulative effect over a window of d rounds is what requires the stability property falsified in Appendix H. An upper bound for the noisy model is Theorem 1. What remains open is one adaptive to E_2 .

G Proofs

The measure E_2 is controlled by the state-loss variation W , and no converse bound holds in general.

Lemma 14 (Delay window). *$E_2 \leq d^2 W$ for $W = \sum_t \|\theta_t - \theta_{t-1}\|_\infty^2$. The ratio $E_2/(d^2 W)$ tends to one when θ drifts monotonically in one coordinate at a constant rate and $T/d \rightarrow \infty$, the first d rounds contributing $\sum_{j < d} j^2$ rather than d^2 each under the convention $m_t = c_1$.*

In words, accumulated squared prediction error is at most the state-loss variation inflated by the square of the delay, and no better bound in W alone is available.

Proof of Lemma 14. Throughout write $\Delta_j = \theta_j - \theta_{j-1}$, set $\theta_0 = \theta_1$, so $\Delta_1 = 0$, and read c_r as c_1 for $r \leq 0$, matching the convention $m_t = c_1$ for $t \leq d$ of Section 3. Then $c_t - c_{t-d} = P \sum_{j=t-d+1}^t \Delta_j$. Because the rows of P are probability vectors, averaging cannot increase the supremum norm, so $\|c_t - c_{t-d}\|_\infty \leq \sum_j \|\Delta_j\|_\infty$. By Cauchy-Schwarz applied to the d summands, $\|c_t - c_{t-d}\|_\infty^2 \leq d \sum_{j=t-d+1}^t \|\Delta_j\|_\infty^2$. Summing over t and using that each increment belongs to at most d windows gives $E_2 \leq d^2 W$. The ratio $E_2/(d^2 W)$ approaches one when every increment has the same magnitude, sign and active coordinate, and some action reaches that coordinate deterministically, since then no cancellation occurs within a full window and P passes the increment through undiminished. Without the second condition the drift can lie in the kernel of P and the ratio can be zero. $\square$

507 **Lemma 15** (Information isolation). *If $h \leq d$ then A_t is independent of the signs governing its own*
 508 *block, and this remains true when the learner is provided with the exact stale action-loss vector m_t .*

509 *Proof of Lemma 15.* For t in block b , every realized outcome usable when choosing A_t comes from
 510 a round $r \leq t - d - 1$. Since $t \leq bh$ and $h \leq d$, we have $r \leq bh - d - 1 < (b - 1)h$, so every usable
 511 realized outcome depends only on signs from blocks strictly before b , whose joint law depends only
 512 on $\sigma_2, \dots, \sigma_{b-1}$. Hence by conditional independence A_t is independent of block b 's signs. For the
 513 oracle, take any $h \leq d$. If $t > d$ then $t - d \leq bh - d \leq (b - 1)h$ by the same computation,
 514 so $m_t = c_{t-d}$ is measurable with respect to the signs of blocks strictly before b . If $t \leq d$ then
 515 $m_t = c_1 = \frac{1}{2}\mathbf{1}$, which is deterministic because block 1 is neutral. In both cases m_t is independent
 516 of block b 's signs, so adjoining it to the learner's information leaves the conclusion intact. $\square$

517 The construction in full. Take states $s_0, \dots, s_q$, a stationary deterministic map sending a_j to s_j for
 518 $j \leq J$ and every other action to s_0 , and independent Rademacher signs $\sigma_b^{(j)}$ for $2 \leq b \leq B = \lfloor T/h \rfloor$.
 519 Set $\theta_t(s_0) = \frac{1}{2}$ throughout, $\theta_t \equiv \frac{1}{2}$ on block 1, and $\theta_t(s_j) = \frac{1}{2} - \varepsilon \sigma_b^{(j)}$ on each later block b . The
 520 neutral first block makes $c_1 = \frac{1}{2}\mathbf{1}$, so the convention $m_t = c_1$ for $t \leq d$ of Section 3 supplies a
 521 constant over exactly the rounds that precede any usable feedback, which is what Lemma 15 needs
 522 at $b = 1$. Rounds after the last complete block carry no drift, costing a constant factor.

523 **Lemma 16** (Rademacher lower tail). *Let S_B be a sum of $B \geq 32$ independent Rademacher signs.*
 524 *For every $1 \leq u \leq \sqrt{B}/32$, $\Pr(S_B \geq u\sqrt{B}) \geq \frac{u}{4}e^{-64u^2}$.*

525 *Proof.* Write $S_B = 2X - B$ with $X \sim \text{Bin}(B, \frac{1}{2})$, so that $\{S_B \geq u\sqrt{B}\} = \{X \geq B/2 + u\sqrt{B}/2\}$.
 526 Passing to X removes the lattice gap, since X takes every integer value in $[0, B]$ whereas S_B takes
 527 only those of the parity of B . Put $n = \lfloor B/2 \rfloor$ and $t = \lceil u\sqrt{B}/2 \rceil + 1$, so $X \geq n + t$ implies the
 528 event, and $t \leq u\sqrt{B}$ because $u\sqrt{B} \geq \sqrt{32} > 4$.

529 For $j \geq 1$ the ratio of consecutive binomial probabilities is $\Pr(X = n + j)/\Pr(X = n + j - 1) =$
 530 $(B - n - j + 1)/(n + j)$. Using $(B - 1)/2 \leq n \leq B/2$, its distance from one is at
 531 most $(2j - 1)/(n + j) \leq (4j - 2)/(B - 1) \leq 8j/B$ for $B \geq 2$. Restricting to $j \leq B/16$ keeps
 532 $8j/B \leq \frac{1}{2}$, where $\log(1 - x) \geq -2x$ holds, so summing over $i \leq j$ gives $\Pr(X = n + j) \geq \Pr(X =$
 533 $n)e^{-16j(j+1)/(2B)} \geq \Pr(X = n)e^{-16j^2/B}$. The central term satisfies $\Pr(X = n) \geq 1/(2\sqrt{B})$ for
 534 every $B \geq 2$.

535 Sum over the t integers $j \in \{t, \dots, 2t - 1\}$, all of which are at most $2t$ and hence at most $B/16$
 536 because $u \leq \sqrt{B}/32$ forces $2t \leq 2u\sqrt{B} \leq B/16$. Each term is at least $e^{-64t^2/B}/(2\sqrt{B})$, and
 537 $t \leq u\sqrt{B}$ makes the exponent at most $64u^2$, so $\Pr(X \geq n + t) \geq t e^{-64u^2}/(2\sqrt{B}) \geq \frac{u}{4}e^{-64u^2}$
 538 using $t \geq u\sqrt{B}/2$. $\square$

539 Three regimes give the maximal inequality the lower bound needs. Each $S_B^{(j)}$ is sub-Gaussian
 540 with variance proxy B , so $\mathbb{E} \max_j (S_B^{(j)})^+ \leq \sqrt{2B \log J} + \sqrt{B}$ throughout. For the matching
 541 lower bound, first suppose $\log J \leq B/8$ and $\log J \geq 128$. Taking $u = \sqrt{\log J/128}$, which
 542 then lies in $[1, \sqrt{B}/32]$, makes $e^{-64u^2} = J^{-1/2}$, so $\Pr(S_B^{(j)} \geq u\sqrt{B}) \geq \frac{u}{4}J^{-1/2}$ and inde-
 543 pendence across j gives $\Pr(\max_j S_B^{(j)} < u\sqrt{B}) \leq e^{-u\sqrt{J}/4}$, at most $\frac{1}{2}$ beyond an absolute J .
 544 Hence $\mathbb{E} \max_j (S_B^{(j)})^+ \geq \frac{1}{2}u\sqrt{B} = \Omega(\sqrt{B \log J})$. Second, if $\log J > B/8$ and $B \geq 1024$, take
 545 $u = \sqrt{B}/32$, so $u\sqrt{B} = B/32$ and $e^{-64u^2} = e^{-B/16}$, and $J \geq e^{B/8}$ makes $J e^{-B/16}$ diverge,
 546 giving $\mathbb{E} \max_j (S_B^{(j)})^+ = \Omega(B)$, which is $\Omega(\sqrt{B \min\{1 + \log J, B\}})$ in that regime. Third, when
 547 $\log J < 128$ or $B < 1024$ the quantity $\sqrt{B \min\{1 + \log J, B\}}$ is $O(\sqrt{B})$, and a single walk already
 548 gives $\mathbb{E}(S_B^{(1)})^+ = \frac{1}{2}\mathbb{E}|S_B| \geq \frac{1}{2}(\mathbb{E}S_B^2)^{3/2}(\mathbb{E}S_B^4)^{-1/2} \geq \frac{1}{2}\sqrt{B/3}$ by Hölder, since $\mathbb{E}S_B^2 = B$ and
 549 $\mathbb{E}S_B^4 \leq 3B^2$. Combining the three regimes, $\mathbb{E} \max_j (S_B^{(j)})^+ = \Theta(\sqrt{B \min\{1 + \log J, B\}})$. This is
 550 the route by which maxima of independent random walks fix the minimax rate for prediction with
 551 expert advice (Cesa-Bianchi and Lugosi, 2006).

552 **Proof idea.** Divide time into blocks of length d and give J actions private states whose losses are
 553 perturbed by independent Rademacher signs within each block. Every observation available when

an action is chosen comes from an earlier block, and the supplied $m_t = c_{t-d}$ depends only on earlier signs, so the learner cannot predict the signs governing its current block. Its expected loss is therefore $T/2$, while the best fixed action in hindsight benefits from the maximum of J independent random walks. Choosing the amplitude to make $E_2 \leq \mathcal{E}$ gives the stated scale.

Proof of Theorem 4. Take $h = d$ and $\varepsilon = \frac{1}{2}\sqrt{\mathcal{E}/T}$. By Lemma 15 every algorithm's expected loss equals $T/2$, since its action is independent of the current block's signs and all states have mean $\frac{1}{2}$ marginally. Its regret is therefore exactly the comparator's fluctuation advantage, $\varepsilon d \mathbb{E}[\max_{j \leq J} (\sum_{2 \leq b \leq B} \sigma_b^{(j)})^+]$ with $B = \lfloor T/d \rfloor$, over the $B - 1$ randomized blocks. Per round $\|c_t - m_t\|_\infty \leq 2\varepsilon$, because the supremum norm charges only the largest coordinate and each coordinate moves by at most 2ε at a block boundary; the gap is ε across the first randomized block, where m_t is still the neutral c_1 , and lies in $\{0, 2\varepsilon\}$ thereafter. Hence $E_2 \leq 4\varepsilon^2 T = \mathcal{E}$ holds deterministically. The budget must bind on the realized sequence, since a budget holding only in expectation would not constrain the instance selected by Yao's principle. By Lemma 16 and the discussion following it, the expected maximum of the positive parts of J independent Rademacher sums of length $B - 1$ is $\Theta(\sqrt{(B - 1) \min\{1 + \log J, B - 1\}})$, which covers both regimes and the case $J = 1$. The hypothesis $d \leq T/2$ gives $B \geq 2$ and hence $B - 1 \geq B/2$, so replacing $B - 1$ by B costs a factor at most $\sqrt{2}$ in each branch of the minimum, including the saturated branch where the minimum is $B - 1$ rather than $1 + \log J$. Substituting $B = \lfloor T/d \rfloor$ and ε therefore gives $\Theta(\sqrt{d\mathcal{E} \min\{1 + \log J, T/d\}})$. Rounds after the last complete block carry no drift and change the constant only. When the oracle is withheld, the $\sqrt{KT}$ term is Esposito et al. (2023, Thm. 5.1), whose construction is stationary and therefore lies in the present model class. $\square$

Proposition 17. Suppose $X_t = c_t(A_t)$. On instances whose drift is carried by a single action, $\Lambda_2 = E_2$, so the bound of Bar-On and Mansour (2025) reads $\tilde{O}(\sqrt{KT} + dE_2)$ and its drift contribution $\sqrt{dE_2}$ matches Theorem 4 up to logarithmic factors.

In words, single-action drift makes the Euclidean and supremum measures agree, so on that family the drift term of the published upper bound is unimprovable in order. Its $\sqrt{KT}$ term is not matched because the oracle supplies an entire stale loss vector, so the oracle-assisted problem need not retain bandit exploration complexity. Neither result provides an E_2 -adaptive upper bound under noisy outcomes.

Proof of Proposition 17. If only one coordinate of $c_t - m_t$ is nonzero then $\|c_t - m_t\|_2 = \|c_t - m_t\|_\infty$. On the family of Section 6 that coordinate has magnitude at most $2\varepsilon \leq 1$, so the truncation at one in Λ_2 is inactive and $\Lambda_2 = E_2$. The bound of Bar-On and Mansour (2025) then reads $\tilde{O}(\sqrt{KT} + dE_2)$, whose $\sqrt{dE_2}$ contribution Theorem 4 matches up to the logarithmic factor at $J = 1$. Its $\sqrt{KT}$ contribution is not matched. $\square$

H The two-action counterexample

Take $K = 2$ with the hybrid regularizer of Zimmert and Seldin (2020), Tsallis parameter $\beta = 0.01$, learning rate $\eta = 0.5$ and residual scale $C = 1$. The iterate solves the stationarity condition $g(x_1) + L_1 = g(x_2) + L_2$ with $x_1 + x_2 = 1$, where

$$g(x) = -\frac{1}{2\beta}x^{-1/2} + \frac{1}{\eta}\log x$$

and L_a is the cumulative estimate for action a . Suppose the cumulative estimates place action 1 at probability x_1 . A single maximally negative centered residual $-C/x_1$ is then delivered to action 1. This is realizable in the model of Section 3, since $m_t(1) = 1$ and $X_t = 0$ give exactly $\check{c}_t(1) - m_t(1) = -C/x_1$ with $C = 1$. Solving the stationarity condition again gives the following.

x_1 before	x_1 after	ratio
9.15×10^{-3}	1.46×10^{-2}	1.6
9.83×10^{-4}	7.39×10^{-3}	7.5
2.19×10^{-4}	9.98×10^{-1}	4561
5.80×10^{-5}	1.00	17226
4.86×10^{-6}	1.00	205641

597 The transition sits at $4\beta^2 C^2 = 4 \times 10^{-4}$, which is where the Tsallis term $x^{-1/2}$ ceases to dominate
598 the entropy term $\log x$ in g , so the one-step linearization underlying a constant-factor bound fails.
599 The violated statement is the local-stability inequality asserting that $x_{t'}(a)/x_t(a)$ stays bounded by
600 a constant for all t' within d rounds of t . A single round refutes it, so no accumulation argument
601 is needed. This invalidates that lemma and the proof route resting on it, and not every argument
602 that might use the same regularizer. It does not refute Conjecture 5, which concerns the regret rate
603 and may be provable by other means. Devices that bound the residual magnitude introduce either a
604 first-order bias dominating the target term or a probability floor narrowing the delay regime.

605 I Numerical detail




Figure 2: Relative gain $(R_{\text{action}} - R_{\text{state}})/R_{\text{action}}$ of pooling through the state, for the practical $m = 1$ variant, at $K = 40$, $|S| = 8$, $T = 5000$ and 12 paired seeds. Blue is a gain, and $\hat{v}$ denotes a Monte Carlo estimate of $\hat{v} = \sup_x v(x)$ obtained from 4000 sampled play distributions of the *raw* map, which for the coarsened panel is not the $\bar{v}(P^g)$ that Theorem 3 charges. The dashed line marks where the bound of Theorem 1, which analyzes $m = d + 1$, falls below $\tilde{O}(\sqrt{(K + d)T})$. It compares two upper bounds for a variant other than the one plotted, so it is a reference and not a prediction.

606 Sixteen studies, alongside Figure 2, which sweeps row overlap against delay in the three settings
607 the studies treat separately, P known, P estimated, and states coarsened. Studies one to thirteen
608 draw P from a Dirichlet and studies fourteen to sixteen do not throughout. STATE-EXP3 is beaten
609 in study fifteen, and in four of six cells of study sixteen when the baseline alone is grid-tuned.
610 Each study reports an observation, an interpretation, and a limitation. Studies 8 and 9 ask whether
611 the two theorems describe the right *dependence* rather than whether a bound holds, which is the
612 question a loose bound leaves open. Study 10 is a control for study 6 and reverses its conclusion,
613 and study 13 reports a rule that does not work. All policies run on the same environment parameters
614 and common random-number streams, so a seed fixes everything except which estimator is formed,
615 and differences are taken within a seed before averaging. Round r 's outcome is consumed before
616 round $r + d + 1$, matching Section 3. Reported $\pm$ is one standard error of the paired difference.
617 Throughout this appendix, *sample-path pseudo-regret* means $\sum_t c_t(A_t) - \min_a \sum_t c_t(a)$, evaluated
618 on one realization of the learner and environment randomness. Averaging it over seeds estimates the
619 expected pseudo-regret that the theorems bound. We use that name rather than “realized regret” to
620 distinguish it from regret computed from the sampled outcomes X_t , which is not what these studies
621 report.

622 **Does the bound hold, and does the estimator help?** With $|S| = 4$, $T = 10000$, $d = 50$, drift band
623 0.05 and 100 seeds:

K	state, $m = d + 1$	state, $m = 1$	action, $m = 1$	paired gain	state better in
8	455.1	368.2	387.6	19.3 ± 3.4	73/100
20	618.3	493.7	547.5	53.8 ± 5.6	87/100
40	734.3	574.5	696.1	121.5 ± 9.4	99/100
80	780.4	614.8	784.0	169.3 ± 9.9	100/100

625 Observed. Sample-path pseudo-regret for the analyzed variant stayed below $\sqrt{2(d+1)|\mathcal{S}|T \log K}$
626 in 100 of 100 seeds at every K , the bound running from 2913 to 4228 against realized 455 to 780.
627 The plain variant beat the action-level baseline at every K , by a margin growing from 19 to 169 as
628 K rose tenfold with $|\mathcal{S}|$ fixed, and it beat the analyzed variant throughout. Sweeping the delay at
629 $K = 40$ gave state-pooled 387, 569 and 731 against action-level 653, 699 and 773 at $d = 10, 50$
630 and 200.

631 Interpretation. The margin growing in K at fixed $|\mathcal{S}|$ is the prediction that separates the two estima-
632 tors, so the improvement is attributable to the state channel the effective dimension attributes to it.
633 Interleaving looks like a price of the analysis rather than of the estimator. The narrowing of the gap
634 as the delay grows is what an additive delay term predicts and a multiplicative one does not, which
635 is the diagnostic evidence behind Conjecture 5.

636 Limitation. This is one synthetic family with a fixed drift band, the plain variant carries no guarantee,
637 and a simulation cannot establish that a bound holds, only that sample-path pseudo-regret stayed
638 below it in the runs performed.

639 **What does an unknown map cost?** What follows evaluates a practical variant, the plug-in estimate
640 with $\hat{P}_t$ updated every round, no restart, $m = 1$ and $\alpha = 0.01$, and not the warm-start algorithm
641 of Theorem 10, which freezes $\hat{P}$ and restarts. Against the same algorithm given P and against the
642 action-level baseline, over 60 paired seeds and at mixing rate γ :

K	γ	known P	plug-in	action-level	plug-in gain
8	0.00	366.2	365.3	385.4	20.0 ± 4.3
20	0.00	480.3	477.9	529.0	51.2 ± 7.6
40	0.00	582.8	585.3	710.2	124.8 ± 11.2
80	0.00	625.2	623.1	802.2	179.1 ± 12.9
40	0.02	590.5	592.8	712.9	120.1 ± 10.4
40	0.10	618.9	623.5	735.5	112.0 ± 10.1

644 Observed. The plug-in landed within 0.5 per cent of known P at every K and beat action-level
645 weighting by 20 to 179 as K rose tenfold. Forced exploration did not help, $\gamma = 0$ being best at
646 every K . At $K = 40$ the sampled maps have $\kappa \approx 1.14$, so the exploration term $(\kappa|\mathcal{S}|)^{2/5}K^{1/5}T^{4/5}$
647 of Theorem 10 evaluates to about 6100 and its leading term to about 2700, against sample-path
648 pseudo-regret 585 and against $\sqrt{KT} = 632$.

649 Interpretation. Estimating P is nearly free on this family, and the gap between the analysis and the
650 behavior is an order of magnitude, so the analysis rather than the method looks conservative.

651 Limitation. The variant run here has no guarantee of its own, so the comparison is against a rate
652 proved for a different algorithm. That $\gamma = 0$ wins is a statement about this family rather than about
653 the algorithm, which the next study tests.

654 **Where does the plug-in fail?** The family above is mild. Five deliberately hostile maps at $K = 40$,
655 $|\mathcal{S}| = 8$, $T = 8000$, $d = 50$ and 30 paired seeds, with κ measured on each: a best action reachable
656 only through a state the others rarely produce, a sensor whose mass sits on one state, a state of
657 probability 10^{-3} everywhere, a deterministic map, and one whose good action's early visits are
658 unrepresentative.

environment	κ	known P	action-level	plug-in at mixing rate γ			
				0	0.01	0.05	0.20
rare optimal action	125.9	950.7	1302.3	948.2	976.8	1087.8	1503.0
unbalanced coverage	152.7	483.4	604.4	474.1	476.2	487.0	531.8
near-unreachable state	108.8	552.9	650.4	559.5	563.3	575.8	624.7
deterministic rows	2.5	556.3	733.7	571.9	585.9	653.9	904.3
misleading early visits	1.5	674.0	847.2	679.9	688.0	716.8	834.3

660 Observed. The plug-in beat action-level weighting on all five, including two with κ above 100, and
661 forced exploration cost regret in every cell. At the Laplace default $\alpha = 1$ the deterministic map is
662 the one case where the plug-in *loses* to action-level weighting, 750.3 against 733.7. Sweeping α

663 over 1, 0.3, 0.1, 0.03, 0.01 gives 750.3, 642.5, 601.4, 581.1, 571.9 there and is monotone on all five,
664 and $\alpha = 0.01$ is what the tables above use.

665 Interpretation. The mixing rate that Theorem 10 tunes is doing analytic work rather than algorithmic
666 work. The single loss has an identifiable mechanism, since one visit to a deterministic row is
667 recorded as $\hat{P}(s \mid a) = 2/(1 + |\mathcal{S}|)$ rather than 1, so rarely played actions are scored as if spread
668 over states they never reach.

669 Limitation. The leading rate is indifferent to fixed α and the algorithm is not, which is worth stat-
670 ing rather than tuning away. These five maps show that hostility to the plug-in is survivable at a
671 small constant, not that it is harmless, and no map adversarial to the smoothing constant itself was
672 constructed.

673 **Does safe blending recover the pooling gain, and if not why not?** The rule of Theorem 12, run
674 in the $m = 1$ variant like everything else in this appendix rather than in the analyzed $m = d + 1$
675 one, was compared against its two ingredients at $K = 40$, $|\mathcal{S}| = 8$, $d = 50$ and 20 paired seeds.
676 Measuring the realized bias against the certificate’s own budget locates the outcome. At $K = 40$,
677 $|\mathcal{S}| = 8$, $T = 8000$, $d = 50$ and 12 seeds the per-round budget $\frac{\eta}{2}(K - 2|\mathcal{S}|)$ sums to 307.3 over the
678 horizon, and

	summed	versus budget	quantity
679	128.7	0.42×	what the regret charges, $\sum_t \langle x_t, b_t \rangle - b_t(a^*) $
	513.7	1.67×	what Proposition 9 charges, $\sum_t \max_a b_t(a) $
	32000.0	104×	what the algorithm can verify, Hoeffding radii
	32000.0	104×	the same with Bernstein radii

680 Observed. The rule reproduced action-level regret to the digit in every cell, 533.2, 754.4 and 1317.0
681 at $T = 8000$ on the Dirichlet, deterministic and rare-optimal families and 2304.6, 2023.5 and 3108.8
682 at $T = 50000$, against a plug-in that scores 464.8, 585.5, 956.5 and 1770.6, 1473.5, 2346.6 on the
683 same runs. The fraction of round-action pairs it cleared for pooling was 0.000 everywhere, at both
684 horizons, at both smoothing constants, and with the confidence logs dropped as well.

685 Interpretation. The implemented certificate selected no pooling in any run, so the policy coincided
686 numerically with the action-level baseline rather than merely matching its guarantee, and the table
687 says where that comes from. On this trajectory the realized bias lies below the best-case variance
688 budget, so pooling would pay for itself here. The bound of Proposition 9 is 1.67 times the budget
689 with a factor of 4.0 available between it and the truth, so a tight bias bound would close the analytical
690 gap. What the algorithm can verify is saturated at the trivial cap of four per round, and Bernstein
691 radii buy nothing over Hoeffding ones, which is where the failure lives.

692 Limitation. Two exact refinements of the bias bound, an oscillation in place of a maximum and a
693 Cauchy–Schwarz split into the comparator’s effective dimension times a chi-square, give 577.9 and
694 5931.6 against a truth of 94.1 over $T = 4000$, where the maximum gives 368.0, so both are worse
695 than what they replace. The budget table is one diagnostic trajectory rather than a worst case, so it
696 shows this certificate inert here and not that no certificate can clear.

697 **Can the denominator be counted instead of predicted?** With x frozen over blocks of 400 and the
698 states of the first half counted, at $K = 40$, $|\mathcal{S}| = 8$, $T = 8000$, $d = 50$ and 12 seeds, the counted
699 denominator gives regret 496.5 against 499.5 for the predicted one and a comparator bias of 99.6
700 against 52.9, both under the budget of 153.7 for the pooled half.

701 Observed. Counting is the noisier of the two on regret and on bias. Computed under uniform play,
702 the predicted relative error is vacuous at every setting tried, $K = 40$ through $K = 4000$, while the
703 counted one reads 0.0244 against a budget of 0.0892 at $K = 1000$ and $T = 10^5$.

704 Interpretation. A prediction pools every action’s history while a count uses one block, which is why
705 counting is noisier, and yet the counted certificate is the better scaled of the two. That setting is the
706 first in which any certificate for ρ_t clears at all.

707 Limitation. The window that buys it is 53 per cent of the horizon, and the frozen-block regret at that
708 window, 1.2×10^7 , exceeds the 2.6×10^4 the pooling was meant to improve, so what clears is the
709 certificate and not a usable algorithm.

710 **Does the effective dimension, rather than $|\mathcal{S}|$, govern the gain?** Holding $|\mathcal{S}| = 8$, $K = 40$,
711 $T = 5000$, $d = 20$ and 24 seeds, and varying only the overlap between the rows of P , the measured
712 mean of v_t moves while $|\mathcal{S}|$ does not:

	row overlap	mean v_t	state-pooled	action-level	paired gain
	0.00	2.719	392.7	599.3	206.6 ± 20.0
713	0.25	2.013	361.1	528.1	167.0 ± 21.3
	0.50	1.532	306.5	406.5	100.0 ± 10.5
	0.75	1.168	199.9	229.2	29.3 ± 3.2
	0.95	1.009	49.2	50.8	1.6 ± 0.5

714 Observed. Both the regret and the advantage over action-level weighting fall monotonically with v_t ,
715 from a paired gain of 206.6 at $v_t = 2.719$ to 1.6 at $v_t = 1.009$. Nothing about the raw state count
716 changes across the sweep.

717 Interpretation. $|\mathcal{S}|$ is constant and cannot explain the pattern, while v_t moves with it, which is the
718 comparison Theorem 1 predicts.

719 Limitation. The two estimators do not become the same object at $v_t \approx 1$. Rather the rows of P agree,
720 so every action induces the same loss and both regrets collapse, making that endpoint uninformative
721 rather than favorable. One parameter of one map family was varied, so the sweep is consistent with
722 the effective dimension governing the gain rather than establishing it.

723 **Is there a bias-variance trade in the representation?** With 16 raw states drawn from four latent
724 clusters, $K = 40$, $T = 6000$, $d = 20$ and 24 seeds, running the algorithm on nested groupings
725 of size 1, 2, 4, 8, 16 gives, for the analyzed $m = d + 1$ variant, regret 1590, 831, 754, 920, 1080
726 at within-cluster heterogeneity 0, with mean v_t rising 1.000, 1.122, 1.435, 1.663, 2.089 and $\delta_t(g)$
727 falling 0.700, 0.275, 0, 0, 0.

728 Observed. The minimum sits at four groups, the number of latent clusters, and stays there at het-
729 erogeneity 0.12, where regret reads 1657, 885, 787, 960, 1113. Sample-path pseudo-regret stayed
730 below the bound of Theorem 3 in every cell. The bound's own optimum agrees at heterogeneity 0,
731 where it reads 9364, 4321, 1155, 1243, 1394 and also selects four groups, but not at 0.12, where it
732 reads 11128, 6073, 3430, 3231, 1385 and selects the finest representation.

733 Interpretation. A trade-off in the representation exists and is located at the latent structure, which is
734 what Theorem 3 describes qualitatively.

735 Limitation. The bias term $2 \sum_t \delta_t(g)$ is the conservative half of the trade, so the theorem certifies
736 a trade-off whose location it does not predict once groups are heterogeneous, and the latent cluster
737 count was known to the experimenter and not to any algorithm.

738 Separately, on the construction of Theorem 4 with $h = d$ the mean pseudo-regret over 400 seeds is
739 44.1 against a predicted fluctuation of 40.5, and of the five policies compared there the four without
740 access to the block signs attain expected loss within 0.05% of $T/2$, as Lemma 15 predicts, while the
741 one given them attains pseudo-regret -954. The code producing every reported number and figure
742 is at <https://github.com/melikabaghi/state-exp3>. It pins the software environment used
743 for the reported numbers, python 3.14.3 with numpy 2.4.2 and matplotlib 3.10.8, states the seed
744 count for each study, lists one command per numbered study, and commits the stored outputs each
745 figure reads so that every figure rebuilds without rerunning an experiment.

746 **Does the drift bound describe the right scaling?** The construction of Theorem 4 at $K = 40$,
747 $T = 8000$, $h = d$ and 12 paired seeds, sweeping d from 10 to 200, the amplitude ε from 0.02 to
748 0.20 so that E_2 runs from 12 to 1193, and the drifting coordinates J from 1 to 16. We run four
749 policies: uniform, action-level EXP3, STATE-EXP3, and the greedy policy on the exact stale vector
750 m_t , the last because Theorem 4 claims to survive it. E_2 is measured on the realized instance rather
751 than predicted, so the normalization uses the quantity the theorem uses.

752 Observed. The predicted scale $\sqrt{d E_2 \min\{1 + \log J, T/d\}}$ ranges from 38 to 377 across the grid
753 and the measured regret tracks it (Figure 3). Normalized, the four policies read 0.230 to 0.429,
754 0.228 to 0.424, 0.238 to 0.417 and 0.237 to 0.437, a spread below 1.9 against a tenfold range in
755 the predictor. The oracle-assisted policy has the largest mean ratio, 0.345 against 0.311, 0.314 and
756 0.310 for the three that never see m_t .




Figure 3: Study 8, the drift channel. **Left**, measured regret against the scale Theorem 4 predicts, over twelve cells in which d varies 20-fold, E_2 100-fold and J 16-fold. The dashed line has slope one. **Right**, the same points divided by that scale, grouped by which parameter the cell sweeps, with the shaded band spanning the observed range. STATE-EXP3 appears here only as one more algorithm on the hard family, where it has no advantage.




Figure 4: Study 9, the state channel. The plotted ratio would be flat if Theorem 1 captured the dependence on all three of T , d and $\hat{v}$. **Left**, horizon sweep at the largest effective dimension, one line per delay. **Right**, every cell with $\hat{v} > 1.05$ against rounds per interleaved copy, where the three delays fall on one curve.

757 Interpretation. The three arguments enter the predicted scale differently and the collapse holds in all
 758 three sweeps, so it is the form of the bound rather than one fitted constant that tracks the data. That
 759 greedy play on the exact m_t sits at the top of the band rather than below it is the concrete content of
 760 the claim that the bound survives the oracle.

761 Limitation. Every cell has $T/d > 1 + \log J$, so only the $1 + \log J$ branch of the minimum is
 762 exercised and the T/d saturation branch, which needs d comparable to T , is not probed. A lower
 763 bound quantifies over all algorithms and four policies cannot establish that. What these runs check
 764 is the scaling of the construction.

765 **Does Theorem 1 describe the right dependence on T , d and $\hat{v}$?** A $4 \times 3 \times 5$ grid in row overlap,
 766 delay and horizon at $K = 40$, $|\mathcal{S}| = 8$ and five paired seeds, running the analyzed $m = d + 1$ variant
 767 and reporting $R / \sqrt{(d+1) \hat{v} T \log K}$. A raw regret table cannot answer this and neither can a bound
 768 that holds with room to spare, since the question is about dependence rather than level.

769 Observed. Over the full grid the ratio spans 35.9 times, from 0.016 to 0.579. Dropping the near-
 770 degenerate column at $\hat{v} = 1.013$ leaves 7.1 times over 45 cells. At $\hat{v} = 3.656$ and $d = 10$ it moves
 771 only 1.35 times across a sixteenfold range of T , and at $d = 200$ it moves 2.57 times. Against rounds
 772 per interleaved copy $T/(d+1)$ the three delays fall on one increasing curve (Figure 4).

773 Interpretation. The data are consistent with the predicted $\sqrt{T}$ dependence once each interleaved
 774 copy receives enough rounds, while much of the remaining variation is associated with the number
 775 of rounds available to each copy. At $d = 200$ and $T = 2500$ each of the 201 copies sees 12 rounds,
 776 so that cell is nowhere near the regime the bound describes. Together with the $m = 1$ comparison in
 777 Study 1, this suggests that much of the observed delay penalty comes from the interleaving device
 778 used in the analysis rather than from the pooled estimator itself.




Figure 5: **Left**, study 10. The same overlap sweep with and without holding the task difficulty fixed. Matching the uniform-play regret reverses the trend, supporting the interpretation that the decline in Study 6 is driven by task-difficulty confounding rather than by the pooling mechanism itself. Bars are one standard error of the paired difference. **Right**, study 11, regret saved against action-level weighting over the smoothing pseudocount and the sensor balance, 10 paired seeds per cell.

779 Limitation. The degenerate column is degenerate for the reason study 6 gives, since at $\hat{v} \approx 1$ the
780 rows of P agree, every action induces nearly the same loss, and both regrets collapse for a reason
781 the bound cannot see. One natural explanation for the remaining looseness is ruled out rather than
782 confirmed, since the realized $\sum_t v_t$ lies within 1.01 to 1.49 of the budget $\hat{v}T$, so the sup over play
783 distributions is not the source.

784 **Does pooling help because the rows overlap, or because overlap made the task easy?** Study 6
785 sweeps the overlap and finds the advantage falling as $\hat{v}$ falls, which read at face value contradicts
786 Theorem 1, whose bound improves as $\bar{v}$ falls. That sweep is confounded, since raising the overlap
787 also flattens $c_t = P\theta_t$ across actions, collapsing the comparator gap so that both regrets vanish
788 together. This study holds $K = 40$, $|\mathcal{S}| = 8$, $T = 10000$, $d = 50$ and the difficulty fixed, and
789 varies only $\hat{v}$. Difficulty is the uniform-play regret $\sum_t \text{mean}_a c_t(a) - \min_a \sum_t c_t(a)$, a property of
790 the instance with no algorithm in it, and the amplitude of θ is bisected in every cell to hit a common
791 target. The map keeps one contrast direction alive at every overlap and θ is aligned with it, which is
792 what lets the gap survive as the rows agree.

793 Observed. Across 20 paired seeds the relative gain rises from 11.7 per cent at $\hat{v} = 3.380$ to 19.3 per
794 cent at $\hat{v} = 1.120$, with absolute gains 85.5 ± 8.9 to 106.2 ± 7.9 and a correlation of -0.929 between
795 $\hat{v}$ and the gain. Study 6’s unmatched sweep falls from 34.5 to 3.1 per cent over a comparable range
796 (Figure 5).

797 Interpretation. The two sweeps disagree in sign, and the matched-difficulty sweep agrees with the
798 overlap dependence suggested by Theorem 1, although this experiment uses the practical $m = 1$
799 variant rather than the theorem’s $m = d + 1$ algorithm. Overlap is what pooling exploits, and
800 study 6 measures overlap confounded with an instance becoming trivial. This also explains study 9’s
801 degenerate column without appealing to it.

802 Limitation. At the two highest overlaps some seeds cannot reach the common target and the realized
803 difficulty is 5 to 13 per cent below it, so the matching is close rather than exact. Since a smaller
804 gap should if anything shrink the advantage, the measured rise is conservative in that direction. The
805 matched gains are also smaller in magnitude than the unmatched ones, which is what fixing the
806 difficulty costs.

807 **Does the matched-difficulty trend survive a second configuration?** Study 10 is a single design
808 point, so the same procedure was re-run at $K = 20$, $|\mathcal{S}| = 12$, $T = 8000$, $d = 20$ and 20 paired
809 seeds, with the common target lowered to 500 so that it is reachable at every overlap. Nothing else
810 changes.




Figure 6: Study 12. **Left**, the three measured arms and the two certificates on one axis, $|\mathcal{S}| = 8$, $T = 5000$, $d = 50$, 12 paired seeds, every measured arm at $m = 1$. The shaded region is what the analysis costs rather than what the method does. **Right**, the same measured arms rescaled with the certificates removed, which is the only way to see that known P and the plug-in coincide. Bars are one standard error.

overlap	$\hat{v}$	R_{unif}	state	action	paired gain
0.00	3.440	500.0	364.7	407.2	42.5 ± 8.2
0.35	2.225	500.0	346.8	401.3	54.6 ± 8.3
0.65	1.518	500.0	309.3	372.4	63.2 ± 6.4
0.85	1.240	493.6	271.8	346.0	74.3 ± 7.2
0.95	1.191	490.1	256.9	327.6	70.7 ± 7.8

Observed. The relative gain rises from 10.4 to 21.6 per cent as $\hat{v}$ falls from 3.440 to 1.191, and the correlation between $\hat{v}$ and the absolute gain is -0.972 against study 10's -0.929 . Realized difficulty spans 490.1 to 500.0, a 1.020 times spread, so the matching is tighter here than in study 10.

Interpretation. The direction study 10 reports is not an artifact of its single configuration. Two independent design points, differing in K , $|\mathcal{S}|$, T and d , give the same sign and a comparable magnitude.

Limitation. Both configurations use the same construction and the same matching routine, so this tests configuration sensitivity and not construction sensitivity. Both run the practical $m = 1$ variant. An earlier attempt at this replication silently failed because `match_scale` binds its target as a default argument, so the bisection saturated and the difficulty was never matched; the run reported here passes the target explicitly.

Is the plug-in fragile in the smoothing constant, and does that depend on coverage? Study 3 reports one map where the plug-in loses at $\alpha = 1$ and recovers as α falls, which is an anecdote about a single instance. Here α is swept over 1 to 0.003 against the sensor balance κ of Theorem 10, controlled from 1.2 to 76 by placing mass ρ on one scarce state and drawing rows around that mean, at $K = 40$, $|\mathcal{S}| = 8$, $T = 8000$, $d = 50$ and 10 paired seeds.

Observed. No cell loses to action-level weighting, the gain running from 28.3 to 66.9 over the 30 cells. The gain roughly doubles as α falls from 1 to 0.03 and is flat below that, while moving along κ at fixed α changes it far less, the worst-to-best plug-in regret within a κ column being 1.07. At the better α the plug-in is indistinguishable from known P , the paired difference clearing two standard errors in one of the five columns and changing sign in another.

Interpretation. Of the two parameters the analysis names, the one that matters here is the one the analysis treats as free. A Laplace default costs about half the available gain at every coverage level tried, and coverage imbalance over a sixtyfold range costs much less than that. Taken with study 2 this is further evidence that the difficulty in the unknown- P setting is certification rather than estimation.

Limitation. This is one map family with κ varied through a single scarce state, so it does not rule out a construction in which κ is the binding parameter, and Theorem 10 still carries κ for a worst case this family need not contain. The flatness below $\alpha = 0.03$ says the leading behavior stops moving, not that any fixed α is safe at every horizon.

842 **How far is the method from its certificate, and the certificate from useful?** Study 2 reports these
843 numbers in separate columns, leaving the reader to assemble them. Here action-level weighting,
844 STATE-EXP3 with P given, and the same estimator with $\hat{P}_t$ updated every round are measured on
845 the same paired seeds across K from 8 to 160, and drawn beside the two certificates. Theorem 1 is
846 the certificate for known P and analyzes $m = d + 1$, while Theorem 10 is the warm-start rate and
847 analyzes an algorithm that freezes $\hat{P}$ and restarts, so both are plotted to show the size of a gap and
848 not as like for like.

849 **Observed.** The plug-in stays within 1.6 per cent of known P at every K , the two curves being
850 visually indistinguishable, while action-level weighting costs 1.4 to 17.4 per cent more than the
851 plug-in as K rises twentyfold. Theorem 1 sits 10.7 to 14.8 times above the measured plug-in and
852 Theorem 10 sits 14.0 to 18.7 times above it. The warm-start certificate also stays above $\sqrt{KT}$ at
853 every K tried, by factors of 19.2 down to 6.6.

854 **Interpretation.** The vertical gap reflects analytical conservatism rather than observed algorithmic
855 performance. Estimating P online is close to free on this family, so the order of magnitude between
856 the measured arms and either certificate is conservatism in the proof. That the warm-start rate never
857 drops below $\sqrt{KT}$ here is the concrete form of the statement in Section 4 that it does not certify the
858 channel it analyzes.

859 **Limitation.** The measured arms run $m = 1$ and Theorem 1 analyzes $m = d + 1$, and the plug-in
860 is not the warm-start algorithm Theorem 10 analyzes, so neither vertical gap is a statement about a
861 bound being loose for the algorithm it covers. A conservative bound and a well-behaved algorithm
862 are consistent with a third possibility, that the family is easy.

863 **Can the coarsening be chosen from data?** Theorem 3 takes the representation as given and study 7
864 uses the latent cluster count, which no algorithm knows. This study supplies a rule and measures
865 what it costs, with no theorem behind it. A prefix of 1200 rounds runs the identity representation
866 and estimates $\hat{\theta}(s)$ by averaging the outcomes seen at each state. Sorting states by $\hat{\theta}$ and merging
867 adjacent pairs greedily gives one candidate at every granularity from 16 down to 1, each scored
868 by the quantity Theorem 3 bounds, $\sqrt{2(d+1)V(g)\log K} + 2\sum_t \hat{\delta}_t(g)$, and the argmin runs the
869 remaining rounds. Against it are the identity, the true clusters, and the best candidate on the ladder
870 scored in hindsight, over 24 paired seeds.

871 **Observed.** The rule fails. It selects 15 groups at heterogeneity 0 and 16 at 0.12, against a true cluster
872 count of four, and lands at 912.9 ± 30.1 and 944.7 ± 31.1 where the identity gives 921.0 ± 30.5 and
873 952.7 ± 32.4 , so it captures under one per cent of the available gain. The true clusters give $843.7 \pm$
874 27.9 and 883.0 ± 31.1 , and hindsight over the same ladder gives 833.5 ± 27.8 and 841.5 ± 23.6 .

875 **Interpretation.** Hindsight nearly matches the true clusters, so the ladder does contain a good repre-
876 sentation and the failure is the criterion rather than the candidate set. Measuring $\hat{\delta}$ at the four-group
877 representation locates it. At heterogeneity 0 the true within-group spread is exactly zero while the
878 prefix estimates it at 0.154, which is sampling noise, and the bound multiplies that by $2(T - T_0)$ to
879 give a phantom bias of 1483 against a real variance saving of 1261. The bias term is the conservative
880 half of Theorem 3, and a noisy plug-in estimate of it is enough to forbid every merge.

881 **Limitation.** One selector on one family, and a negative result about this rule is not a negative result
882 about data-driven selection. The obvious repairs, debiasing $\hat{\delta}$ for its sampling noise or scoring
883 candidates on held-out loss instead of on the bound, are not evaluated here. The prefix is also charged
884 to every arm equally, so the comparison understates what a selector that spent less on estimation
885 might do.

886 **Is “many actions resolved by few states” a natural regime?** Every study so far draws P from
887 a Dirichlet, which is the assumption under test rather than evidence for it. This one builds an
888 environment to the shape of a recommendation funnel instead. Actions are $K = 200$ catalogue
889 items, states are $|\mathcal{S}| = 6$ ordered engagement depths, and the outcome is a delayed conversion.
890 Three structural facts are imposed because funnels have them, namely that items fall into a modest
891 number of categories sharing an engagement profile, that item quality is heavy-tailed, and that the
892 loss falls with engagement depth and drifts seasonally rather than adversarially. The generated
893 catalogue puts 74 per cent of its mass on no engagement and 1.4 per cent on the deepest.




Figure 7: Study 14, a recommendation funnel. **Left**, the effective dimension stays near two while the catalogue grows thirty-two-fold, so the gap between K and $\hat{v}$ comes from the catalogue having categories, not from a chosen K . **Right**, the three arms at $T = 20000$ and 8 paired seeds, labelled by whether the regime condition of Section 4 holds. Bars are one standard error. No real data is used anywhere in this study.




Figure 8: Study 15. **Left**, regret relative to the best arm in each family, so 1 is that column’s winner, over five policies, nine settings and 10 paired seeds. The best-state rule leads seven columns and STATE-EXP3 leads two. **Right**, growth in T on the one family built so the rule’s assumption is false, with fitted exponents. The left panel is a worst case over families chosen here; the right panel is why that is not a worst case over the model.

894 Observed. The effective dimension is $\hat{v} = 1.97$ against $K = 200$, a hundredfold gap, and it stays
895 between 1.41 and 1.81 as the catalogue grows from 25 to 800 items (Figure 7). Against action-level
896 weighting the state-pooled variant saves 79.3, 63.6 and 38.9 per cent at $d = 10, 50$ and 200, the
897 largest margins anywhere in this appendix. The plug-in matches known P at every delay, the paired
898 difference never reaching two standard errors. The regime condition $\hat{v}(d+1) \log K \leq K + d$ holds
899 at $d = 10$ and fails at $d = 50$ and $d = 200$.

900 Interpretation. The separation between K and $\hat{v}$ arrives from catalogue structure rather than from
901 a tuned parameter, and it does not close as the catalogue grows, which is the claim the Dirichlet
902 families cannot make. That the condition Theorem 1 needs fails at the delays a funnel has, while
903 the practical variant still saves most of the regret, is the same message as study 1 in the setting that
904 motivates the paper. It is evidence for Conjecture 5 rather than for the theorem.

905 Limitation. *No real data is used.* There is no interaction log in this environment, nothing here
906 is fitted to one, and this is not a semi-synthetic benchmark in the usual sense of that phrase. The
907 structure is imposed by hand from qualitative properties, so it shows that the regime is describable
908 and self-consistent, not that any deployed system occupies it. The arms run $m = 1$, which carries no
909 guarantee, and the environment has no adversarial component, so it exercises the pooling channel
910 and not the drift channel.

911 **Which assumption produces which benefit?** Comparing against action-level weighting alone
912 leaves open whether the gain comes from the state structure or from an implementation detail. We
913 run five policies on the same paired seeds over nine settings. Besides action-level weighting and

STATE-EXP3 they are EXP3-IX (Neu, 2015), the implicit-exploration variant of EXP3, an importance weight; *naive sharing*, which credits $P(S_t | a)X_t/\bar{p}(S_t)$ using the unconditional state distribution and so isolates the value of the denominator $q_t = x_t^\top P$ from the value of sharing at all; and a *best-state rule*, which estimates $\hat{\theta}(s)$ online, takes the best state, and plays the action most likely to reach it. Two settings are added to the earlier families, one where the best action mixes over several good states rather than maximizing the single best one, and one that is obviously non-stationary. Running a bandit over states and mapping the values back through P is not a separate baseline, since $\sum_s P(s | a)\mathbf{1}\{S_t = s\}X_t/q_t(s)$ equals $P(S_t | a)X_t/q_t(S_t)$ identically, so that method is STATE-EXP3.

Observed. STATE-EXP3 does not win most columns. The best-state rule leads seven of the nine and STATE-EXP3 leads two, and over these nine settings the rule is also better on the median, 1.00 against 1.47 times the column winner, and on the worst case, 3.08 against 3.39 (Figure 8). EXP3-IX is 1.5 to 17.1 per cent worse than plain action-level weighting everywhere, and naive sharing ranges from 18.7 per cent better to 22.7 per cent worse than it. On the spread family the ordering inverts, STATE-EXP3 leading by 47.6 and 22.2 per cent while the best-state rule falls 61.5 and 43.0 per cent behind action-level weighting. Sweeping the horizon there, the rule’s regret grows with fitted exponent 1.00 and STATE-EXP3’s with 0.58.

Interpretation. The best-state rule buys its lead with an assumption, that the best action is the one most likely to reach the single best state. Where that holds it is very strong, and in the funnel it holds in 10 of 10 seeds because engagement depth is ordered, which is why it leads there by a factor of two. Where it fails the rule commits to a fixed action and pays regret linear in T , which no exploration rate repairs, and the exponent of 1.00 is that happening. The comparison therefore separates the assumptions rather than ranking the methods. STATE-EXP3 assumes nothing about the shape of θ . Across these nine selected settings its largest regret ratio to the column winner is 3.39, and the rule’s suite maximum is slightly smaller at 3.08, but the rule’s fitted growth on the spread family is linear. Action-level weighting also carries a sublinear guarantee, so what separates the arms here is the constant rather than the rate. STATE-EXP3’s paired standard errors are the smallest of the five at 5.8 per cent of the mean against 23.6 for the rule. That naive sharing can be worse than no sharing is the separate point that the denominator, not the sharing, is what Proposition 7 buys.

Limitation. Nine settings chosen here are not a worst case over the model, and the median and worst-case columns above should be read as describing this set only. A practitioner whose problem has the ordered structure of the funnel should use the best-state rule, and this appendix does not identify when that structure can be verified in advance. On the drift family no arm’s paired difference from action-level weighting reaches two standard errors, the largest being the best-state rule’s $+24.6 \pm 12.4$, so that column is reported as an absence of signal rather than a result.

Does pooling still help against the strongest published delayed-bandit algorithm? Every comparison so far is against action-level weighting, EXP3-IX, naive sharing, or the best-state rule. None is the strongest published method for delayed feedback: that is the hybrid Tsallis-plus-entropy FTRL of Zimmert and Seldin (2020), whose delay term is additive rather than multiplicative. We run it on both headline families, with the same environments, seeds and delay convention, under three tunings.

family, cell	STATE-EXP3	Zimmert and Seldin	same, grid-best	STATE-EXP3, grid-best
funnel, $d = 10$	949.0	2959.4	2413.8	27.0
funnel, $d = 50$	1728.1	3557.0	2556.4	95.9
funnel, $d = 200$	3244.8	4734.7	2807.7	199.7
overlap 0.00	661.5	731.8	610.4	202.4
overlap 0.65	584.1	673.1	544.0	51.2
overlap 0.95	461.4	540.2	419.4	34.3

Observed. With STATE-EXP3 at the single tuning used everywhere else here and Zimmert and Seldin (2020) at the best of a six-point grid around its own value, STATE-EXP3 wins every cell, cutting regret by 31.5 to 67.9 per cent on the funnel and 9.6 to 14.6 per cent on the matched sweep. With both at their grid-best, over seven multipliers for Zimmert and Seldin (2020) and six for STATE-EXP3, STATE-EXP3 wins every cell again and by more, 66.8 to 98.9 per cent. With Zimmert and Seldin (2020) at its grid-best and STATE-EXP3 left at its paper tuning, the third column against the first, Zimmert and Seldin (2020) wins the funnel at $d = 200$ and all three matched cells, four of six.

Interpretation. No column is a symmetric protocol. The two grids are finite and unequal, and STATE-EXP3's grid-best sits at its top multiplier in every cell, so its own minimum may lie outside the range searched; the truncation runs against STATE-EXP3. Where it is not held to the worse tuning of the two, the state channel helps against the strongest published method, which is expected, since that method is action-level by construction and cannot read the state at all. Grid tuning helps STATE-EXP3 far more than it helps Zimmert and Seldin (2020), which is consistent with pooling cutting estimator variance and lower variance admitting a larger learning rate.

Limitation. The third column against the first is a real loss and is reported as one. These are our own synthetic families, and a method built for adversarial worst cases is not shown at its best on benign instances. Every grid-best multiplier is chosen by reading these benchmarks, so none of the last three columns is a procedure a practitioner could run; only the first is. Eight paired seeds on the funnel and ten on the matched sweep.

Taken together the sixteen studies cover both channels. Studies one, six, seven, nine, ten, fourteen and fifteen probe the state channel. They are consistent with the effective dimension, rather than $|\mathcal{S}|$ or K , governing what pooling buys, with the qualification that study ten had to control for task difficulty before study six's sweep pointed the way Theorem 1 does, and study nine shows where the theorem stops describing the dependence and attributes that to the interleaving. Studies two to five, eleven and twelve probe what happens when the action-to-state matrix must be estimated, and locate the difficulty in certifying the estimate rather than in forming it, study eleven adding that the smoothing constant matters more than the coverage imbalance the analysis carries and study twelve putting the measured arms and their certificates on one axis. Study thirteen closes the representation thread by showing that the bound of Theorem 3 is too conservative to serve as its own selection rule. Study fourteen is the only one whose action-to-state matrix is not drawn from a Dirichlet, and it reports the largest margins here, though on a structure imposed by hand rather than measured. Study fifteen widens the baseline set and finds STATE-EXP3 beaten in seven of nine settings by a heuristic that assumes the best action is the one reaching the single best state, which holds on several of these families and fails with linear regret when it does not. Study sixteen places STATE-EXP3 against the strongest published delayed-bandit method. It is ahead in every cell when both are tuned alike, and behind in four of the six when that method alone is grid-tuned. Study eight is the drift channel, and its collapse across d , E_2 and J is the strongest evidence here that Theorem 4 has the right form. No study runs a method that uses both channels, which is the gap Section 8 describes.

J Why the delay does not become additive

Run STATE-EXP3 with $m = 1$, so $x_t \propto \exp(-\eta L_t)$ with $L_t = \sum_{r \leq t-d-1} \hat{c}_r$. Let $\tilde{x}_r$ be the iterate that has also absorbed the outstanding estimates $G_r = \sum_{s=r-d}^{r-1} \hat{c}_s \geq 0$. Two terms appear.

The delay term is $\mathbb{E}[\langle x_r - \tilde{x}_r, \hat{c}_r \rangle]$. Since $\tilde{x}_r(a) \geq x_r(a)e^{-\eta G_r(a)}$ and $1 - e^{-z} \leq z$, it is at most $\eta \mathbb{E}[\sum_a x_r(a) G_r(a) \hat{c}_r(a)] = \eta \sum_{s=r-d}^{r-1} \mathbb{E}[\langle x_r, \hat{c}_s \rangle]$, and x_r depends only on rounds $r-d-1$ and earlier, hence on rounds before s , so $\mathbb{E}[\langle x_r, \hat{c}_s \rangle] = \mathbb{E}[\langle x_r, c_s \rangle] \leq 1$. This term is $\eta d T$ for any unbiased estimator, which is the additive form.

The quadratic term is where it stops. The telescoping scores $\sum_a x_{r+d}(a) \hat{c}_r(a)^2$, with x_{r+d} the iterate that has absorbed rounds up to $r-1$, while Proposition 7 bounds the play-weighted sum, the one against x_r , by v_r . Carrying out the substitution, $\mathbb{E}[\sum_a x_{r+d}(a) \hat{c}_r(a)^2 \mid \cdot] \leq \sum_s q_{r+d}(s)/q_r(s) \leq |\mathcal{S}|/Z_r$ for $Z_r = \mathbb{E}_{a \sim x_r}[e^{-\eta G_r(a)}]$, and $Z_r \geq e^{-\eta \langle x_r, G_r \rangle}$, so the swap costs an exponential moment of $\langle x_r, G_r \rangle$.

No bound on that moment holds uniformly over state distributions. For this estimator $\langle x_r, \hat{c}_s \rangle = q_r(S_s)X_s/q_s(S_s)$, so the quantity is governed by $1/q_s(S_s)$, the reciprocal probability of the state actually observed. Its mean is $\sum_s q_s(s)/q_s(s) = |\mathcal{S}|$, which is why the delay term is fine. For a fixed q of full support the exponential moment is finite, $1/q(S)$ being bounded, and what fails is uniformity. Already at $|\mathcal{S}| = 2$ with $q = (\varepsilon, 1 - \varepsilon)$, $\mathbb{E}[e^{\eta/q(S)}] \geq \varepsilon e^{\eta/\varepsilon} \rightarrow \infty$ as $\varepsilon \rightarrow 0$ while the mean stays at 2, so no bound in $|\mathcal{S}|$ and η alone is available, and q_s is set by the algorithm's own past play. Measured at $|\mathcal{S}| = 4$, the mean stays at 4.0 while $\mathbb{E}[e^{0.05/q(S)}]$ runs 1.22, 1.43 and then 1.3×10^{21} as the smallest state probability falls through 0.25, 0.019 and 0.0009.

None of this is needed: Appendix K controls the same swap without any moment of $1/Z_r$, by bounding the state-marginal ratio directly. The rest of this section records why the route through $1/Z_r$ fails, since it is the natural first attempt.

Flooring q restores the moment and costs the product back. Mixing with the uniform action at rate γ gives $q_t(s) \geq \gamma \bar{p}(s)$, so $\eta \langle x_r, G_r \rangle \leq 1$ needs $\gamma \asymp \eta d \kappa |\mathcal{S}|$ under the balanced-sensor condition $\bar{p}(s) \geq 1/(\kappa |\mathcal{S}|)$, and the mixing cost γT then reproduces $\tilde{O}(\sqrt{d \kappa |\mathcal{S}| T \log K})$. So the obstruction is not an artifact of one device. Empirically it does not bite, $1/Z_r$ staying below 1.33 at the tuned η over $d \in \{10, 50, 200\}$, which is what suggested that the obstruction was an artifact of the device rather than of the problem, as Appendix K confirms.

K Additive delay without interleaving

Interleaving cannot give an additive bound. Copy i , which absorbs a share V_i of the budget $V = \sum_i V_i$, has regret $\sqrt{2V_i \log K}$, and Cauchy–Schwarz over the m copies gives $\sqrt{2mV \log K}$, tight when the V_i agree, so the $\sqrt{m}$ in Theorem 1 is not an artifact of the analysis. An additive bound has to come from a single copy.

Throughout, $L_t = \sum_{r \leq t-d-1} \hat{c}_r$ and $x_t \propto e^{-\eta L_t}$, so $L_{t+1} - L_t = \hat{c}_{t-d}$, with $\hat{c}_r = 0$ for $r < 1$. The filtration $\mathcal{F}_t$ is the environment’s, including outcomes not yet delivered, as in Proposition 7; under it x_t is $\mathcal{F}_{t-d-1}$ -measurable, so x_{r+d} is $\mathcal{F}_{r-1}$ -measurable.

The drift lemma. Appendix J priced the change of measure from x_r to x_{r+d} through $1/Z_r$. That is avoidable. The pooled estimate satisfies

$$\langle x_t, \hat{c}_t \rangle = \sum_a x_t(a) \frac{P(S_t | a) X_t}{q_t(S_t)} = X_t \leq 1$$

identically, on every path, because the denominator is the state marginal of the same distribution that supplies the numerator.

Lemma 18. *If $\eta \leq 1/(e(d+1))$ then $x_{t+1}(a) \leq (1 + 1/d)x_t(a)$ for every t and every a , surely, and hence $x_{t+d}(a) \leq e x_t(a)$.*

Proof. Strong induction on t . For $t \leq d$ nothing is delivered and $x_{t+1} = x_t$. For $t > d$, assume the claim for every $s < t$ and compose it over the d steps from $t-d$ to t , giving $x_t(a) \leq (1 + 1/d)^d x_{t-d}(a) \leq e x_{t-d}(a)$. Contracting with the fixed nonnegative vector $\hat{c}_{t-d}$ and using the identity above, $\langle x_t, \hat{c}_{t-d} \rangle \leq e \langle x_{t-d}, \hat{c}_{t-d} \rangle = e X_{t-d} \leq e$. With $Z_t = \sum_b x_t(b) e^{-\eta \hat{c}_{t-d}(b)}$ and $e^{-z} \geq 1 - z$, this gives $Z_t \geq 1 - \eta \langle x_t, \hat{c}_{t-d} \rangle \geq 1 - \eta e \geq d/(d+1)$, so $x_{t+1}(a) = x_t(a) e^{-\eta \hat{c}_{t-d}(a)} / Z_t \leq x_t(a) / Z_t \leq (1 + 1/d)x_t(a)$. $\square$

What the induction bounds is $\langle x_t, \hat{c}_{t-d} \rangle = X_{t-d} q_t(S_{t-d}) / q_{t-d}(S_{t-d})$, the state-marginal ratio Appendix J could not reach. Individual $\hat{c}_r(a)$ remain unbounded; only their x_t -weighted mass is controlled. Cesa-Bianchi et al. (2019) prove the action-level version under $\eta \leq 1/(Ke(d+1))$ and Thune et al. (2019) remove the K , using that their estimate is carried by the single action played, so the weight cancels into a same-action ratio. A pooled estimate is spread over every action and does not collapse that way; the identity above replaces the cancellation.

Centring. Let $\zeta_t = \hat{c}_t - X_t \mathbf{1}$, with $\zeta_r = 0$ for $r < 1$. Exponential weights is unchanged by a shift that is constant across actions, so x_t is the same iterate and this is a change of analysis, not of algorithm (Eldowa et al., 2024). Keeping $\gamma(s) = \mathbb{E}[X_t^2 | S_t = s]$ inside the state sum,

$$\mathbb{E}[\langle x_t, \zeta_t^2 \rangle | \mathcal{F}_{t-1}] = \sum_s \gamma(s) \left[\frac{\sum_a x_t(a) P(s | a)^2}{q_t(s)} - q_t(s) \right] \leq v_t - 1,$$

each bracket being nonnegative by Jensen and $\gamma(s) \leq 1$, the final step being Remark 6. Bounding $\langle x_t, \zeta_t^2 \rangle$ by v_t first and then subtracting does not work: it leaves $v_t - \mathbb{E}[X_t^2]$.

Proof of Theorem 2. Since $\zeta_r(a) \geq -X_r \geq -1$ we have $\eta \zeta_r(a) \geq -\eta$, and Lagrange’s remainder gives $e^{-z} \leq 1 - z + (e^\eta/2)z^2$ for $z \geq -\eta$. With $\log(1+y) \leq y$, the potential

1056 $W_t = \sum_a \exp(-\eta \sum_{r \leq t-d-1} \zeta_r(a))$ obeys $\log(W_{t+1}/W_t) \leq -\eta \langle x_t, \zeta_{t-d} \rangle + (\eta^2 e^\eta / 2) \langle x_t, \zeta_{t-d}^2 \rangle$.
 1057 Telescoping from $W_1 = K$ and using $\log W_{T+1} \geq -\eta \sum_{r \leq T-d} \zeta_r(a^*)$ for the comparator a^* ,

$$\sum_{r=1}^{T-d} \langle x_{r+d}, \zeta_r \rangle \leq \sum_{r=1}^{T-d} \zeta_r(a^*) + \frac{\log K}{\eta} + \frac{\eta e^\eta}{2} \sum_{r=1}^{T-d} \langle x_{r+d}, \zeta_r^2 \rangle.$$

1058 The shift cancels between the two sides. As x_r is $\mathcal{F}_{r-d-1}$ -measurable, $\mathbb{E}[\langle x_r, \hat{c}_r \rangle] = \mathbb{E}[\langle x_r, c_r \rangle]$
 1059 and $\mathbb{E}[\hat{c}_r(a^*)] = c_r(a^*)$, and the d rounds undelivered by $T+1$ contribute at most d , so

$$\mathbb{E}[R_T] \leq \frac{\log K}{\eta} + \frac{\eta e^\eta}{2} \sum_r \mathbb{E}[\langle x_{r+d}, \zeta_r^2 \rangle] + \sum_r \mathbb{E}[\langle x_r - x_{r+d}, \hat{c}_r \rangle] + d.$$

1060 Lemma 18 gives $x_{r+d} \leq e x_r$ surely, and x_{r+d} is $\mathcal{F}_{r-1}$ -measurable, so the second sum is at most
 1061 $e \sum_r \mathbb{E}[v_r - 1] \leq eV^-$.

1062 For the third, write $x_{r+d}(a) = x_r(a)e^{-\eta G_r(a)}/Z'$ with $G_r = \sum_{u=r-d}^{r-1} \hat{c}_u \geq 0$ and normaliser
 1063 $Z' \leq 1$, distinct from the one-step Z_t of the Lemma. Then $x_r(a) - x_{r+d}(a) \leq x_r(a)(1 -$
 1064 $e^{-\eta G_r(a)}) \leq \eta x_r(a) G_r(a)$ for every a , trivially where the left side is negative. Since $\hat{c}_r \geq 0$ and
 1065 x_r, G_r are $\mathcal{F}_{r-1}$ -measurable, the conditional expectation is at most $\eta \langle x_r, G_r \rangle$, and $\mathbb{E}[\langle x_r, G_r \rangle] =$
 1066 $\sum_{u=r-d}^{r-1} \mathbb{E}[\langle x_r, \hat{c}_u \rangle] \leq d$: for each such u , x_r is $\mathcal{F}_{r-d-1}$ -measurable and $r-d-1 \leq u-1$, so
 1067 it leaves the conditional expectation and $\mathbb{E}[\langle x_r, \hat{c}_u \rangle] = \langle x_r, c_u \rangle \leq 1$. This is measurability, not
 1068 independence, and at $u = r-d$ the two σ -fields coincide, so the step rests on the timing convention
 1069 of Section 3. The sum is at most ηdT .

1070 Finally $\eta \leq 1/(2e)$ gives $e^{1+\eta} < 3.3$, so $\mathbb{E}[R_T] \leq \log K/\eta + \frac{\eta}{2}(3.3V^- + 2dT) + d$. For $f(\eta) =$
 1071 $A/\eta + B\eta$ restricted to $\eta \leq \eta_{\max}$ one has $\min f \leq 2\sqrt{AB} + A/\eta_{\max}$, in the interior and in the
 1072 binding case alike; with $A = \log K$, $B = (3.3V^- + 2dT)/2$ and $\eta_{\max} = 1/(e(d+1))$ this is the
 1073 stated bound. $\square$

1074 The cap on η binds only when $d \gtrsim T/(e^2 \log K)$, and not at the settings of Appendix I: the funnel
 1075 study's own rates are 0.0058, 0.0031 and 0.0016 at $d = 10, 50, 200$ against caps 0.0334, 0.0072 and
 1076 0.0018, so Theorem 2 covers those runs as they were executed. Over three seeds and 20000 rounds
 1077 the induction has room to spare, $\max_a x_{t+1}(a)/x_t(a)$ reaching 1.006, 1.003 and 1.002 against the
 1078 permitted 1.100, 1.020 and 1.005.

1079 The two bounds are numerically close on the funnel, where $\bar{v} \approx 2$: the additive form carries $3.3(\bar{v} -$
 1080 $1) + 2d$ and the interleaved one $(d+1)\bar{v}$, and these cross near $\bar{v} = 2$. The additive form gains as $\bar{v}$
 1081 rises toward $|S|$, and it describes the algorithm that is actually run, which is the reason to prefer it.

1082 L Stationary sensors

1083 *Observation 19.* The multiplicative term $\sqrt{\min\{|S|, d\}T}$ is not forced by the known construction,
 1084 and is absent at one endpoint of the present model class.

1085 The $\Omega(\sqrt{\min\{|S|, d\}T})$ construction of Esposito et al. (2023, Thm. 5.2) re-splits states across ac-
 1086 tions adversarially per block, so its action-to-state map is time-varying and lies outside the model
 1087 class of Section 3. At $E_2 = 0$ with stationary P the problem is a stochastic delayed bandit, where
 1088 delay enters additively (Joulani et al., 2013). One route to it is therefore unavailable and the term is
 1089 absent at one endpoint, but neither a matching upper bound for all $E_2 > 0$ nor a lower bound ruling
 1090 it out is established here.